

Table of Contents

Apache Kafka Introduction	2
Import configuration parameter.....	17
Producer API	20
Apache kafka Producer workflow	24
Callback and Acknowledgment.....	24
Custom Partition	27
Custom Serializer	33
Producer Config	39
Consumer Group.....	46
Consumer API.....	54
Offset Management.....	58
Rebalance Listener.....	70
Exactly one processing.....	78
Schema Evolution.....	87
Confluent Platform	98
Confluent Schema Registry	99
Installation Step on Centos7	100

Apache Kafka Introduction

https://www.youtube.com/watch?v=udnX21_SuU&list=PLkz1SCf5iB4enAR00Z46JwY9GGkaS2NON&index=3

Apache Kafka Tutorial

What is a producer?

An application that sends messages to Kafka.

Apache Kafka Tutorial

What is a message?

Small to medium sized piece of data.

Apache Kafka Tutorial

What is a Message?

A text file



Apache Kafka Tutorial

What is a Message?

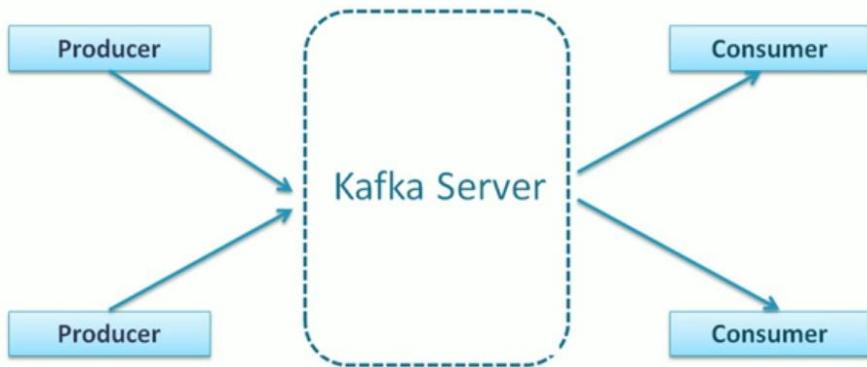
ID	Value	Description	
1	Producer	An application that sends data to Kafka	Message 1
2	Consumer	An application that receives data from Kafka	
3	Broker	Kafka Server	
4	Cluster	Group of computers	
5	Topic	A name for a Kafka stream	Message 5
6	Partitions	Part of a topic	
7	Offset	Unique id for a message within a partition	
8	Consumer groups	A group of consumers acting as a single logical unit.	Message 8

Apache Kafka Tutorial

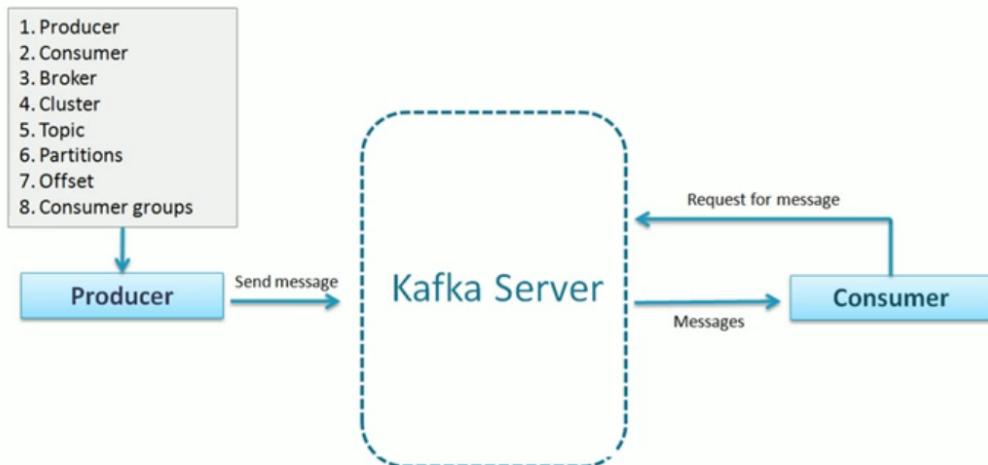
What is a consumer?

An application that reads data from Kafka

Apache Kafka Tutorial



Apache Kafka Tutorial



Apache Kafka Tutorial

What is a cluster?

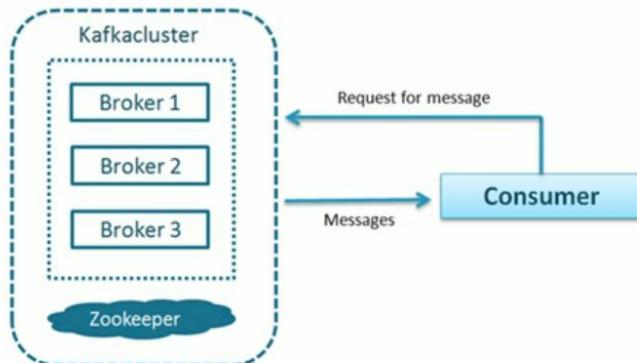
A group of computers sharing workload for a common purpose

Apache Kafka Tutorial

- 1. Producer
- 2. Consumer
- 3. Broker
- 4. Cluster
- 5. Topic
- 6. Partitions
- 7. Offset
- 8. Consumer groups

Producer

Send message

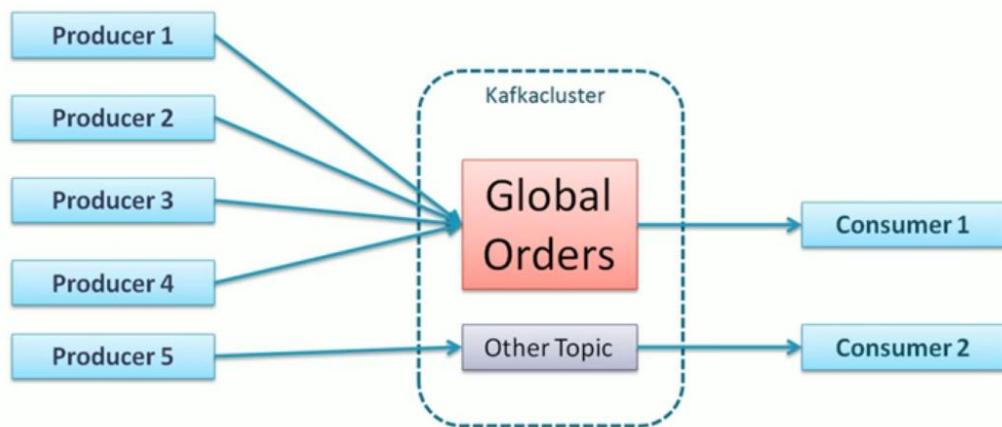


Apache Kafka Tutorial

What is a Topic?

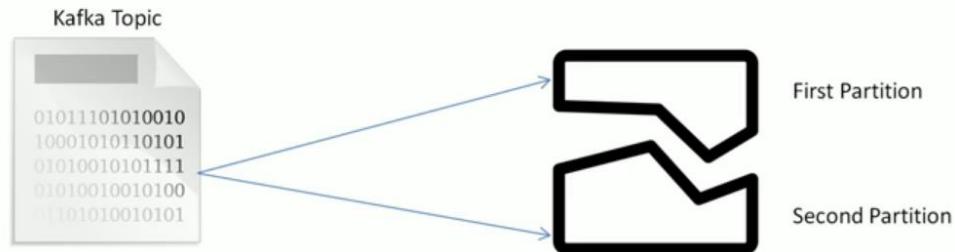
A topic is a unique name for Kafka stream

Apache Kafka Tutorial



Apache Kafka Tutorial

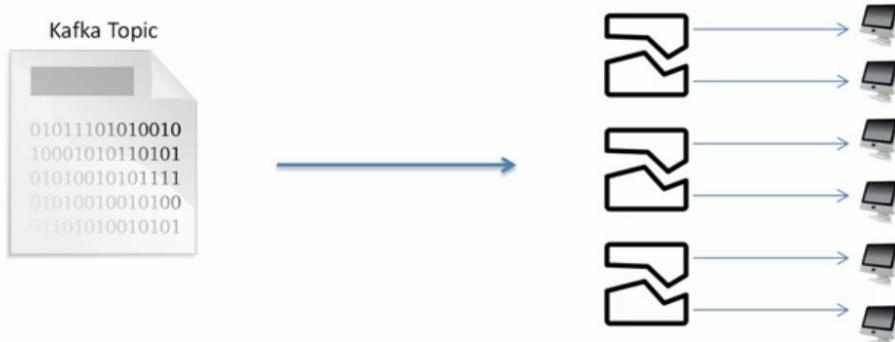
What is a Partition?



If data size is too big one computer can't store it so kafka split the data according to partition given and store it in different machines

Apache Kafka Tutorial

How many partitions?

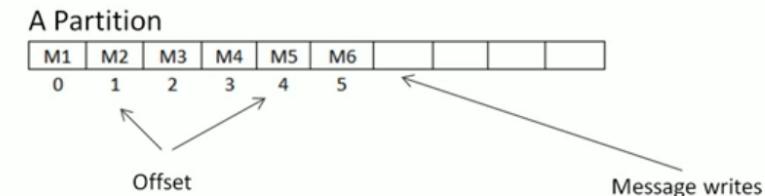


We need to decide it while creating topic

Apache Kafka Tutorial

What is offset?

A sequence id given to messages as they arrive in a partition.



Apache Kafka Tutorial

Global unique identifier of a message?

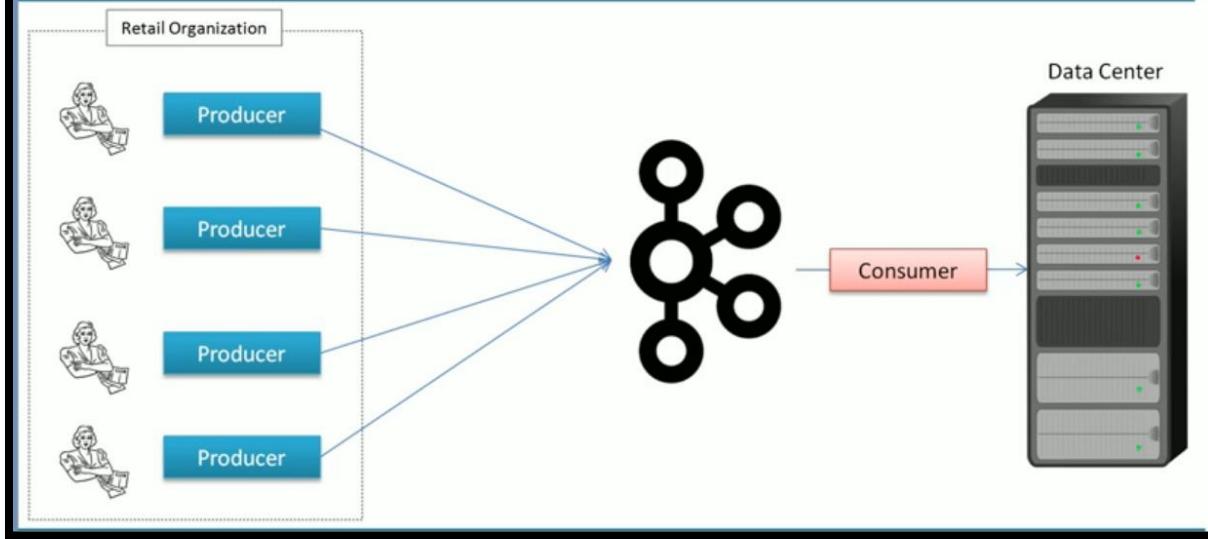
Topic Name → Partition Number → Offset

Apache Kafka Tutorial

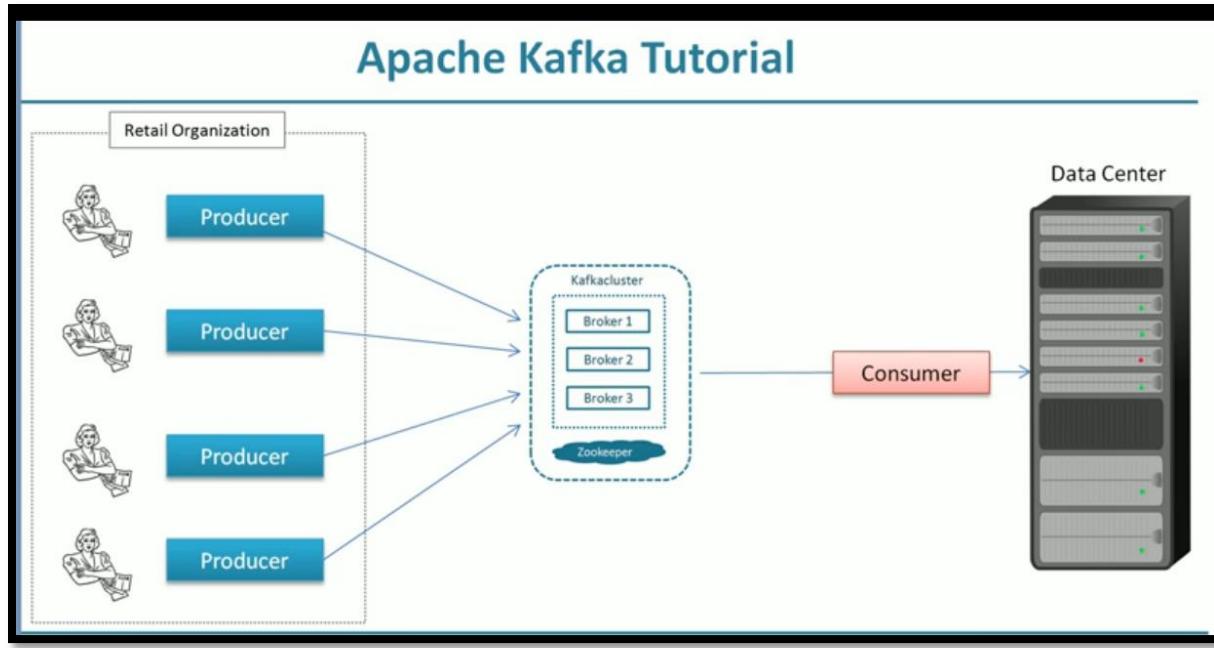
What is a consumer group?

A group of consumers acting as a single logical unit

Apache Kafka Tutorial

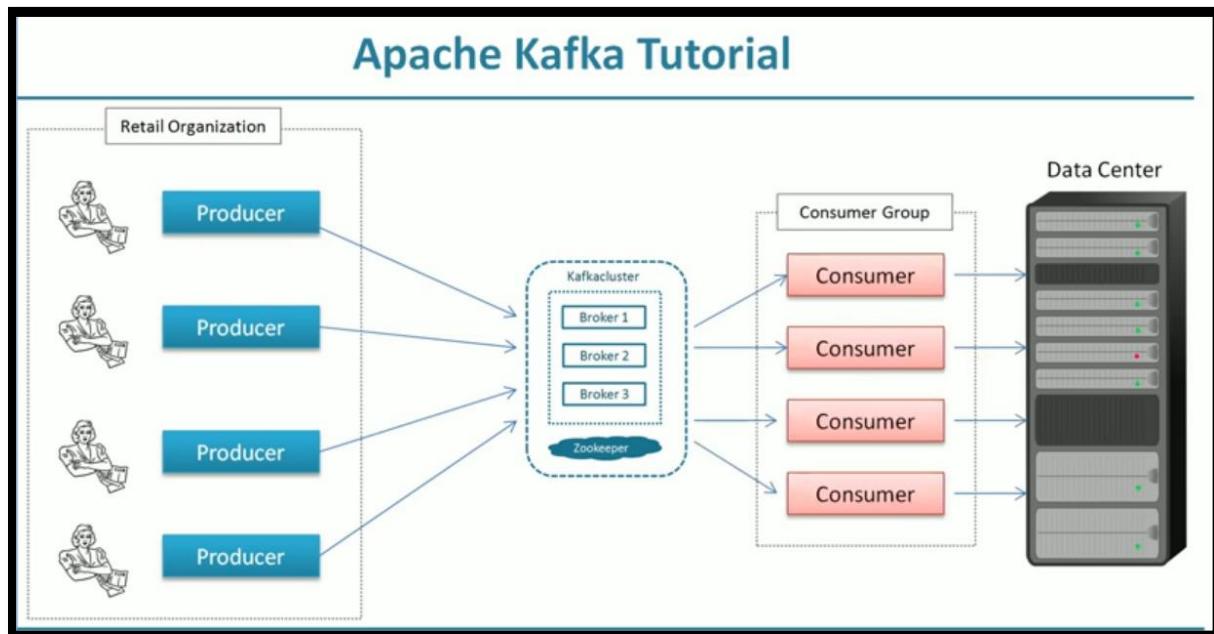


Hundreds of producers pushing data into single topic how we will handle the data volume and velocity



Kafka cluster and partitioned the topic

Now source side we have many producers but at destination side just one consumer to receive data and process data

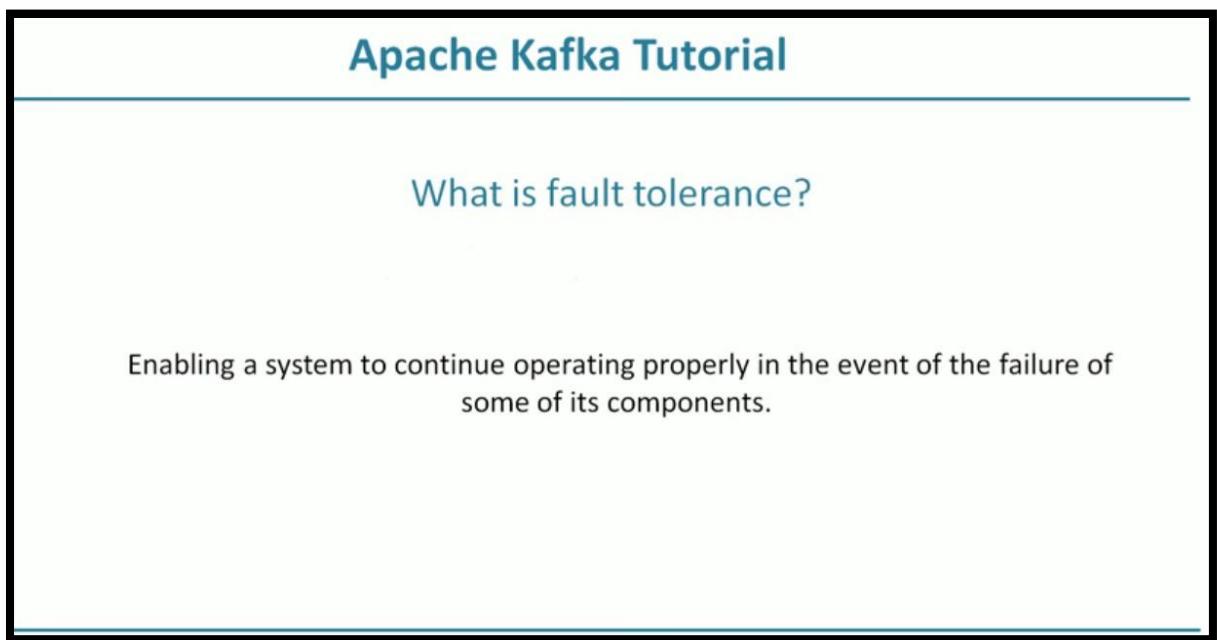
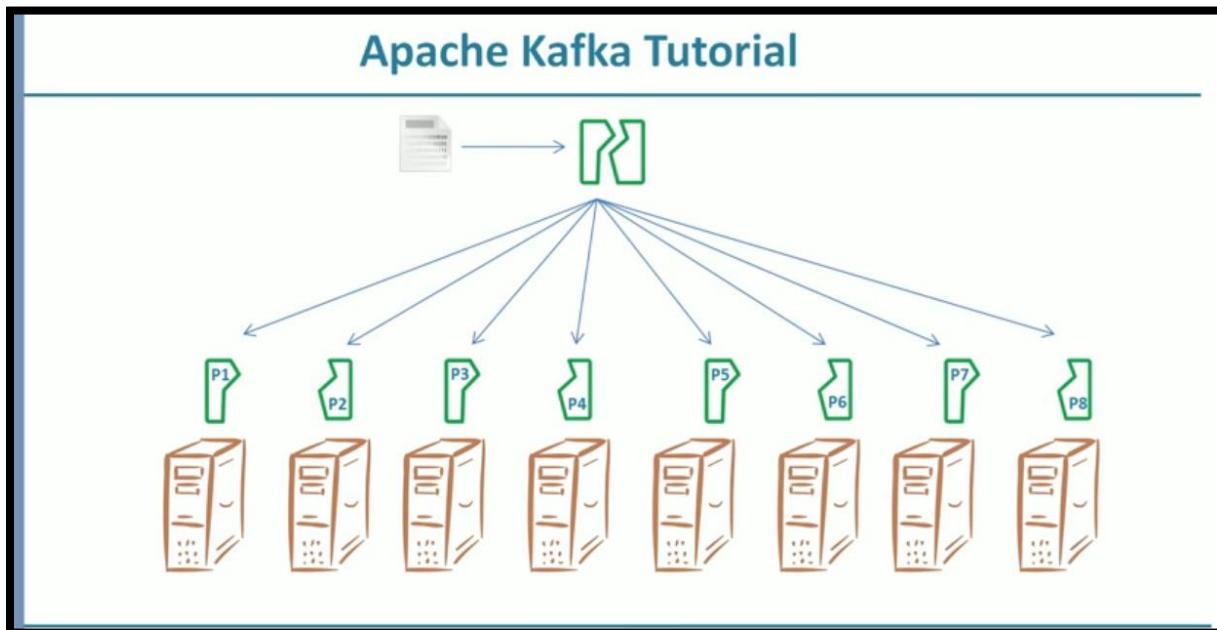


Multiple consumer divide works among them, each consumer can work on some partition

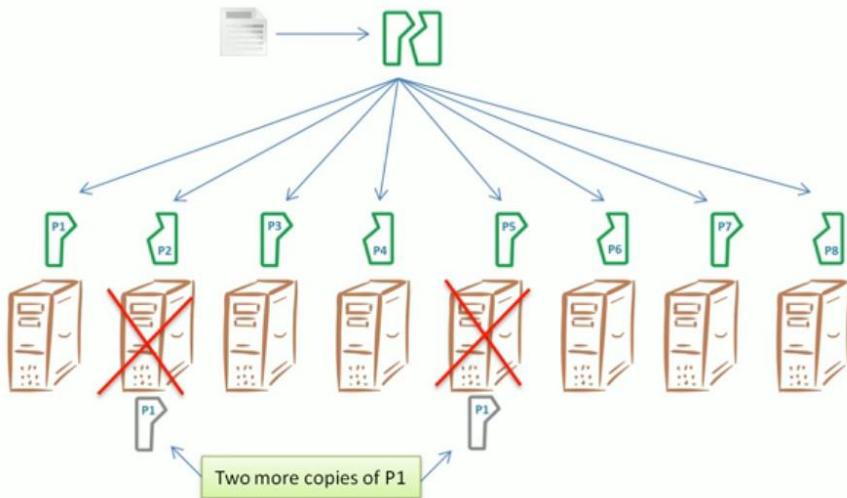
Suppose we have 12 partition and 2 kafka consumer under Consumer group , they will both read 6 partition each means they will read different set of partitions=different set of messages

Max no of consumer in a group = total no of partition you have on the topic

Kafka doesn't allow two consumers to read from the same partition simultaneously to avoid double reading of record



Apache Kafka Tutorial



Make multiple copies of data and store them in different system

Apache Kafka Tutorial

What is replication factor?

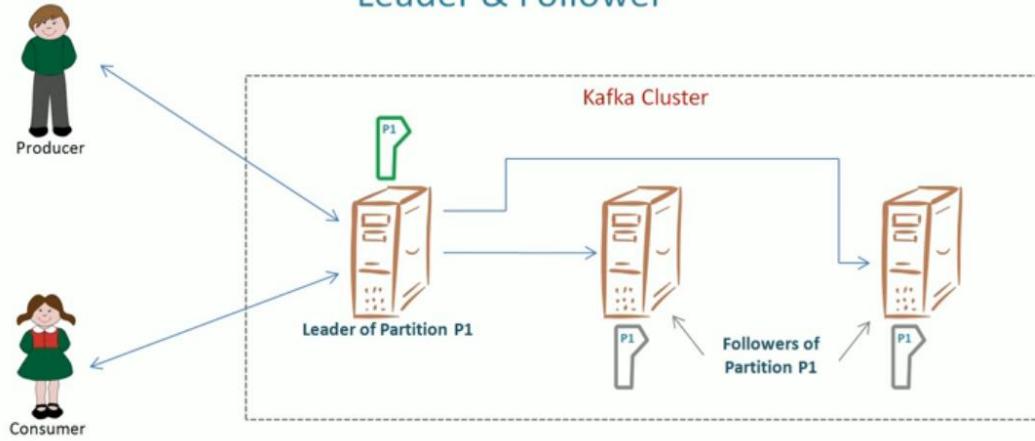
Number of total copies.



Replication factor is define for the topics not for partitions

Apache Kafka Tutorial

Leader & Follower



Kafka implements a leader and follower model so for every partition one broker is elected as a leader and leader takes care of all client interaction

It is leader's responsibility to receive the message, store it in local disk and send back an acknowledgement to the producer. Similarly when a consumer is willing to read data it sends a request to the leader, its responsibility to send requested data back to the consumer.

Follower maintains the copies of partition, these followers copy the data from the leader they don't talk to producer or consumer

Start three brokers in single machine just for demo purpose

Apache Kafka Tutorial

First Broker

bin/kafka-server-start.sh config/server.properties

Second Broker

bin/kafka-server-start.sh config/server-1.properties

Third Broker

bin/kafka-server-start.sh config/server-2.properties

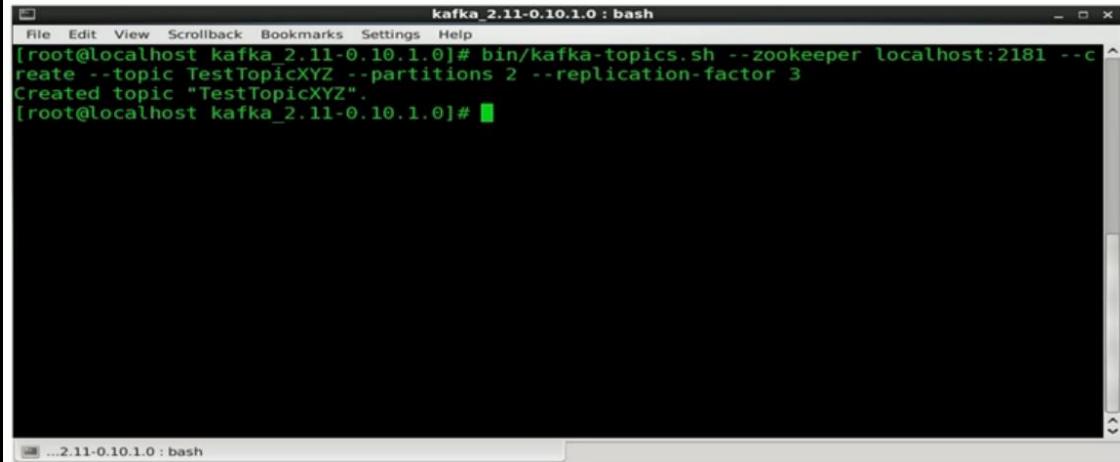
Apache Kafka Tutorial

1. broker.id
2. Port
3. log.dirs



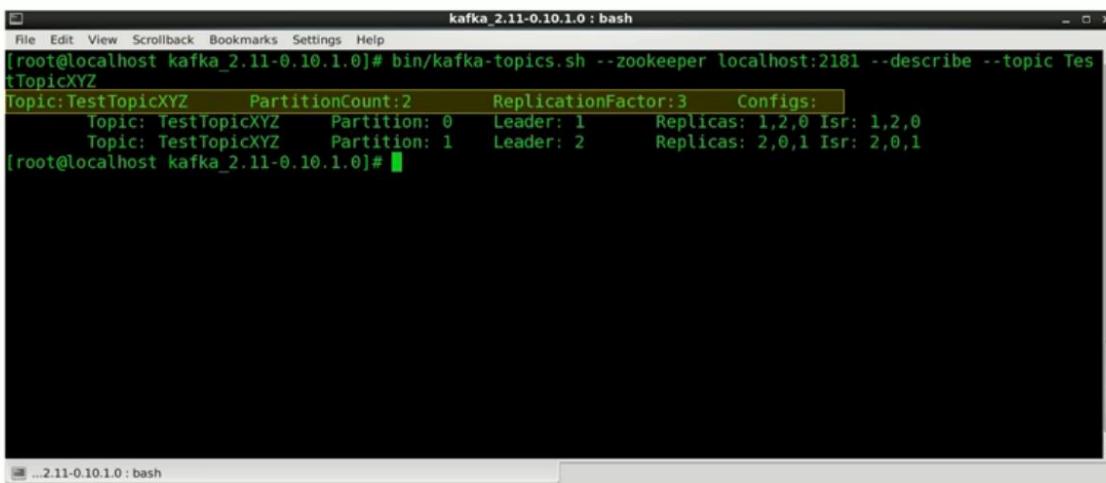
Change these three properties in each server.properties file

Apache Kafka Tutorial



```
[root@localhost kafka_2.11-0.10.1.0]# bin/kafka-topics.sh --zookeeper localhost:2181 --create --topic TestTopicXYZ --partitions 2 --replication-factor 3
Created topic "TestTopicXYZ".
[root@localhost kafka_2.11-0.10.1.0]#
```

Apache Kafka Tutorial



```
[root@localhost kafka_2.11-0.10.1.0]# bin/kafka-topics.sh --zookeeper localhost:2181 --describe --topic TestTopicXYZ
Topic:TestTopicXYZ    PartitionCount:2    ReplicationFactor:3    Configs:
          Topic: TestTopicXYZ    Partition: 0    Leader: 1    Replicas: 1,2,0 ISR: 1,2,0
          Topic: TestTopicXYZ    Partition: 1    Leader: 2    Replicas: 2,0,1 ISR: 2,0,1
[root@localhost kafka_2.11-0.10.1.0]#
```

- For first partition broker1 is the leader and for second broker 2 is the leader
- Replicas shows the three copies of partitioned maintained by different brokers
- ISR is list of in synch replicas

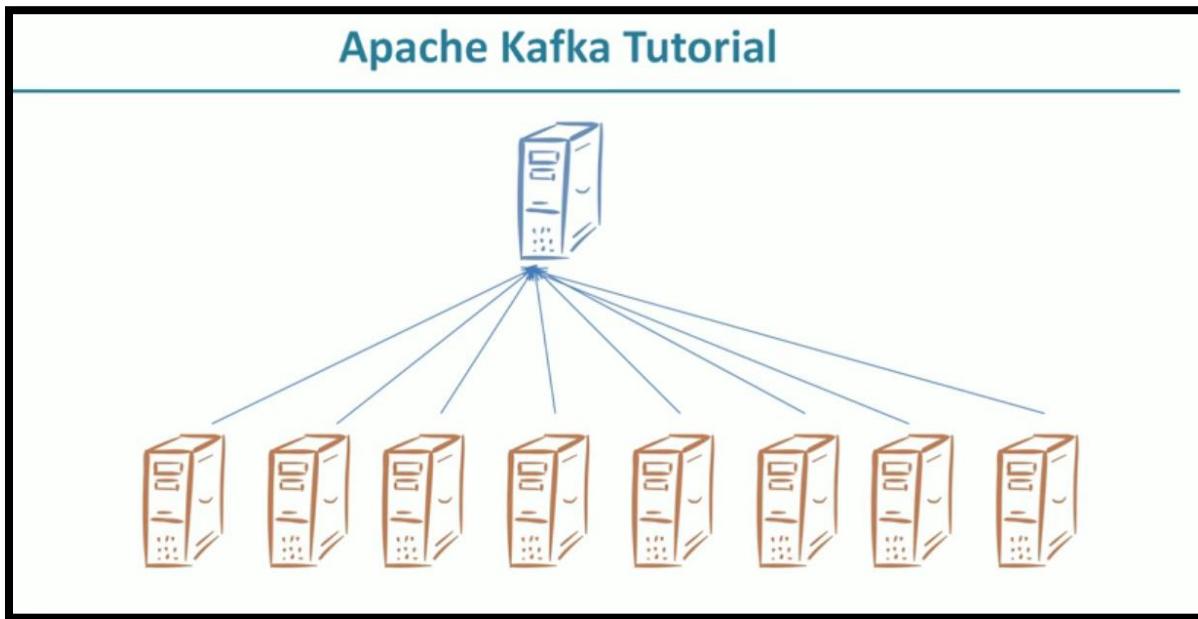
Import configuration parameter

Apache Kafka Tutorial



1. zookeeper.connect

Zookeeper address that manages Kafka cluster



Apache Kafka Tutorial



1. `delete.topic.enable`

You cannot remove a topic, by default value of `delete.topic.enable` is false, but we can change this to true in development and testing environment to delete the topic by command line tool

Apache Kafka Tutorial

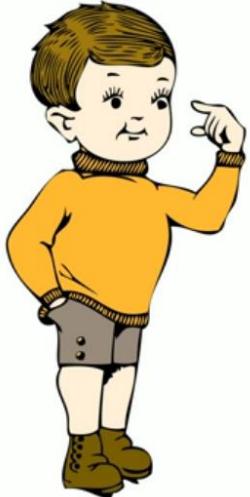


1. `auto.create.topics.enable`

If this parameter is set to true Kafka create topics automatically if producers sending message on non-existence topic

Set this parameter to false then kafka will stop creating topic automatically

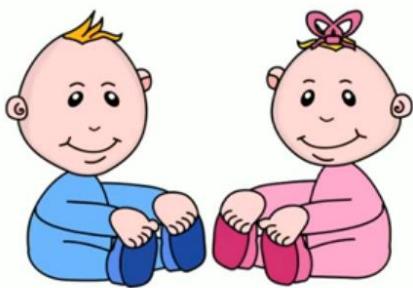
Apache Kafka Tutorial



1. default.replication.factor
2. num.partitions

Default value for both of them is one, if kafka creates topics automatically replication factor and partition are set to 1

Apache Kafka Tutorial



1. log.retention.ms
2. log.retention.bytes

Kafka retain the data for certain amount of time and size, we can configure these two properties to configure them, default retention period is 7 days (in latest version we have **log.retention.hours**)
log.retention.bytes applies to partition level (not a topic size) so if size of the partition reach up to this parameter kafka delete that partition

Producer API

Apache Kafka Tutorial

➡ <https://kafka.apache.org/documentation#api>

2.1 Producer API

The Producer API allows applications to send streams of data to topics in the Kafka cluster.

Examples showing how to use the producer are given in the javadocs.

To use the producer, you can use the following maven dependency:

```
<dependency>
  <groupId>org.apache.kafka</groupId>
  <artifactId>kafka-clients</artifactId>
  <version>0.10.1.0</version>
</dependency>
```

2.2 Consumer API

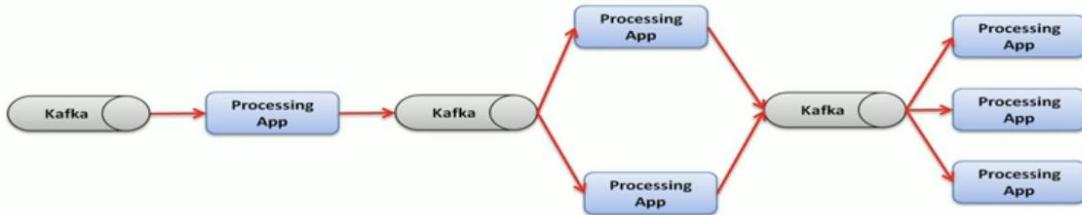
The Consumer API allows applications to read streams of data from topics in the Kafka cluster.

Examples showing how to use the consumer are given in the javadocs.

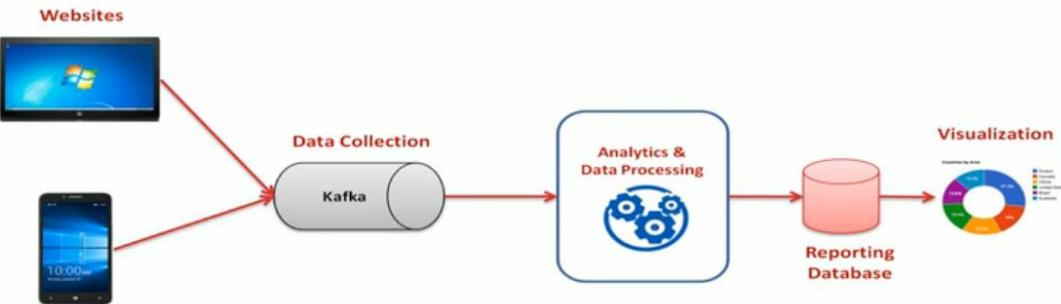
To use the consumer, you can use the following maven dependency:

```
<dependency>
  <groupId>org.apache.kafka</groupId>
  <artifactId>kafka-clients</artifactId>
  <version>0.10.1.0</version>
</dependency>
```

Apache Kafka Tutorial



Apache Kafka Tutorial



Apache Kafka Tutorial

<https://kafka.apache.org/0101/javadoc/index.html?org/apache/kafka/clients/producer/KafkaProducer.html>

Overview Package **Class** Tree Deprecated Index Help

Prev Class Next Class Frames No Frames
Summary: Nested | Field | Constr | Method Detail: Field | Constr | Method

org.apache.kafka.clients.producer

Class KafkaProducer<K,V>

java.lang.Object
org.apache.kafka.clients.producer.KafkaProducer<K,V>

All Implemented Interfaces:

Closeable, AutoCloseable, Producer<K,V>

Get the code from below URL

<https://github.com/LearningJournal/ApacheKafkaTutorials/tree/master/ProducerExamples>

```
import java.util.*;  
import org.apache.kafka.clients.producer.*;  
public class SimpleProducer {
```

```

public static void main(String[] args) throws Exception{

    String topicName = "SimpleProducerTopic";
    // generally we send key and value however key is not mandatory
    String key = "Key1";
    String value = "Value-1";

    Properties props = new Properties();
    //list of kafka brokers
    props.put("bootstrap.servers", "localhost:9092,localhost:9093");
    // type of key
    props.put("key.serializer", "org.apache.kafka.common.serialization.StringSerializer");
    // type of value
    props.put("value.serializer", "org.apache.kafka.common.serialization.StringSerializer");

    // type of key.value
    Producer<String, String> producer = new KafkaProducer <>(props);
    //this object is our message
    ProducerRecord<String, String> record = new ProducerRecord<>(topicName, key, value);
    producer.send(record);
    producer.close();

    System.out.println("SimpleProducer Completed.");
}

}

```

Apache Kafka Tutorial

Constructor Summary

Constructors
Constructor and Description
KafkaProducer(Map<String, Object> configs)
A producer is instantiated by providing a set of key-value pairs as configuration.
KafkaProducer(Map<String, Object> configs, Serializer<K> keySerializer, Serializer<V> valueSerializer)
A producer is instantiated by providing a set of key-value pairs as configuration, a key and a value Serializer .
KafkaProducer(Properties properties)
A producer is instantiated by providing a set of key-value pairs as configuration.
KafkaProducer(Properties properties, Serializer<K> keySerializer, Serializer<V> valueSerializer)
A producer is instantiated by providing a set of key-value pairs as configuration, a key and a value Serializer .

Apache Kafka Tutorial

Constructor Summary

Constructors

Constructor and Description

`ProducerRecord(String topic, Integer partition, K key, V value)`

Creates a record to be sent to a specified topic and partition

`ProducerRecord(String topic, Integer partition, Long timestamp, K key, V value)`

Creates a record with a specified timestamp to be sent to a specified topic and partition

`ProducerRecord(String topic, K key, V value)`

Create a record to be sent to Kafka

`ProducerRecord(String topic, V value)`

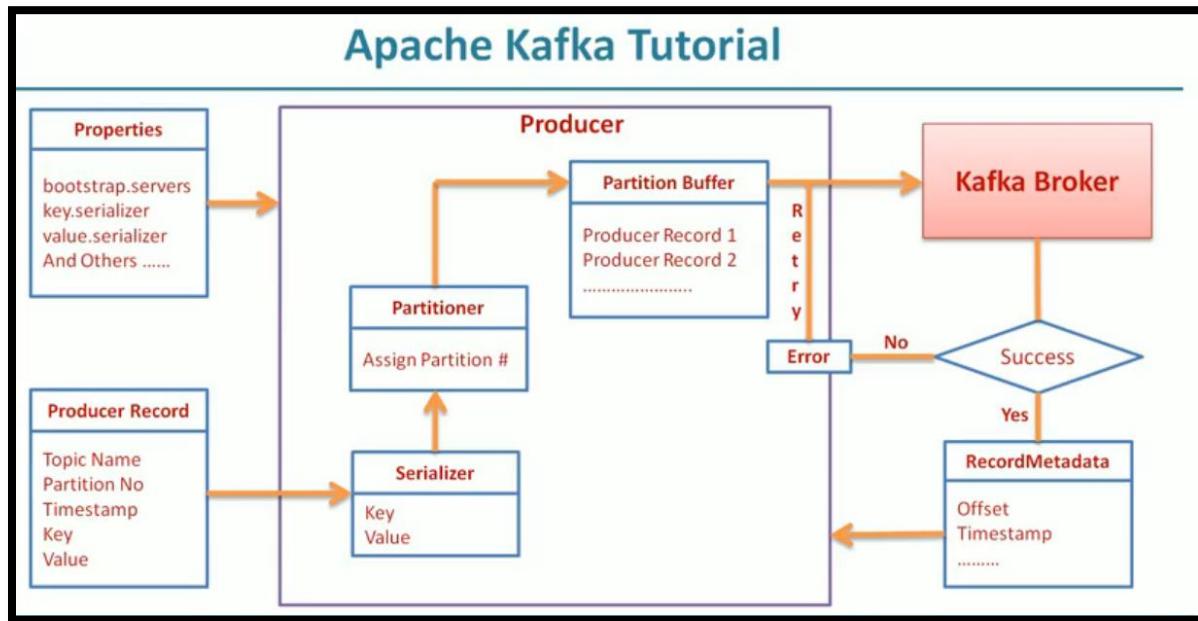
Create a record with no key

If you send multiple messages with same message key then they all will land to same partition but as key is an optional field in message so if we don't send keys with message , default partition eventually distributed messages across the partition.

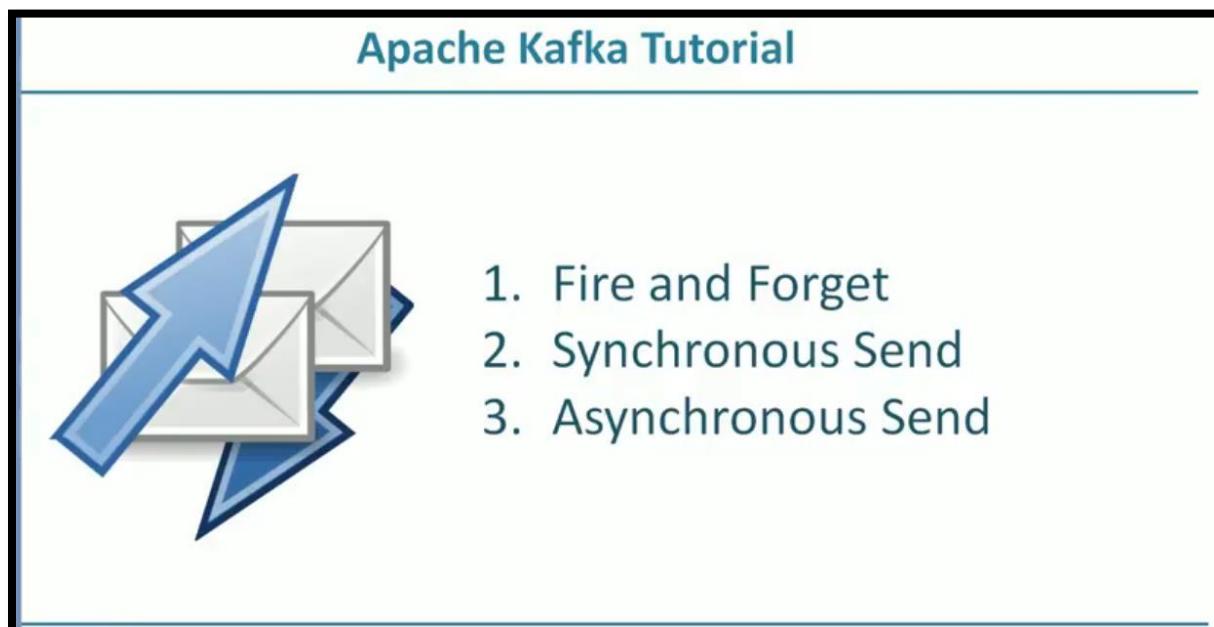
If we set the partition number in producerRecord, it disable default partitioner

We can set timestamp when we are sending the message if we don't set the timestamp would be the time when broker receive the message

Apache Kafka Producer workflow



Callback and Acknowledgment



Three approaches to send message to Kafka broker

1. **Fire and forget** – send message to broker and don't care if it is successful received by broker or not, in this approach you may miss some messages. Like in sentimental analysis we don't care about if we miss some tweet

2. **Synchronous send** – send message and wait for response, either get metadata record or exception, here we only care for exception as we may want to log the exception. If messages are critical and can't afford to miss anything, it will limit your throughput as we are waiting for response.

```
import java.util.*;
import org.apache.kafka.clients.producer.*;
public class SynchronousProducer {

    public static void main(String[] args) throws Exception{

        String topicName = "SynchronousProducerTopic";
        String key = "Key1";
        String value = "Value-1";

        Properties props = new Properties();
        props.put("bootstrap.servers", "localhost:9092,localhost:9093");
        props.put("key.serializer", "org.apache.kafka.common.serialization.StringSerializer");
        props.put("value.serializer", "org.apache.kafka.common.serialization.StringSerializer");

        Producer<String, String> producer = new KafkaProducer <>(props);

        ProducerRecord<String, String> record = new ProducerRecord<>(topicName,key,value);

        try{
            RecordMetadata metadata = producer.send(record).get();
            System.out.println("Message is sent to Partition no " + metadata.partition() + " and offset " +
metadata.offset());
            System.out.println("SynchronousProducer Completed with success.");
        }catch (Exception e) {
            e.printStackTrace();
            System.out.println("SynchronousProducer failed with an exception");
        }finally{
            producer.close();
        }
    }
}
```

3. **Asynchronous send** – in this approach we send message and provide a callback function to receive an acknowledgement, we don't wait for success or failure, the producer will callback our function with recordmeta and an exception object.

```

import java.util.*;
import org.apache.kafka.clients.producer.*;

public class AsynchronousProducer {

    public static void main(String[] args) throws Exception{
        String topicName = "AsynchronousProducerTopic";
        String key = "Key1";
        String value = "Value-1";

        Properties props = new Properties();
        props.put("bootstrap.servers", "localhost:9092,localhost:9093");
        props.put("key.serializer", "org.apache.kafka.common.serialization.StringSerializer");
        props.put("value.serializer", "org.apache.kafka.common.serialization.StringSerializer");

        Producer<String, String> producer = new KafkaProducer <>(props);

        ProducerRecord<String, String> record = new ProducerRecord<>(topicName, key, value);

        producer.send(record, new MyProducerCallback());
        System.out.println("AsynchronousProducer call completed");
        producer.close();

    }

}

class MyProducerCallback implements Callback{

    @Override
    public void onCompletion(RecordMetadata recordMetadata, Exception e) {
        if (e != null)
            System.out.println("AsynchronousProducer failed with an exception");
        else
            System.out.println("AsynchronousProducer call Success:");
    }
}

```

max.in.flight.requests.per.connection – tell how many message you can send to broker without receiving acknowledgement, the default value is 5

Custom Partition

Apache Kafka Tutorial



Default Partitioner

1. If a partition is specified in the record, use it
2. If no partition is specified but a key is present choose a partition based on a hash of the key
3. If no partition or key is present choose a partition in a round-robin fashion

Apache Kafka Tutorial

Hashing

Hash two Keys SSP8 & SSP3 using below code taken from Default Partitioner.

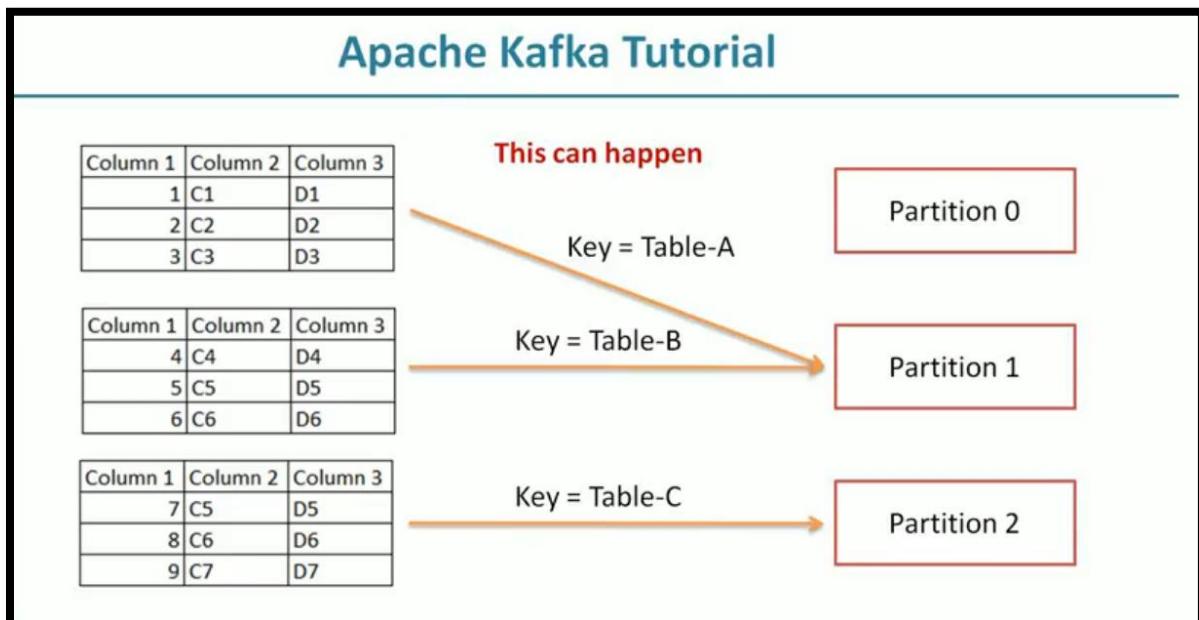
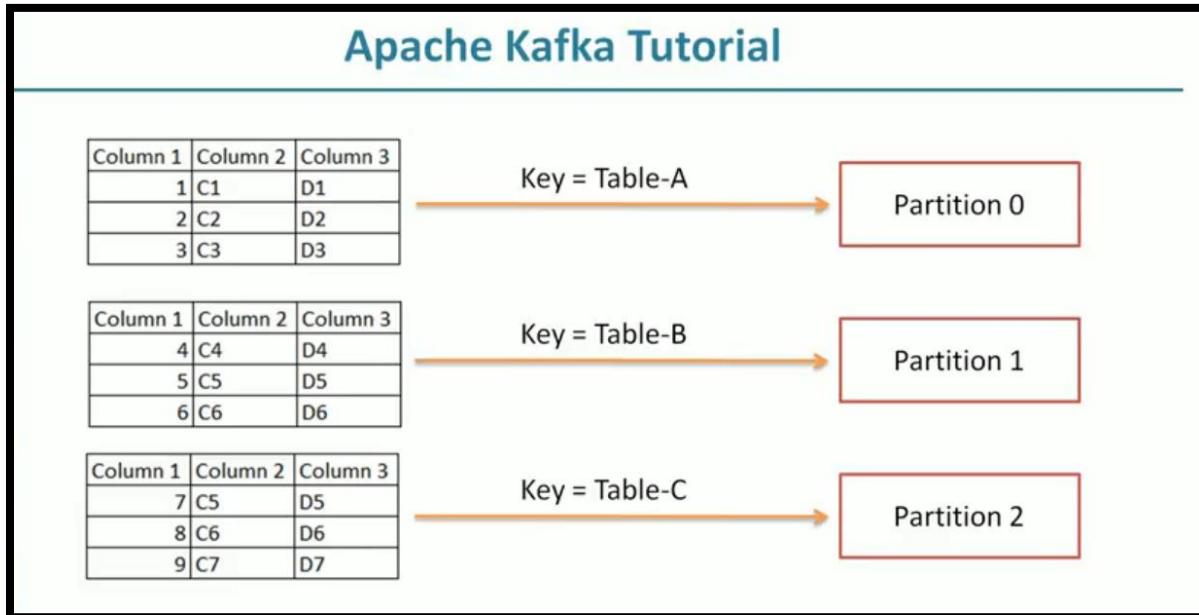
```
Utils.toPositive(Utils.murmur2(keyBytes))
```

Hopefully, you will get the same number

Same key will go to the same partition but this also may happen that two distinguish keys can go to same partition

If we want to create partitions in such a way that data from Table-A go to Partition-0, Table-B to Partition-1 and so on.

Here we are sending table name as a key



Apache Kafka Tutorial

This is safer

Column 1	Column 2	Column 3
1 C1		D1
2 C2		D2
3 C3		D3

Partition = 0

Partition 0

Column 1	Column 2	Column 3
4 C4		D4
5 C5		D5
6 C6		D6

Partition = 1

Partition 1

Column 1	Column 2	Column 3
7 C5		D5
8 C6		D6
9 C7		D7

Partition = 2

Partition 2

So better to send hardcoded partition number or we can create custom partition.

Apache Kafka Tutorial

Code snippet from Default Partitioner

```
// hash the keyBytes to choose a partition  
return Utils.toPositive(Utils.murmur2(keyBytes)) % numPartitions;
```

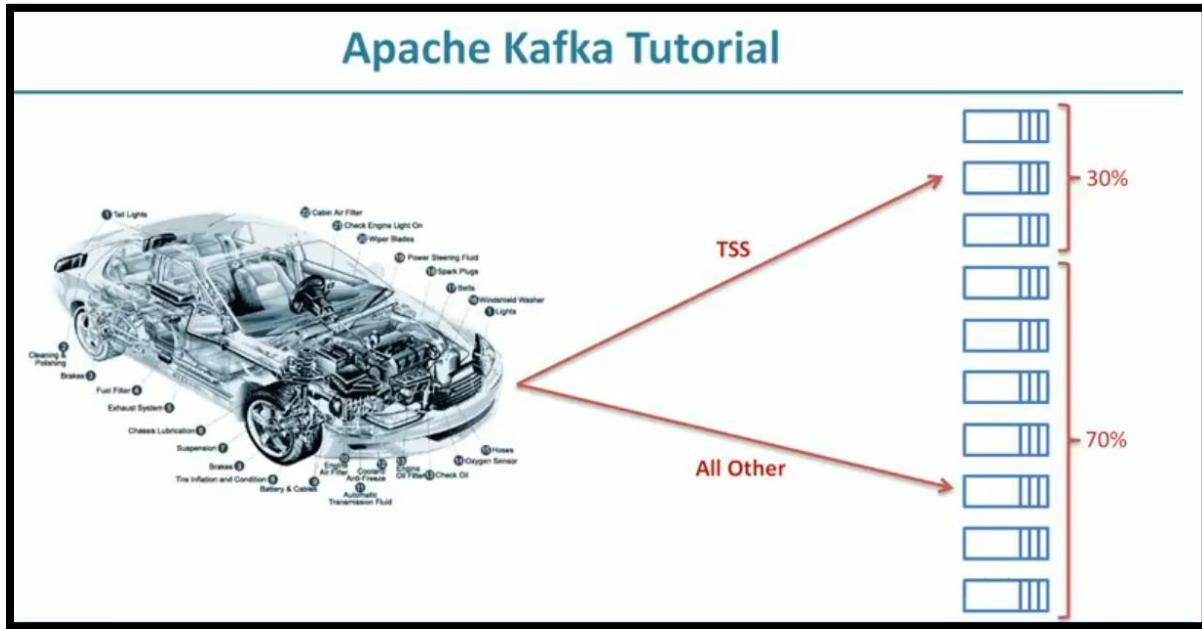
↑
Total Number of
Partitions on the Topic

Default partition return partition based on key modulo number of partition

If you are increasing the number of partition later default partition will start returning some different numbers.

So key is not a good use for making partition

Use case for custom partition



We are receiving data from different sensors and we planned 10 partitions to a single topic but we want to 3 partitions dedicated for sensor named TSS and 7 partitions for rest of the sensors.

```
import java.util.*;
import org.apache.kafka.clients.producer.*;
public class SensorProducer {

    public static void main(String[] args) throws Exception{

        String topicName = "SensorTopic";

        Properties props = new Properties();
        props.put("bootstrap.servers", "localhost:9092,localhost:9093");
        props.put("key.serializer", "org.apache.kafka.common.serialization.StringSerializer");
        props.put("value.serializer", "org.apache.kafka.common.serialization.StringSerializer");
        props.put("partitioner.class", "SensorPartitioner");
        props.put("speed.sensor.name", "TSS"); // it is custom property not a kafka property

        Producer<String, String> producer = new KafkaProducer <>(props);

        for (int i=0 ; i<10 ; i++)
            producer.send(new ProducerRecord <>(topicName, "SSP"+i, "500"+i));

        for (int i=0 ; i<10 ; i++)
            producer.send(new ProducerRecord <>(topicName, "TSS", "500"+i));
    }
}
```

```

        producer.close();

        System.out.println("SimpleProducer Completed.");
    }
}

```

```

import java.util.*;
import org.apache.kafka.clients.producer.*;
import org.apache.kafka.common.*;
import org.apache.kafka.common.utils.*;
import org.apache.kafka.common.record.*;

public class SensorPartitioner implements Partitioner {

    private String speedSensorName;

    // initialization method
    public void configure(Map<String, ?> configs) {
        speedSensorName = configs.get("speed.sensor.name").toString();
    }

    public int partition(String topic, Object key, byte[] keyBytes, Object value, byte[] valueBytes, Cluster
cluster) {

        List<PartitionInfo> partitions = cluster.partitionsForTopic(topic);
        int numPartitions = partitions.size();
        // calculate 30% of total partition for TSS sensor which is 3 in our case as total partitions are 10
        int sp = (int) Math.abs(numPartitions * 0.3);
        int p=0;

        if ( (keyBytes == null) || (!(key instanceof String)) )
            throw new InvalidRecordException("All messages must have sensor name as key");

        if ( ((String)key).equals(speedSensorName) )
            //here we are using value for hashing because in that case key is constant (TSS) if we give the key it will
            //always return single partition so all the TSS data goes to single partition
            p = Utils.toPositive(Utils.murmur2(valueBytes)) % sp;
        else
            p = Utils.toPositive(Utils.murmur2(keyBytes)) % (numPartitions-sp) + sp ;

        System.out.println("Key = " + (String)key + " Partition = " + p );
        return p;
    }
    // cleanup method
    public void close() {}
}

```

```
}
```

We created a topic with ten partitions

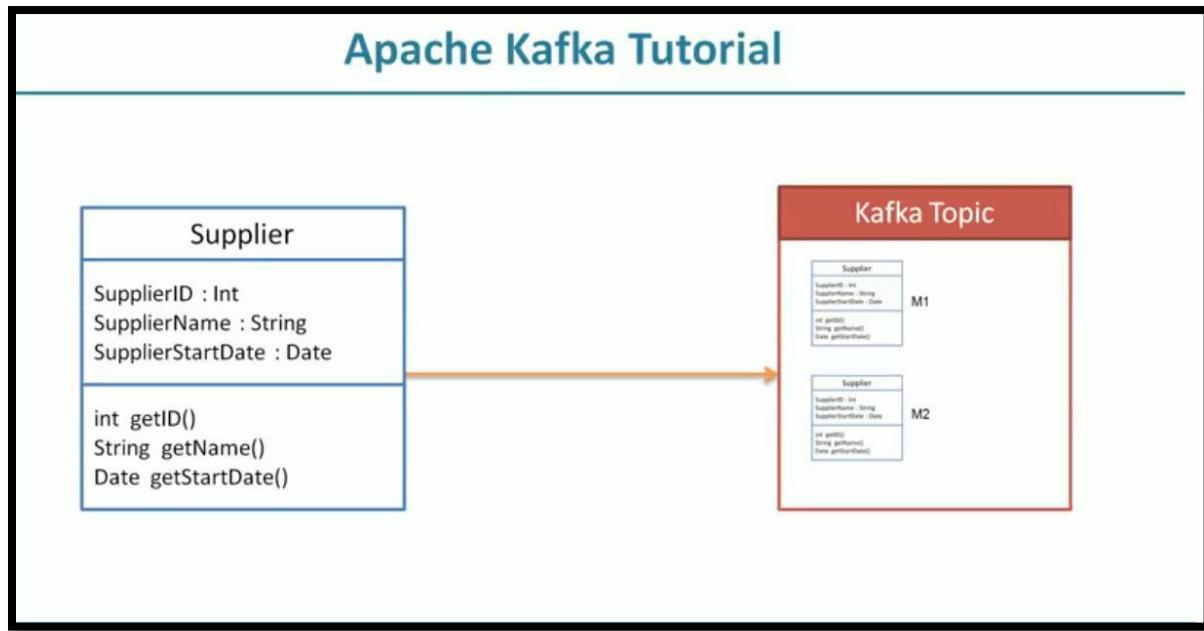
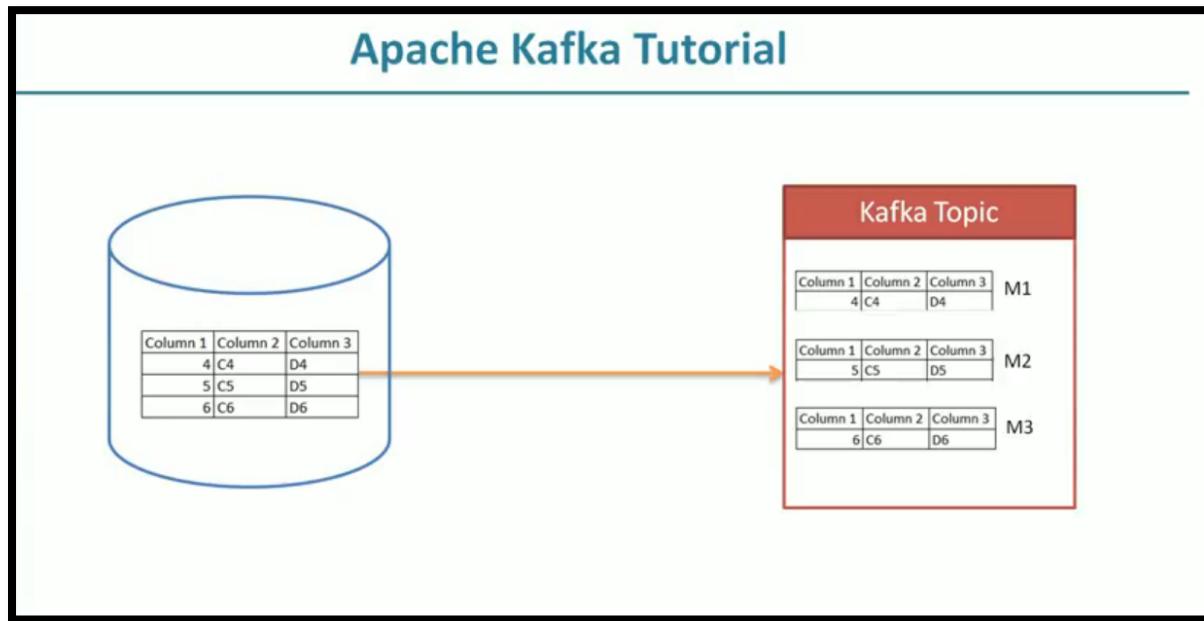
```
[root@localhost javaProj]# ~/kafka/kafka_2.11-0.10.1.0/bin/kafka-topics.sh --zookeeper localhost:2181 --describe --topic SensorTopic
Topic:SensorTopic      PartitionCount:10      ReplicationFactor:1      Configs:
  Topic: SensorTopic    Partition: 0    Leader: 1    Replicas: 1    Isr: 1
  Topic: SensorTopic    Partition: 1    Leader: 2    Replicas: 2    Isr: 2
  Topic: SensorTopic    Partition: 2    Leader: 0    Replicas: 0    Isr: 0
  Topic: SensorTopic    Partition: 3    Leader: 1    Replicas: 1    Isr: 1
  Topic: SensorTopic    Partition: 4    Leader: 2    Replicas: 2    Isr: 2
  Topic: SensorTopic    Partition: 5    Leader: 0    Replicas: 0    Isr: 0
  Topic: SensorTopic    Partition: 6    Leader: 1    Replicas: 1    Isr: 1
  Topic: SensorTopic    Partition: 7    Leader: 2    Replicas: 2    Isr: 2
  Topic: SensorTopic    Partition: 8    Leader: 0    Replicas: 0    Isr: 0
  Topic: SensorTopic    Partition: 9    Leader: 1    Replicas: 1    Isr: 1
[root@localhost javaProj]#
```

```
[1] AsynchronousProducer
[2] SensorProducer
[3] SimpleProducer
[4] SynchronousProducer
^JEnter number: 2

[info] Running SensorProducer
SLF4J: Failed to load class "org.slf4j.impl.StaticLoggerBinder".
SLF4J: Defaulting to no-operation (NOP) logger implementation
SLF4J: See http://www.slf4j.org/codes.html#StaticLoggerBinder for further details.
Key = SSP0 Partition = 4
Key = SSP1 Partition = 9
Key = SSP2 Partition = 5
Key = SSP3 Partition = 6
Key = SSP4 Partition = 3
Key = SSP5 Partition = 8
Key = SSP6 Partition = 3
Key = SSP7 Partition = 9
Key = SSP8 Partition = 6
Key = SSP9 Partition = 7
Key = TSS Partition = 1
Key = TSS Partition = 0
Key = TSS Partition = 0
Key = TSS Partition = 2
Key = TSS Partition = 0
Key = TSS Partition = 1
Key = TSS Partition = 0
Key = TSS Partition = 0
Key = TSS Partition = 1
Key = TSS Partition = 1
```

TSS is distributed in partition 0, 1 and 2 and rests are distributed 3 to 9

Custom Serializer



Sometime we need to send whole object as a message to Kafka broker instead of plain text

For that we can use custom serializer, there are some generic serializers available like Avro and protocol buffer, most of the time we will use generic serializer like Avro

Apache Kafka Tutorial

Example – Serialize & Deserialize

- ➡ • Create a supplier class.
- Create a producer.
- Create a serializer.
- Create a deserializer.
- Create a consumer.

```
import java.util.Date;
public class Supplier{
    private int supplierId;
    private String supplierName;
    private Date supplierStartDate;

    public Supplier(int id, String name, Date dt){
        this.supplierId = id;
        this.supplierName = name;
        this.supplierStartDate = dt;
    }

    public int getID(){
        return supplierId;
    }

    public String getName(){
        return supplierName;
    }

    public Date getStartDate(){
        return supplierStartDate;
    }
}
```

```
import org.apache.kafka.common.serialization.Serializer;
import org.apache.kafka.common.errors.SerializationException;
import java.io.UnsupportedEncodingException;
import java.util.Map;
import java.nio.ByteBuffer;

public class SupplierSerializer implements Serializer<Supplier> {
    private String encoding = "UTF8";

    @Override
    public void configure(Map<String, ?> configs, boolean isKey) {
        // nothing to configure
    }

    @Override
    public byte[] serialize(String topic, Supplier data) {

        int sizeOfName;
        int sizeOfDate;
        byte[] serializedName;
        byte[] serializedDate;

        try {
            if (data == null)
                return null;
            serializedName = data.getName().getBytes(encoding);
            sizeOfName = serializedName.length;
            serializedDate = data.getStartDate().toString().getBytes(encoding);
            sizeOfDate = serializedDate.length;

            ByteBuffer buf = ByteBuffer.allocate(4+4+sizeOfName+4+sizeOfDate);
            buf.putInt(data.getID());
            buf.putInt(sizeOfName);
            buf.put(serializedName);
            buf.putInt(sizeOfDate);
            buf.put(serializedDate);

            return buf.array();
        } catch (Exception e) {
            throw new SerializationException("Error when serializing Supplier to byte[]");
        }
    }

    @Override
    public void close() {
```

```

        // nothing to do
    }
}

import java.nio.ByteBuffer;
import java.util.Date;
import java.text.DateFormat;
import java.text.SimpleDateFormat;
import org.apache.kafka.common.errors.SerializationException;
import org.apache.kafka.common.serialization.Deserializer;
import java.io.UnsupportedEncodingException;
import java.util.Map;

public class SupplierDeserializer implements Deserializer<Supplier> {
    private String encoding = "UTF8";

    @Override
    public void configure(Map<String, ?> configs, boolean isKey) {
        //Nothing to configure
    }

    @Override
    public Supplier deserialize(String topic, byte[] data) {

        try {
            if (data == null){
                System.out.println("Null received at deserialize");
                return null;
            }
            ByteBuffer buf = ByteBuffer.wrap(data);
            int id = buf.getInt();

            int sizeOfName = buf.getInt();
            byte[] nameBytes = new byte[sizeOfName];
            buf.get(nameBytes);
            String serializedName = new String(nameBytes, encoding);

            int sizeOfDate = buf.getInt();
            byte[] dateBytes = new byte[sizeOfDate];
            buf.get(dateBytes);
            String dateString = new String(dateBytes,encoding);

            DateFormat df = new SimpleDateFormat("EEE MMM dd HH:mm:ss Z yyyy");

            return new Supplier(id,serializedName,df.parse(dateString));
        }
    }
}

```

```

        } catch (Exception e) {
            throw new SerializationException("Error when deserializing byte[] to Supplier");
        }
    }

    @Override
    public void close() {
        // nothing to do
    }
}

import java.util.*;
import java.text.DateFormat;
import java.text.SimpleDateFormat;
import org.apache.kafka.clients.producer.*;
public class SupplierProducer {

    public static void main(String[] args) throws Exception{

        String topicName = "SupplierTopic";

        Properties props = new Properties();
        props.put("bootstrap.servers", "localhost:9092,localhost:9093");
        props.put("key.serializer", "org.apache.kafka.common.serialization.StringSerializer");
        props.put("value.serializer", "SupplierSerializer");

        Producer<String, Supplier> producer = new KafkaProducer <>(props);

        DateFormat df = new SimpleDateFormat("yyyy-MM-dd");
        Supplier sp1 = new Supplier(101,"Xyz Pvt Ltd.",df.parse("2016-04-01"));
        Supplier sp2 = new Supplier(102,"Abc Pvt Ltd.",df.parse("2012-01-01"));

        producer.send(new ProducerRecord<String,Supplier>(topicName,"SUP",sp1).get();
        producer.send(new ProducerRecord<String,Supplier>(topicName,"SUP",sp2).get();

        System.out.println("SupplierProducer Completed.");
        producer.close();

    }

    import java.util.*;
    import org.apache.kafka.clients.consumer.KafkaConsumer;
    import org.apache.kafka.clients.consumer.ConsumerRecords;
    import org.apache.kafka.clients.consumer.ConsumerRecord;

    public class SupplierConsumer{

```

```

public static void main(String[] args) throws Exception{

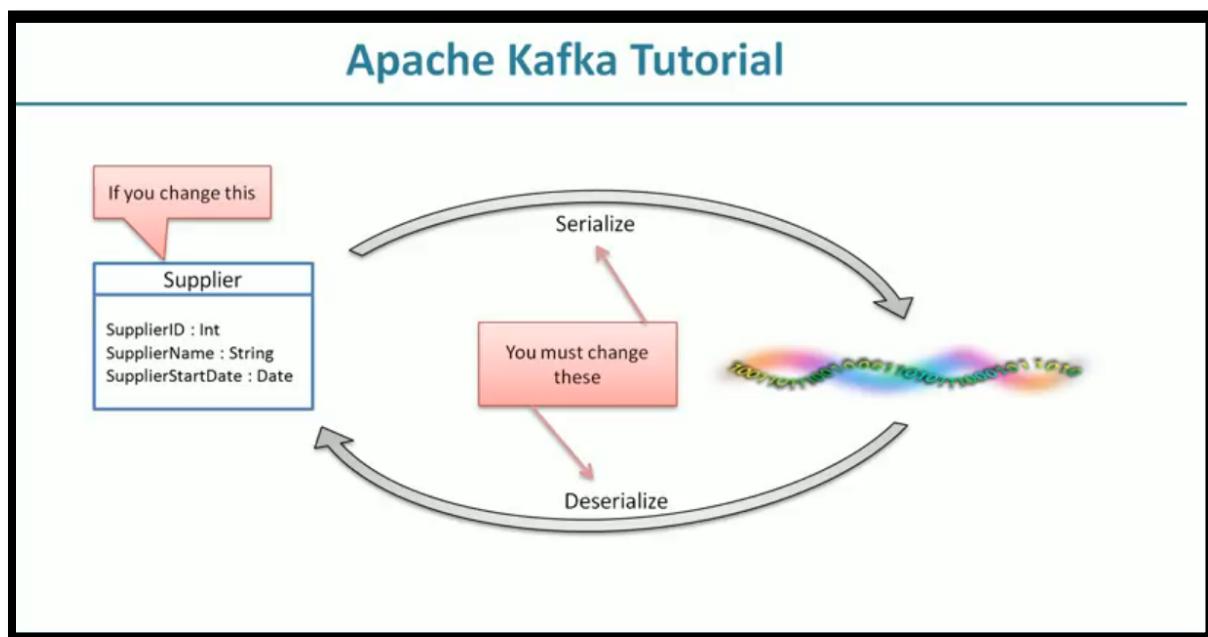
    String topicName = "SupplierTopic";
    String groupName = "SupplierTopicGroup";

    Properties props = new Properties();
    props.put("bootstrap.servers", "localhost:9092,localhost:9093");
    props.put("group.id", groupName);
    props.put("key.deserializer", "org.apache.kafka.common.serialization.StringDeserializer");
    props.put("value.deserializer", "SupplierDeserializer");

    KafkaConsumer<String, Supplier> consumer = new KafkaConsumer<>(props);
    consumer.subscribe(Arrays.asList(topicName));

    while (true){
        ConsumerRecords<String, Supplier> records = consumer.poll(100);
        for (ConsumerRecord<String, Supplier> record : records){
            System.out.println("Supplier id= " + String.valueOf(record.value().getID()) + " Supplier
Name = " + record.value().getName() + " Supplier Start Date = " +
record.value().getStartDate().toString());
        }
    }
}

```



If we change the supplier by adding some new fields we need to change our serializer and deserializer class but problem of this approach is we can't read previous serialized object , so generic serializer like Avro is good option for this case.

Producer Config

Follow this link to get full list of configuration parameters for producer

<https://kafka.apache.org/documentation/#producerconfigs>

Apache Kafka Tutorial

Kafka Producer Configuration



- bootstrap.servers
- key.serializer
- value.serializer
- partitioner.class

Apache Kafka Tutorial

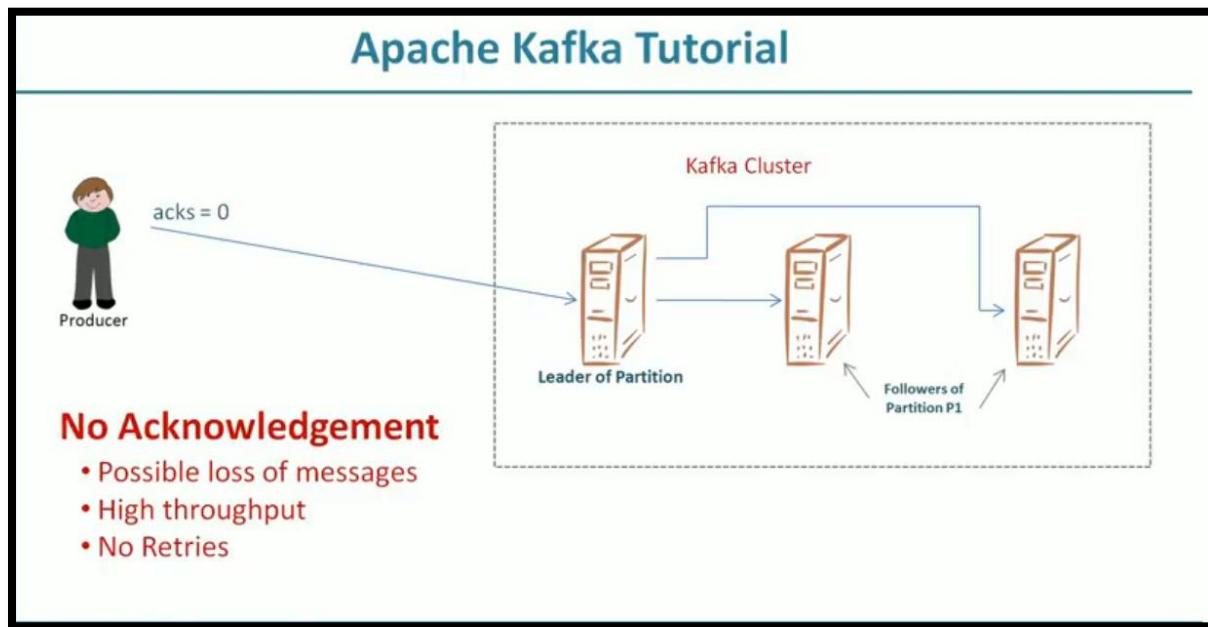
Kafka Producer Configuration



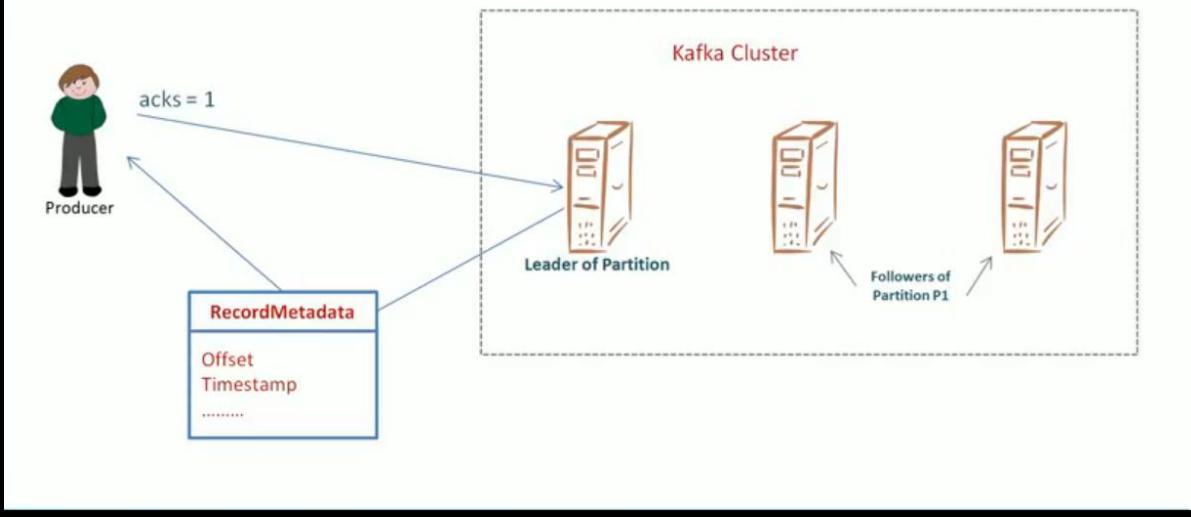
- acks
- retries
- max.in.flight.requests.per.connection

#acks – The number of acknowledgments the producer requires the leader to have received before considering a request complete. This controls the durability of records that are sent. The following settings are allowed:

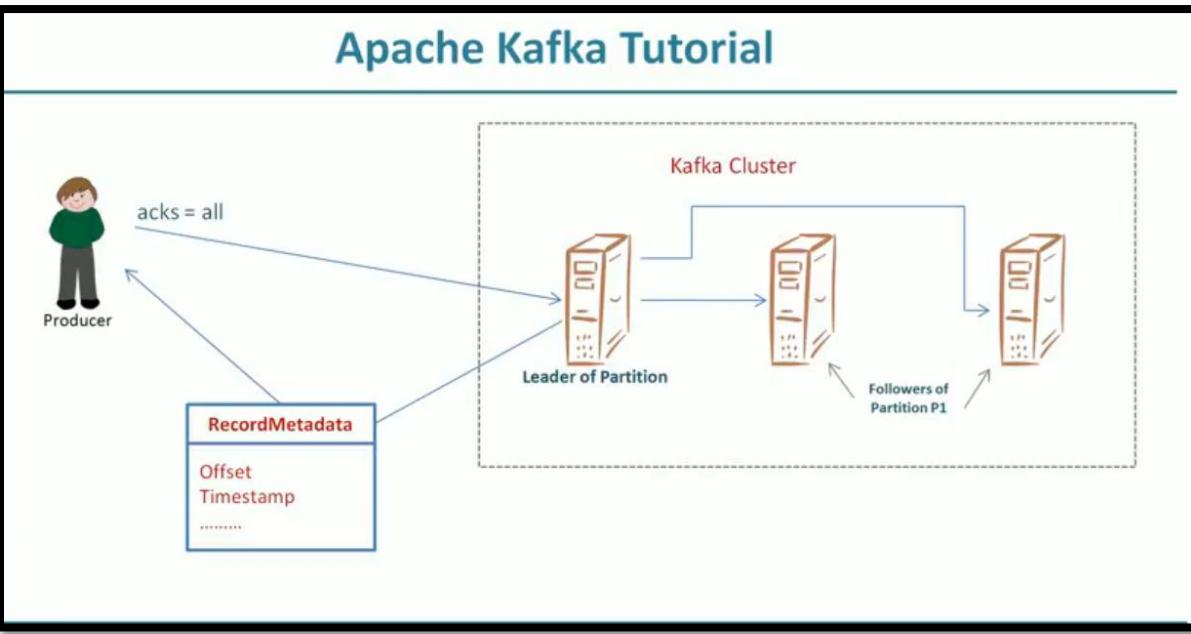
- **acks=0** If set to zero then the producer will not wait for any acknowledgment from the server at all. The record will be immediately added to the socket buffer and considered sent. No guarantee can be made that the server has received the record in this case, and the retries configuration will not take effect (as the client won't generally know of any failures). The offset given back for each record will always be set to -1.
- **acks=1** This will mean the leader will write the record to its local log but will respond without awaiting full acknowledgement from all followers. In this case if the leader fails immediately after acknowledging the record but before the followers have replicated it then the record will be lost.
- **acks=all** This means the leader will wait for the full set of in-sync replicas to acknowledge the record. This guarantees that the record will not be lost as long as at least one in-sync replica remains alive. This is the strongest available guarantee. This is equivalent to the **acks=-1** setting.



Apache Kafka Tutorial



Leader sends acknowledgement before creating replicas, so there is chance to lose the data if leader crashes after acknowledgement



Leader acknowledge only after it receive acknowledge from all of the replicas

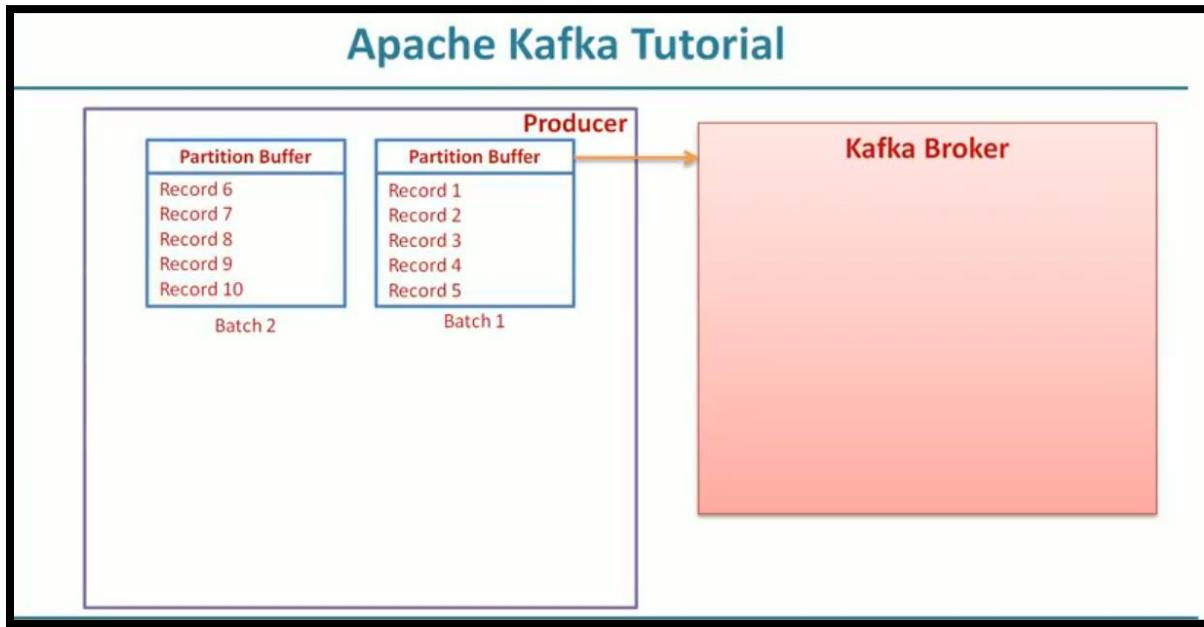
High reliable but high latency

#retries - Setting a value greater than zero will cause the client to resend any record whose send fails with a potentially transient error. Note that this retry is no different than the client resent the record

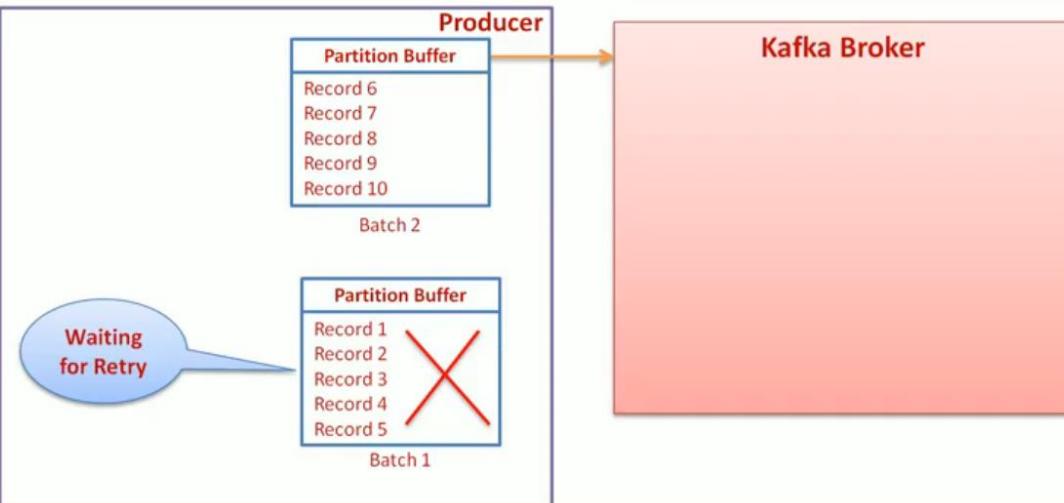
upon receiving the error. Allowing retries without setting **max.in.flight.requests.per.connection** to 1 will potentially change the ordering of records because if two batches are sent to a single partition, and the first fails and is retried but the second succeeds, then the records in the second batch may appear first.

#max.in.flight.requests.per.connection - The maximum number of unacknowledged requests the client will send on a single connection before blocking. Note that if this setting is set to be greater than 1 and there are failed sends, there is a risk of message re-ordering due to retries (i.e., if retries are enabled).

Assume you are delivering two batches to same partition in asynchronous mode

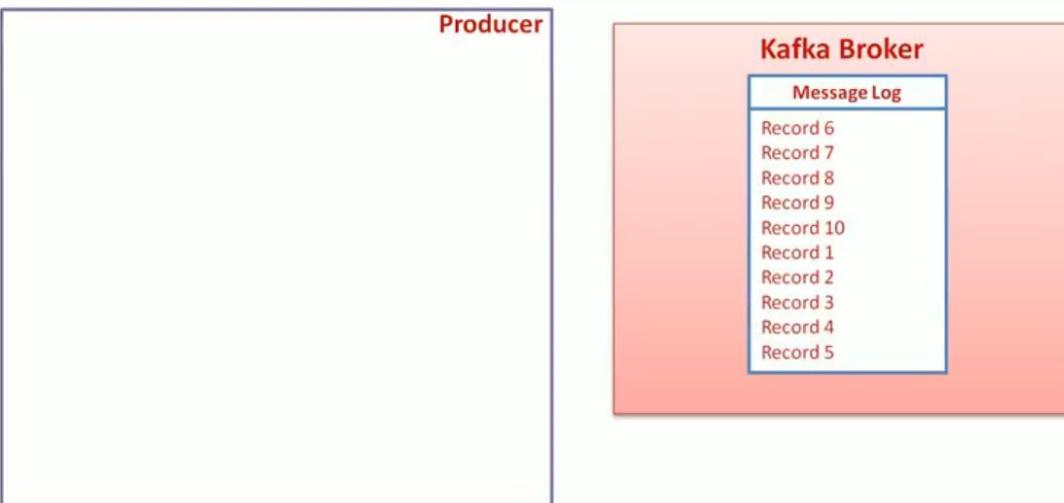


Apache Kafka Tutorial



Let's suppose if batch1 fails and batch2 succeed and after retries the batch1 again and if it is successful you lose your order

Apache Kafka Tutorial



If order is critical you can use following two option

Apache Kafka Tutorial

Kafka Producer Configuration



- • Use Synchronous send
- `max.in.flight.requests.per.connection = 1`

For financial transaction order is critical in that case we can use them

Some more important properties

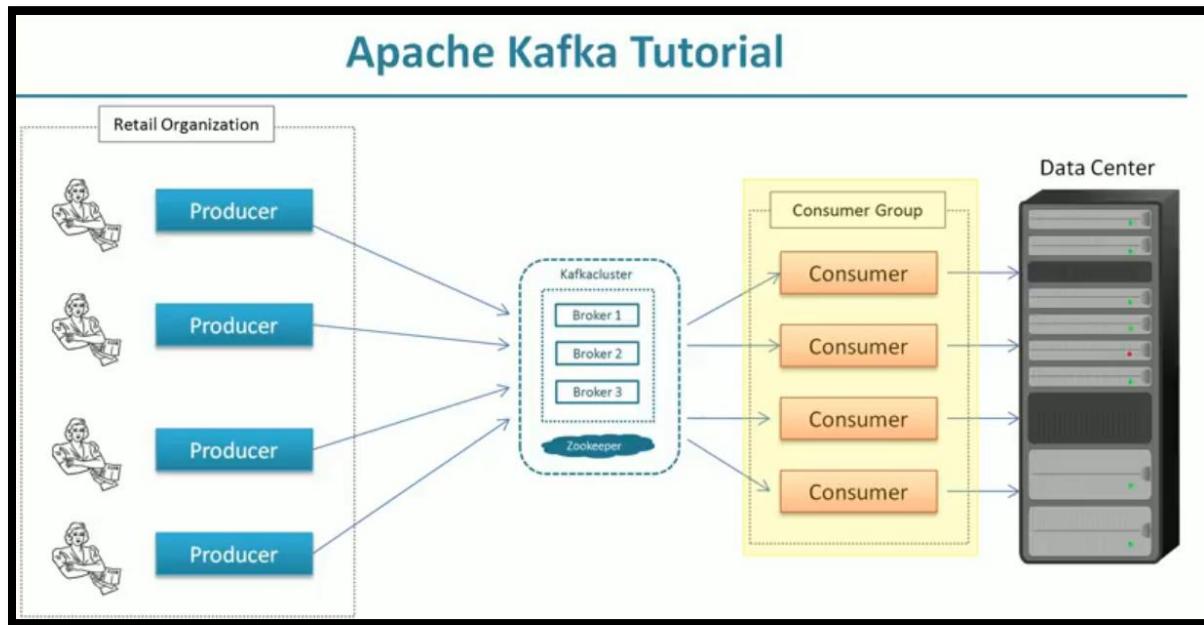
Apache Kafka Tutorial

Kafka Producer Configuration

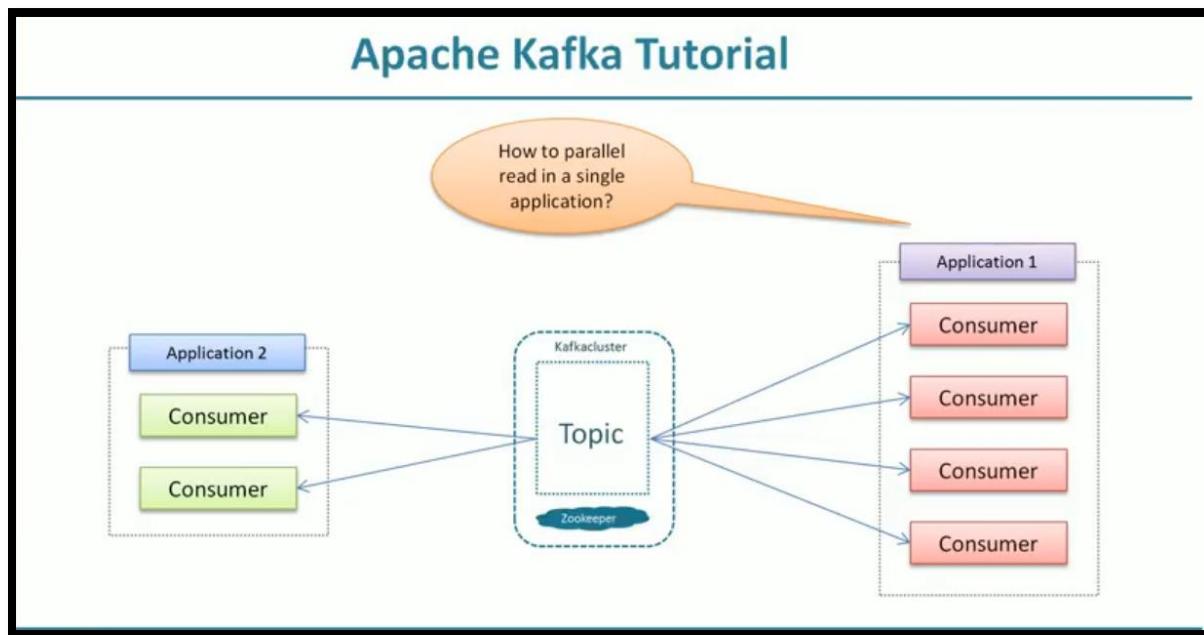


- buffer.memory
- compression.type
- batch.size
- linger.ms
- client.id
- max.request.size

Consumer Group

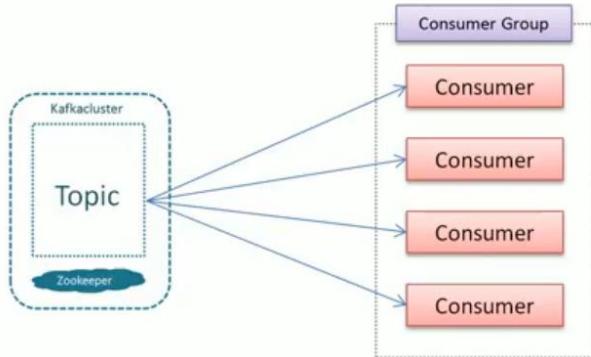


How do we read messages in parallel?



We can read by creating consumer group

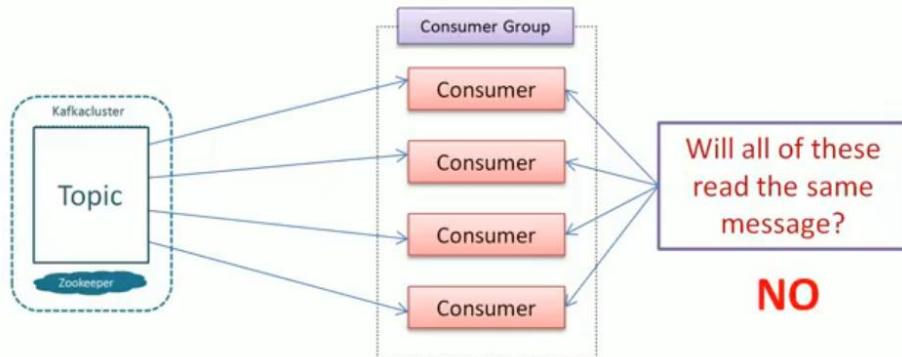
Apache Kafka Tutorial



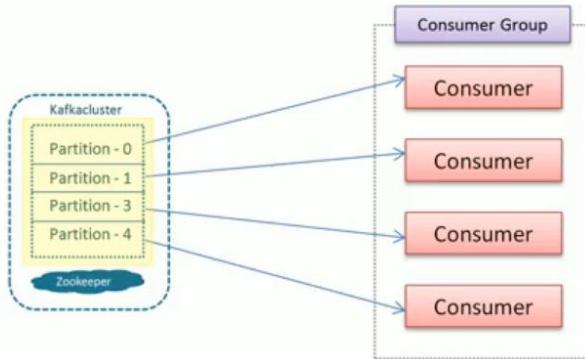
What about duplicate read?

Apache Kafka Tutorial

What about duplicate reads?

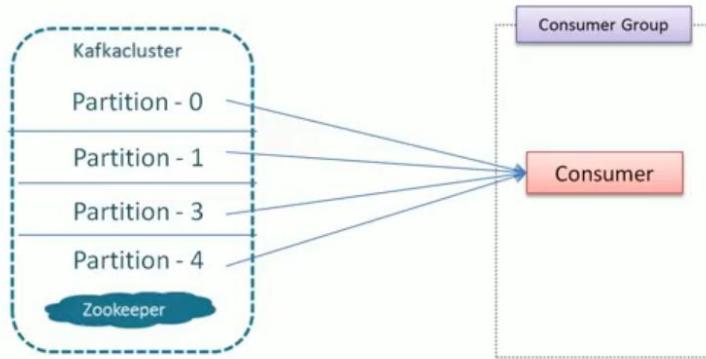


Apache Kafka Tutorial

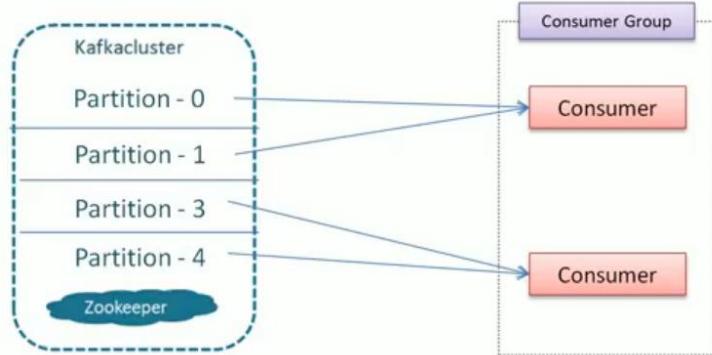


Only one consumer owns a partition at any point in time

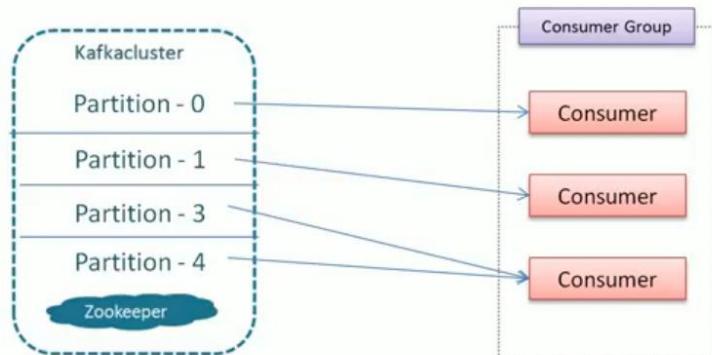
Apache Kafka Tutorial



Apache Kafka Tutorial

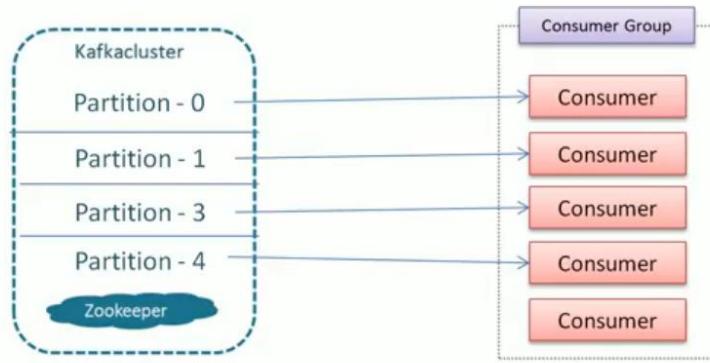


Apache Kafka Tutorial



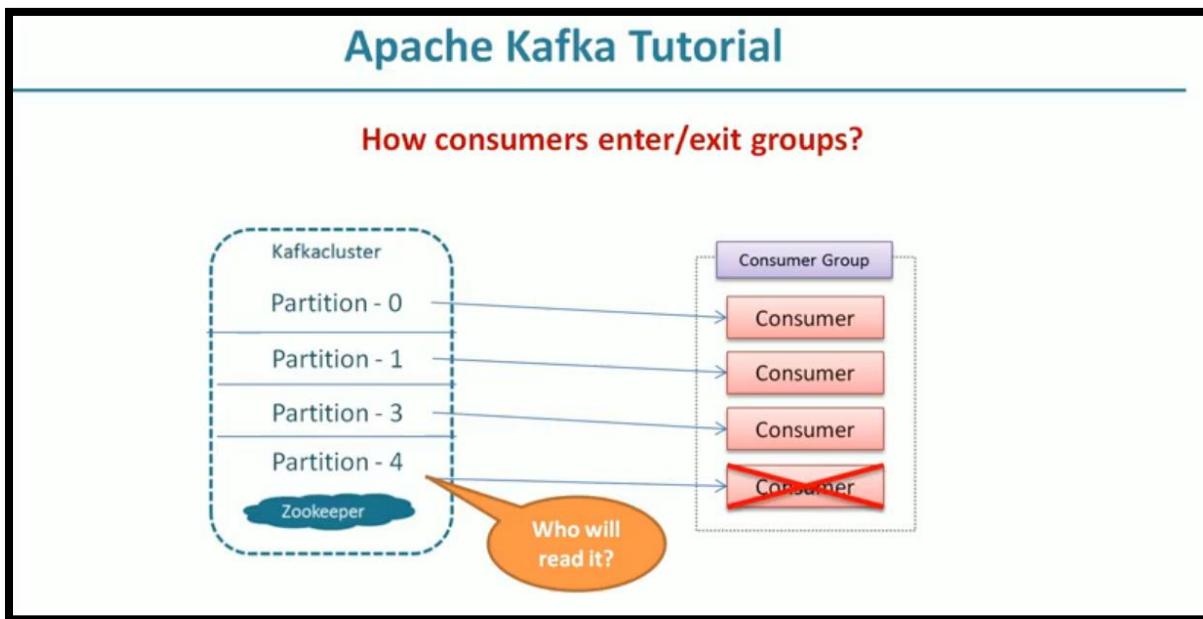
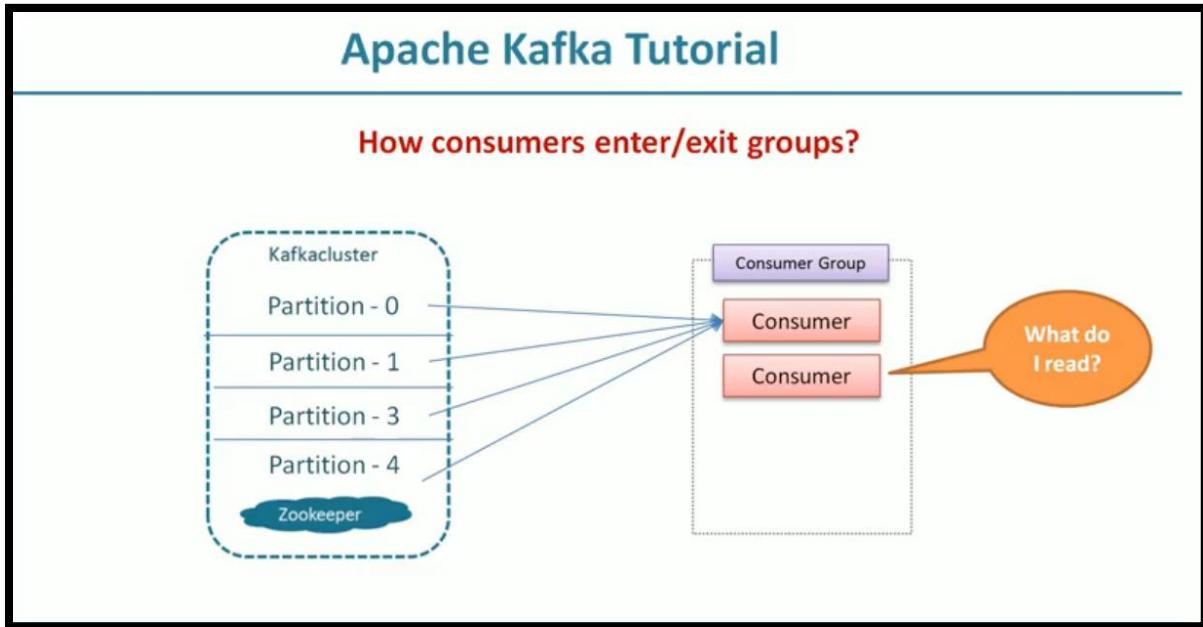
Consumers do not share a partition

Apache Kafka Tutorial



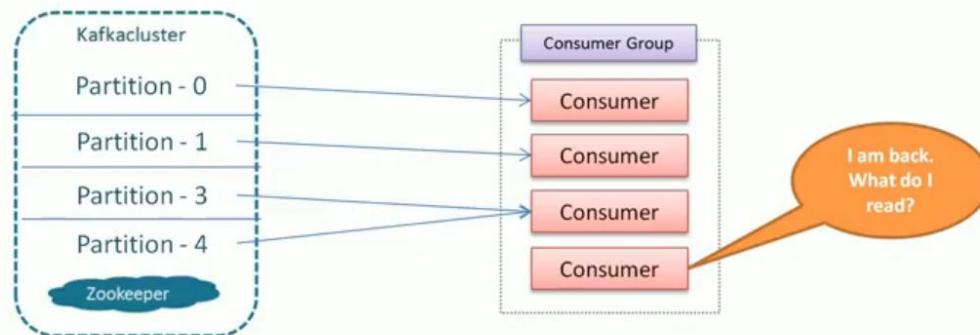
If consumers are greater than partition rest consumer will remain idle

Some problems in assignments/reassignments consumers in a group



Apache Kafka Tutorial

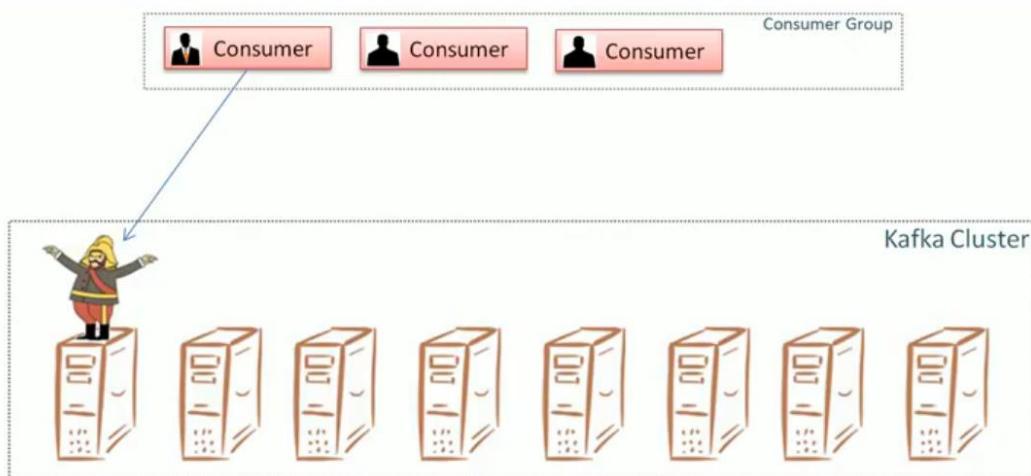
How consumers enter/exit groups?



Group coordinator maintain all of these scenario

One of the kafka broker elected as a group coordinator when consumer wants to join a consumer group it sends a request to group coordinator , first consumer join group becomes leader and all other consumer joining later in the group becomes members of the group

Apache Kafka Tutorial



Apache Kafka Tutorial



Coordinator

- Manage a list of group members
- Initiates a rebalance activity (Block the read for all members)
- Communicate about new assignment to consumers



Leader

- Executes the rebalance activity
- Sends new partition assignment to coordinator

Apache Kafka Tutorial

Summary

- **Consumers Groups** – used to read and process data in parallel.
- **Partitions** – are never shared among members of same group at the same time.
- **Group Coordinator** – maintains a list of active consumers.
- **Rebalance** – initiated when the list of consumers is modified.
- **Group leader** – executes a rebalance activity.

Consumer API

```
import java.util.*;
import org.apache.kafka.clients.consumer.KafkaConsumer;
import org.apache.kafka.clients.consumer.ConsumerRecords;
import org.apache.kafka.clients.consumer.ConsumerRecord;

public class SupplierConsumer{

    public static void main(String[] args) throws Exception{

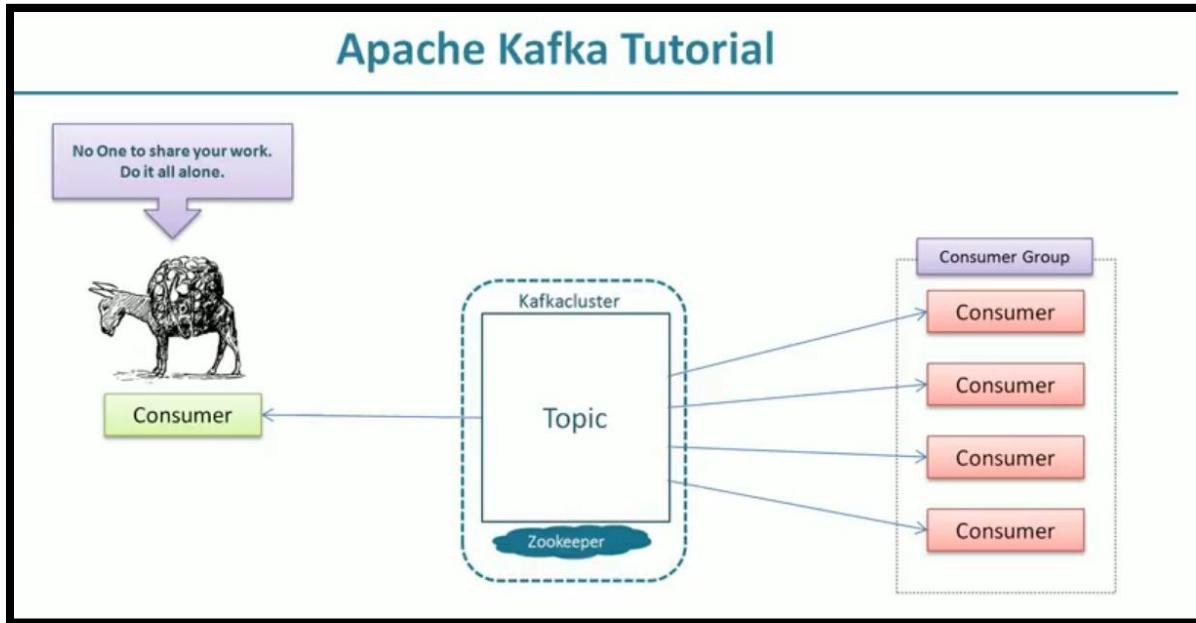
        String topicName = "SupplierTopic";
        String groupName = "SupplierTopicGroup";

        Properties props = new Properties();
        props.put("bootstrap.servers", "localhost:9092,localhost:9093");
// consumer group name as the value of this property , we need not to create consumer group , we just
need to specify the name of the group, if we don't specify this property, this consumer alone will
process all the message for this topic
        props.put("group.id", groupName);
        props.put("key.deserializer", "org.apache.kafka.common.serialization.StringDeserializer");
        props.put("value.deserializer", "SupplierDeserializer");

        KafkaConsumer<String, Supplier> consumer = new KafkaConsumer<>(props);
// consumer can subscribe to multiple topics
        consumer.subscribe(Arrays.asList(topicName));

        while (true){
            // 100 is timeout
            ConsumerRecords<String, Supplier> records = consumer.poll(100);
            for (ConsumerRecord<String, Supplier> record : records){
                System.out.println("Supplier id= " + String.valueOf(record.value().getID()) + " Supplier
Name = " + record.value().getName() + " Supplier Start Date = " +
record.value().getStartDate().toString());
            }
        }
    }
}
```

If consumer is not part of consumer group



Apache Kafka Tutorial

```
32 lines (23 sloc) | 1.37 KB
1 import java.util.*;
2 import org.apache.kafka.clients.consumer.KafkaConsumer;
3 import org.apache.kafka.clients.consumer.ConsumerRecords;
4 import org.apache.kafka.clients.consumer.ConsumerRecord;
5
6 public class SupplierConsumer{
7
8     public static void main(String[] args) throws Exception{
9
10         String topicName = "SupplierTopic";
11         String groupName = "SupplierTopicGroup";
12
13         Properties props = new Properties();
14         props.put("bootstrap.servers", "localhost:9092,localhost:9093");
15         props.put("group.id", groupName);
16         props.put("key.deserializer", "org.apache.kafka.common.serialization.StringDeserializer");
17         props.put("value.deserializer", "SupplierDeserializer");
18
19
20         KafkaConsumer<String, Supplier> consumer = new KafkaConsumer<>(props);
21         consumer.subscribe(Arrays.asList(topicName));
22
23         while (true){
24             ConsumerRecords<String, Supplier> records = consumer.poll(100);
25             for (ConsumerRecord<String, Supplier> record : records){
26                 System.out.println("Supplier id= " + String.valueOf(record.value().getID()) + " Supplier Name = " + record
27             }
28         }
29
30     }
31 }
```

Raw Blame History

- Connect to group Coordinator
- Join the group
- Receives partition assignment
- Sends heartbeat
- Fetches you messages
- And many more things

Apache Kafka Tutorial

```
32 lines (23 sloc) | 1.37 KB
1 import java.util.*;
2 import org.apache.kafka.clients.consumer.KafkaConsumer;
3 import org.apache.kafka.clients.consumer.ConsumerRecords;
4 import org.apache.kafka.clients.consumer.ConsumerRecord;
5
6 public class SupplierConsumer{
7
8     public static void main(String[] args) throws Exception{
9
10         String topicName = "SupplierTopic";
11         String groupName = "SupplierTopicGroup";
12
13         Properties props = new Properties();
14         props.put("bootstrap.servers", "localhost:9092,localhost:9093");
15         props.put("group.id", groupName);
16         props.put("key.deserializer", "org.apache.kafka.common.serialization.StringDeserializer");
17         props.put("value.deserializer", "SupplierDeserializer");
18
19
20         KafkaConsumer<String, Supplier> consumer = new KafkaConsumer<>(props);
21         consumer.subscribe(Arrays.asList(topicName));
22
23         while (true){
24             ConsumerRecords<String, Supplier> records = consumer.poll(100);
25             for (ConsumerRecord<String, Supplier> record : records){
26                 System.out.println("Supplier id= " + String.valueOf(record.value().getID()) + " Supplier Name = " + record
27             }
28         }
29     }
30 }
31 }
```

Make sure each iteration completes in less than 3 seconds.

```
32 lines (23 sloc) | 1.37 KB
1 import java.util.*;
2 import org.apache.kafka.clients.consumer.KafkaConsumer;
3 import org.apache.kafka.clients.consumer.ConsumerRecords;
4 import org.apache.kafka.clients.consumer.ConsumerRecord;
5
6 public class SupplierConsumer{
7
8     public static void main(String[] args) throws Exception{
9
10         String topicName = "SupplierTopic";
11         String groupName = "SupplierTopicGroup";
12
13         Properties props = new Properties();
14         props.put("bootstrap.servers", "localhost:9092,localhost:9093");
15         props.put("group.id", groupName);
16         props.put("key.deserializer", "org.apache.kafka.common.serialization.StringDeserializer");
17         props.put("value.deserializer", "SupplierDeserializer");
18
19
20         KafkaConsumer<String, Supplier> consumer = new KafkaConsumer<>(props);
21         consumer.subscribe(Arrays.asList(topicName));
22
23         while (true){
24             ConsumerRecords<String, Supplier> records = consumer.poll(100);
25             for (ConsumerRecord<String, Supplier> record : records){
26                 System.out.println("Supplier id= " + String.valueOf(record.value().getID()) + " Supplier Name = " + record
27             }
28         }
29     }
30 }
31 }
```

Set `heartbeat.interval.ms` and `session.timeout.ms` to higher if you can't poll every 3 seconds.

We can also use scheduler to start the consumer instead of infinite while loop

Ideally we should not put properties inside code we should maintain a property file and load that file in code

```
import java.util.*;
import java.io.*;
import org.apache.kafka.clients.consumer.KafkaConsumer;
import org.apache.kafka.clients.consumer.ConsumerRecords;
import org.apache.kafka.clients.consumer.ConsumerRecord;
```

```
public class NewSupplierConsumer{

    public static void main(String[] args) throws Exception{

        String topicName = "SupplierTopic";
        String groupName = "SupplierTopicGroup";
        Properties props = new Properties();
        //props.put("bootstrap.servers", "localhost:9092,localhost:9093");
        //props.put("group.id", groupName);
        //props.put("key.deserializer", "org.apache.kafka.common.serialization.StringDeserializer");
        //props.put("value.deserializer", "SupplierDeserializer");

        InputStream input = null;
        KafkaConsumer<String, Supplier> consumer = null;

        try {
            input = new FileInputStream("SupplierConsumer.properties");
            props.load(input);
            consumer = new KafkaConsumer<>(props);
            consumer.subscribe(Arrays.asList(topicName));

            while (true){
                ConsumerRecords<String, Supplier> records = consumer.poll(100);
                for (ConsumerRecord<String, Supplier> record : records){
                    System.out.println("Supplier id= " + String.valueOf(record.value().getID()) + " Supplier Name
= " + record.value().getName() + " Supplier Start Date = " + record.value().getStartDate().toString());
                }
            }
        }catch(Exception ex){
            ex.printStackTrace();
        }finally{
            input.close();
            consumer.close();
        }
    }
}
```

Offset Management

Apache Kafka Tutorial

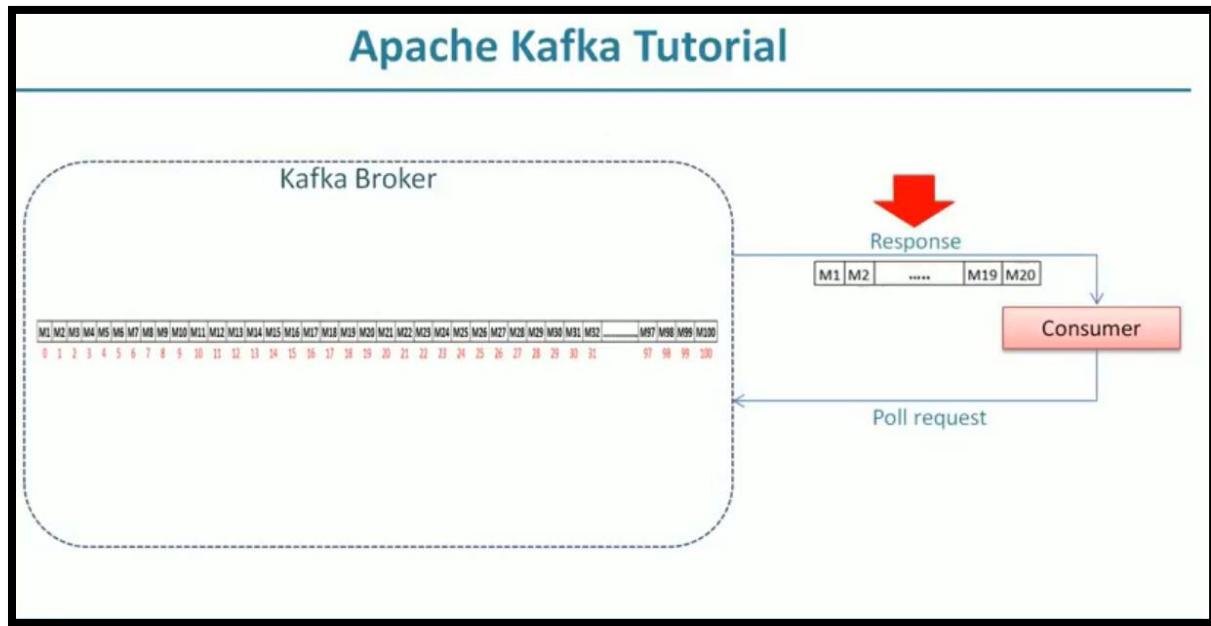
Kafka Offsets

- Current Offset
- Committed Offset

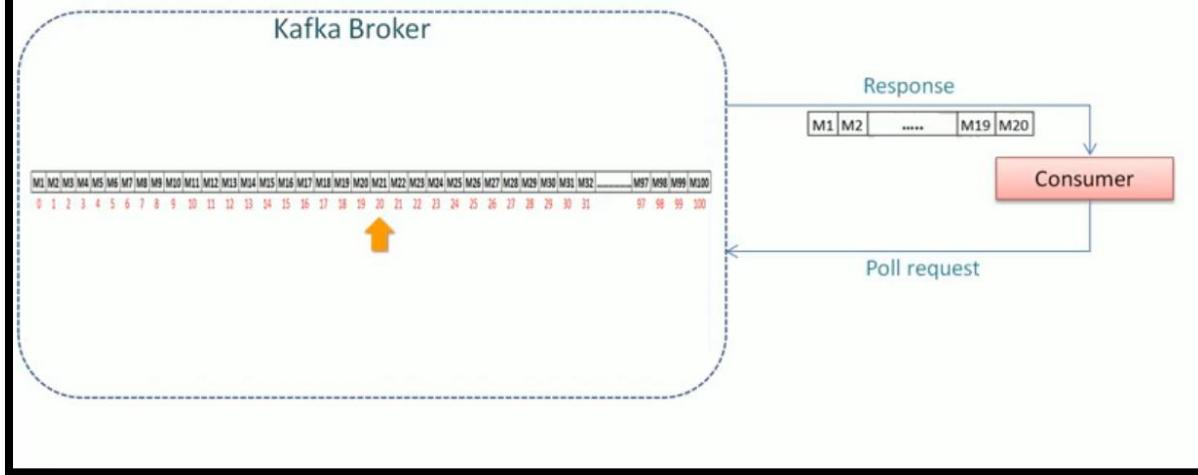
Current offset - is used to avoid resending same record again to the same consumer

Committed Offset – is used to avoid resending the same record to new consumer in the event of partition rebalance.

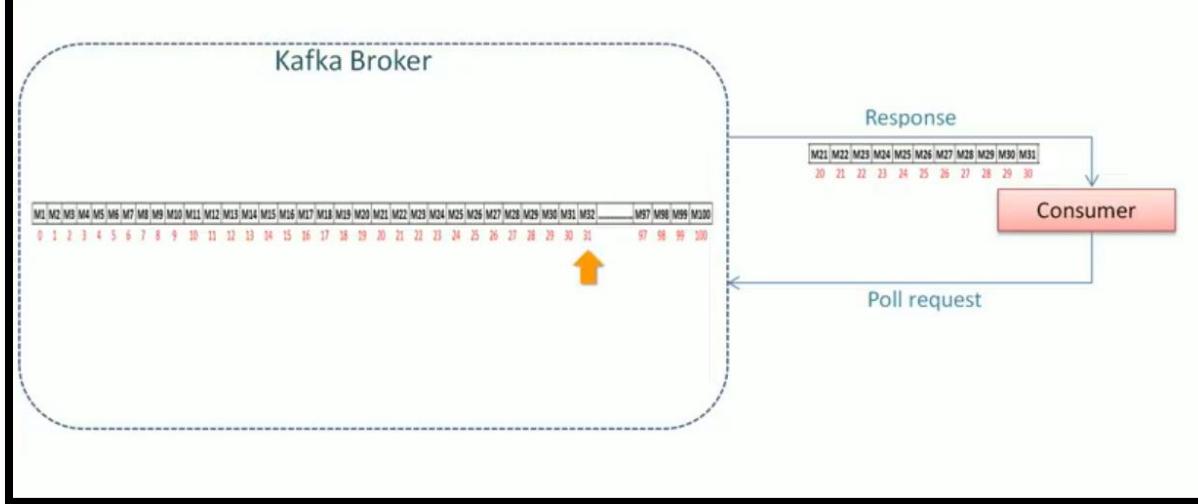
Kafka broker manages current and committed offset for each consumer



Apache Kafka Tutorial



Apache Kafka Tutorial

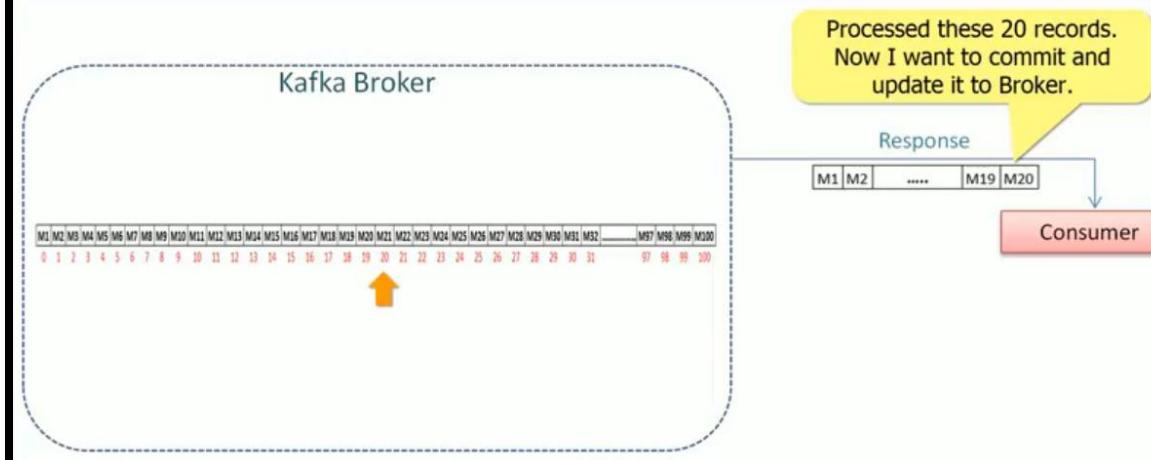


Apache Kafka Tutorial

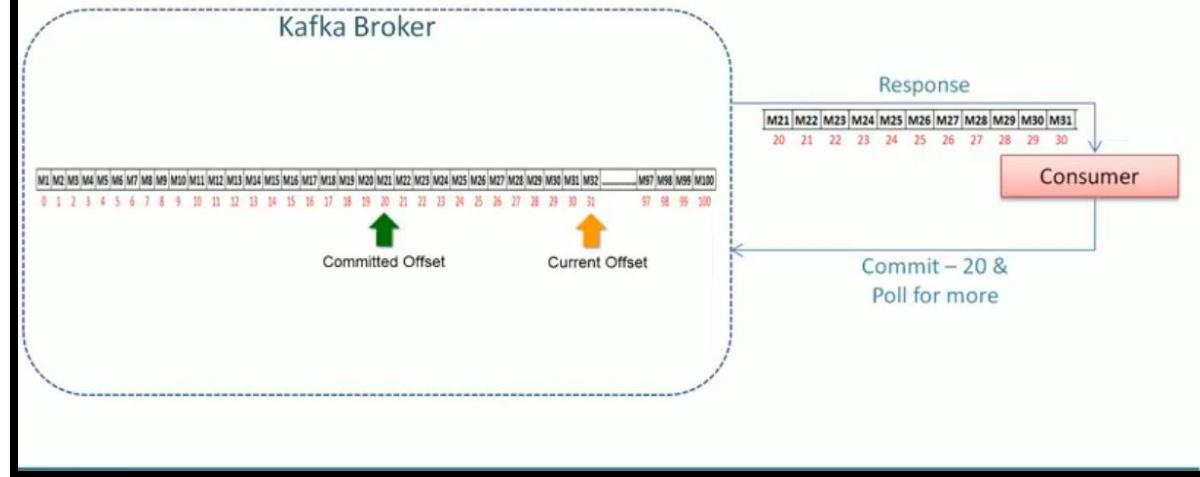
Committed Offset

The offset position that is already processed by a consumer.

Apache Kafka Tutorial



Apache Kafka Tutorial



Committed offset is the last record that consumer has successfully processed

In the event of rebalancing when a new consumer assign the same partition, it should now where to start reading means what records are already processed by the previous owner

Apache Kafka Tutorial

Offsets

- **Current offset** – Sent Records
- **Committed offset** – Processed Records

How to commit an offset?

- Auto Commit
- Manual Commit

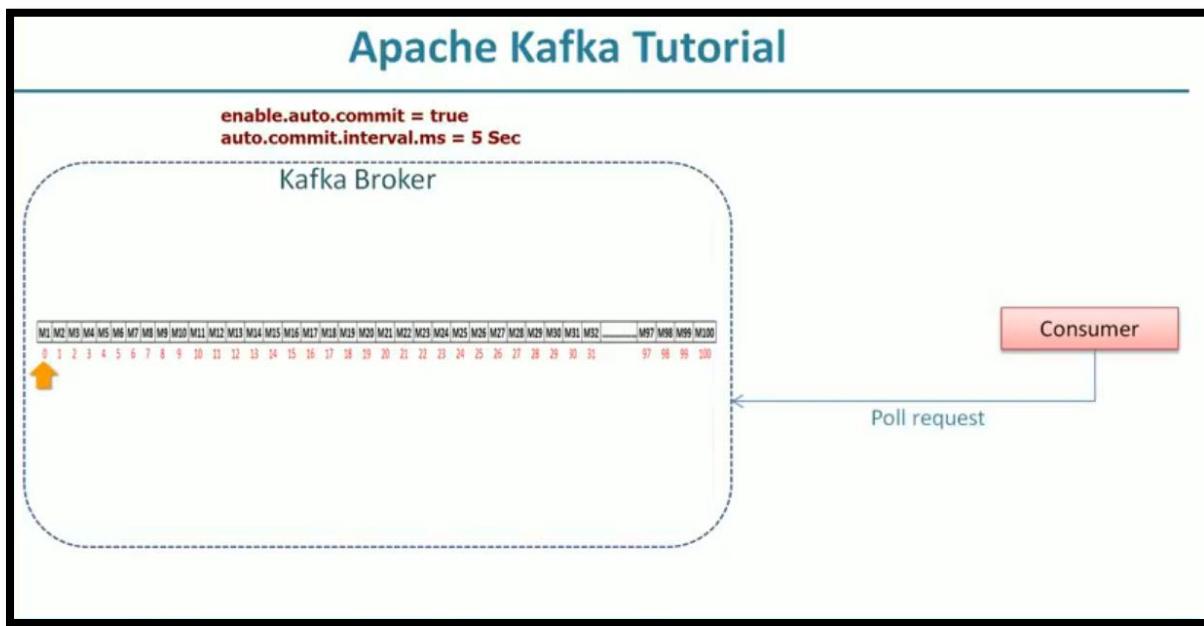
Apache Kafka Tutorial

Auto Commit

- enable.auto.commit
- auto.commit.interval.ms

enable.auto.commit is by default set to true

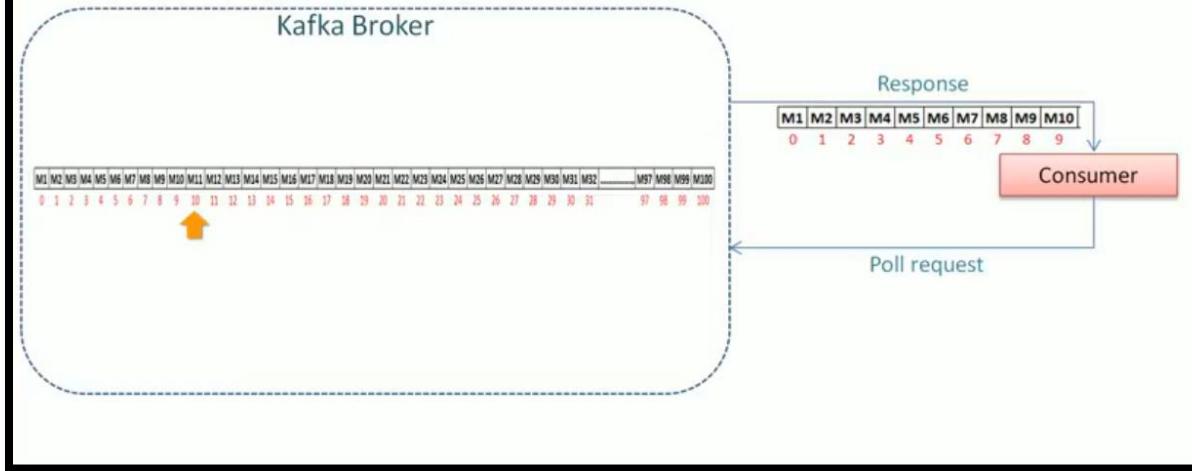
auto.commit.interval.ms defines interval of auto commit default value is 5 sec



You receive 10 messages and consumer increases the offset to 10

Apache Kafka Tutorial

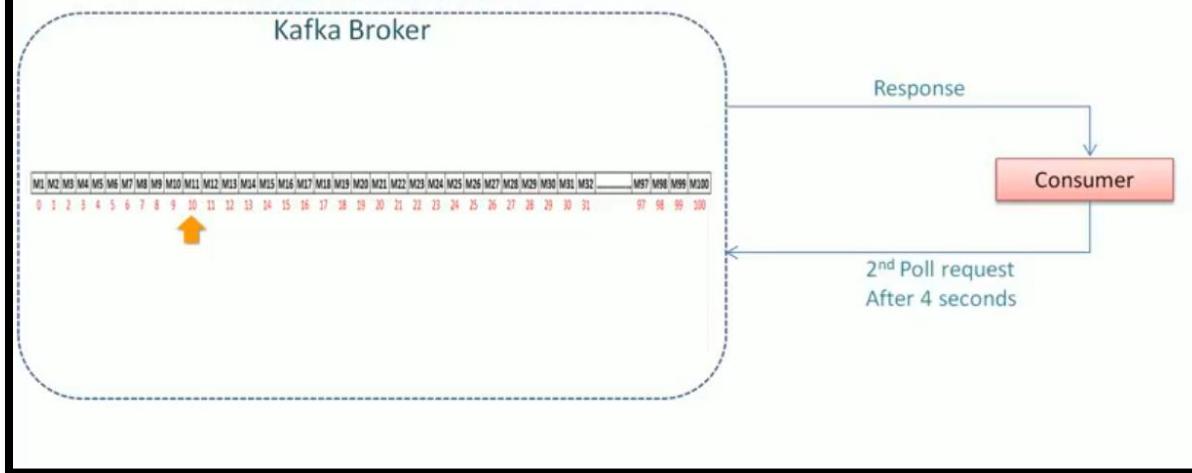
enable.auto.commit = true
auto.commit.interval.ms = 5 Sec



You take 4 seconds to process the message and make a new call

Apache Kafka Tutorial

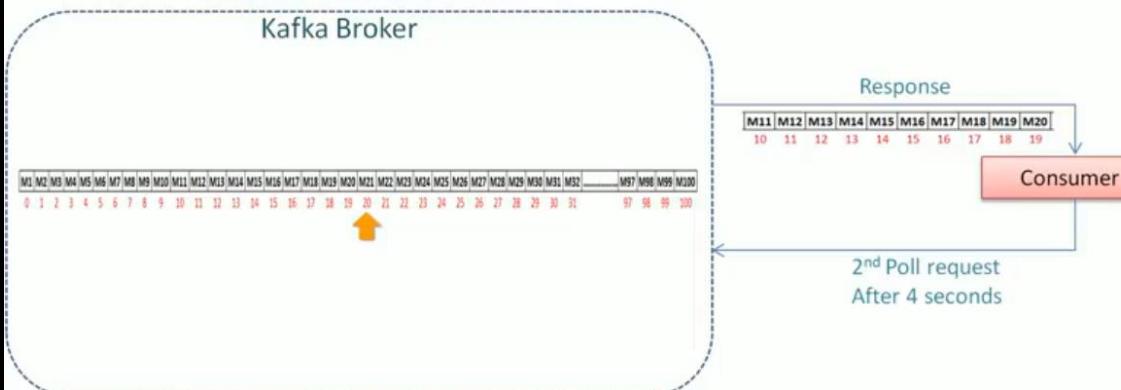
enable.auto.commit = true
auto.commit.interval.ms = 5 Sec



Since you haven't passed five seconds consumer will not commit the offset, you receive another set of records

Apache Kafka Tutorial

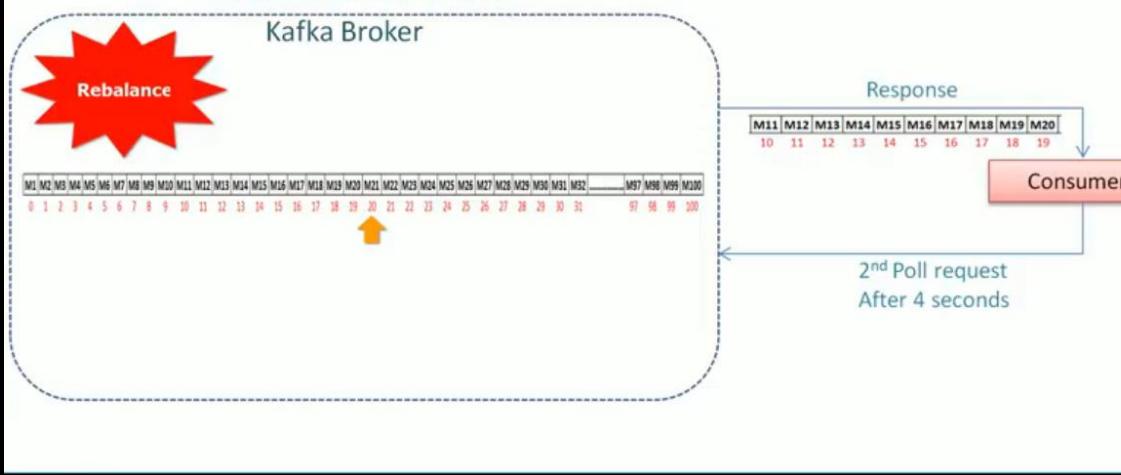
enable.auto.commit = true
auto.commit.interval.ms = 5 Sec



For some reason rebalance occurs at this movement

Apache Kafka Tutorial

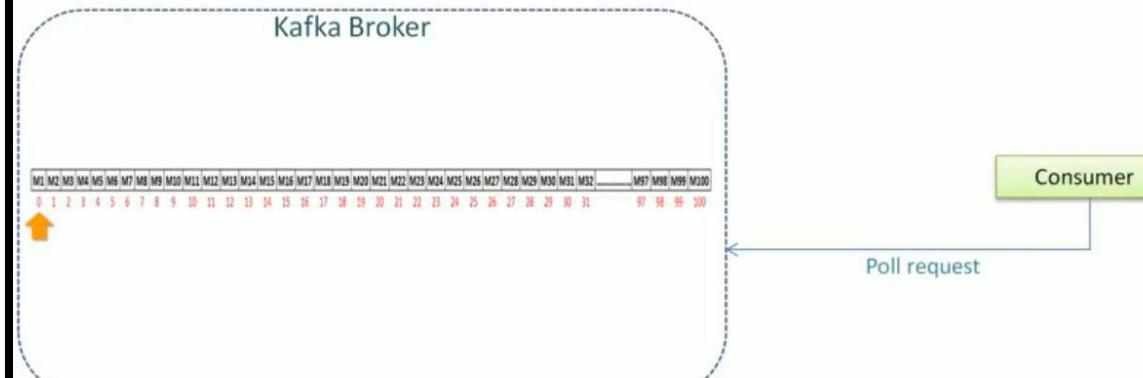
enable.auto.commit = true
auto.commit.interval.ms = 5 Sec



In this scenario first 10 records are already processed but nothing is committed yet

Apache Kafka Tutorial

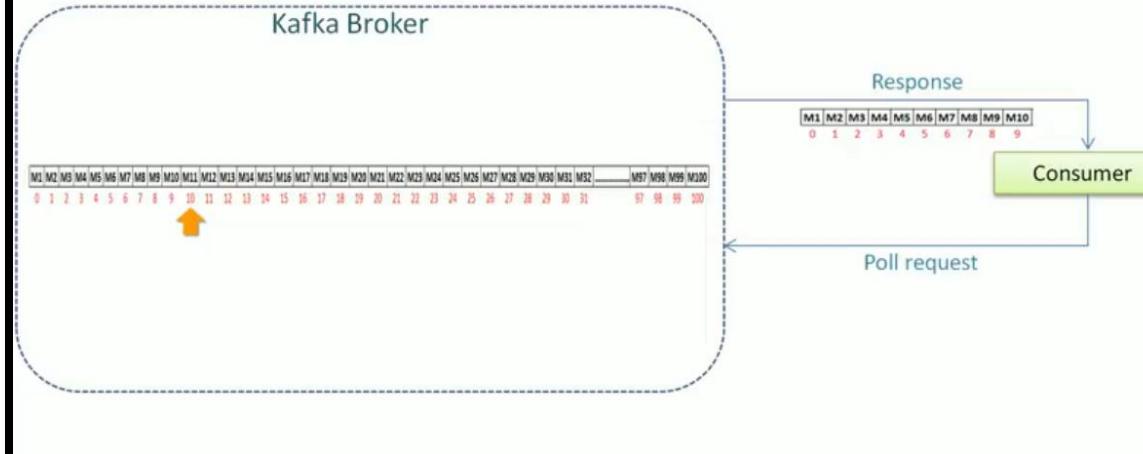
enable.auto.commit = true
auto.commit.interval.ms = 5 Sec



Now partition goes to a different consumer because of rebalance event since we don't have the committed offset the new owner of this partition start processing from the beginning and process first 10 records once again

Apache Kafka Tutorial

enable.auto.commit = true
auto.commit.interval.ms = 5 Sec



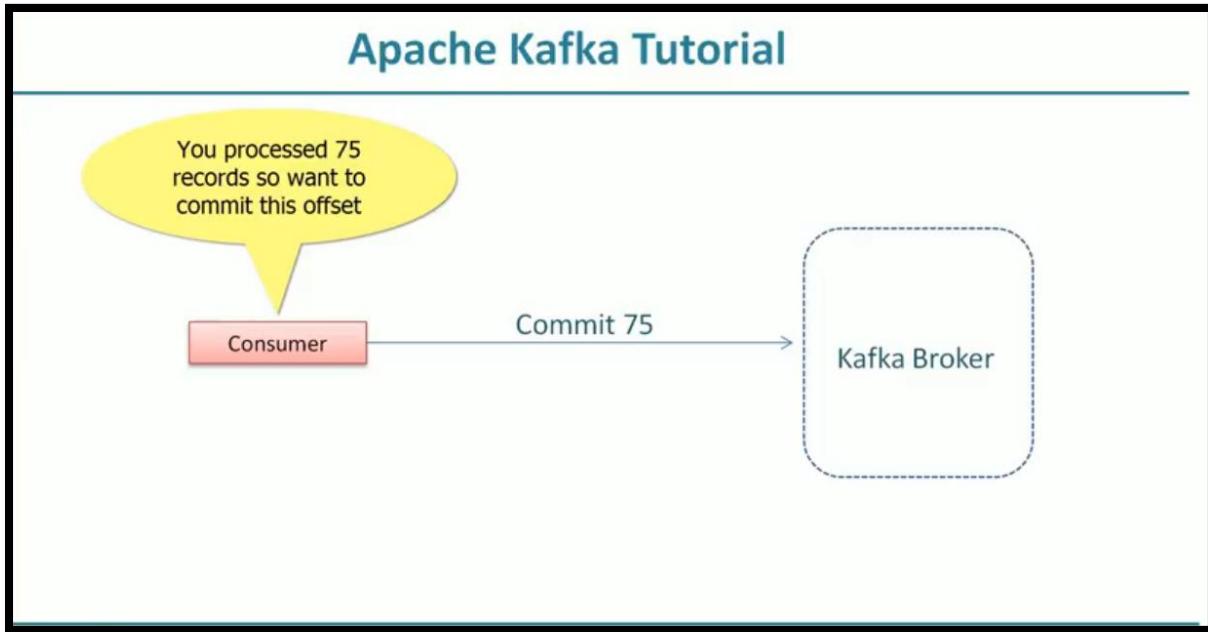
The solution to this particular problem is the manual commit set auto commit to false and manual commit after processing the records

Manual commit has two approaches

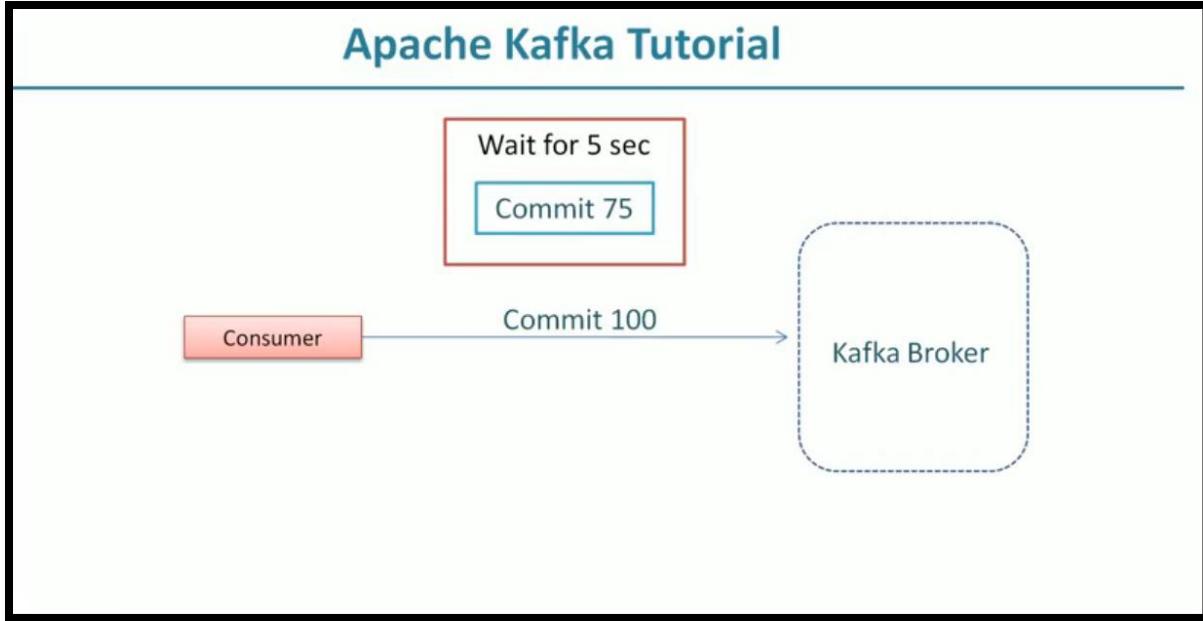
- Commit Sync – blocking method it will block your call for completing a commit operation it will also retries if there are any recoverable errors like rebalancing
- Commit Async – unblocking method but commit Async will not retries

Why Aync Commit won't retry?

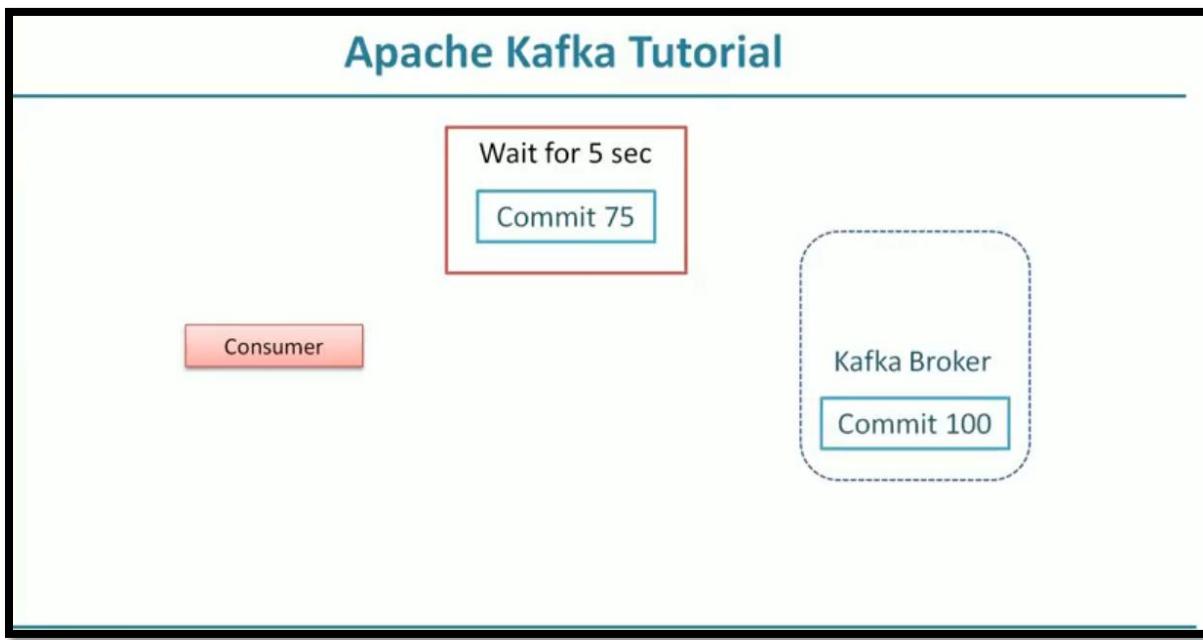
Assume you are trying to commit offset as 75, it failed for some recoverable reason and you want to retry it after few seconds



Since this was an asynchronous call so without knowing that your previous commit is waiting you initiated another commit, this time it is to commit offset 100



Your commit 100 is successful while commit 75 is waiting for retry



In this case you don't want to commit 75 after commit 100 that may cause issues so they designed asynchronous commit to not to retry

Code

```
import java.util.*;
import java.io.*;
import org.apache.kafka.clients.consumer.KafkaConsumer;
import org.apache.kafka.clients.consumer.ConsumerRecords;
import org.apache.kafka.clients.consumer.ConsumerRecord;

public class ManualConsumer{

    public static void main(String[] args) throws Exception{

        String topicName = "SupplierTopic";
        String groupName = "SupplierTopicGroup";

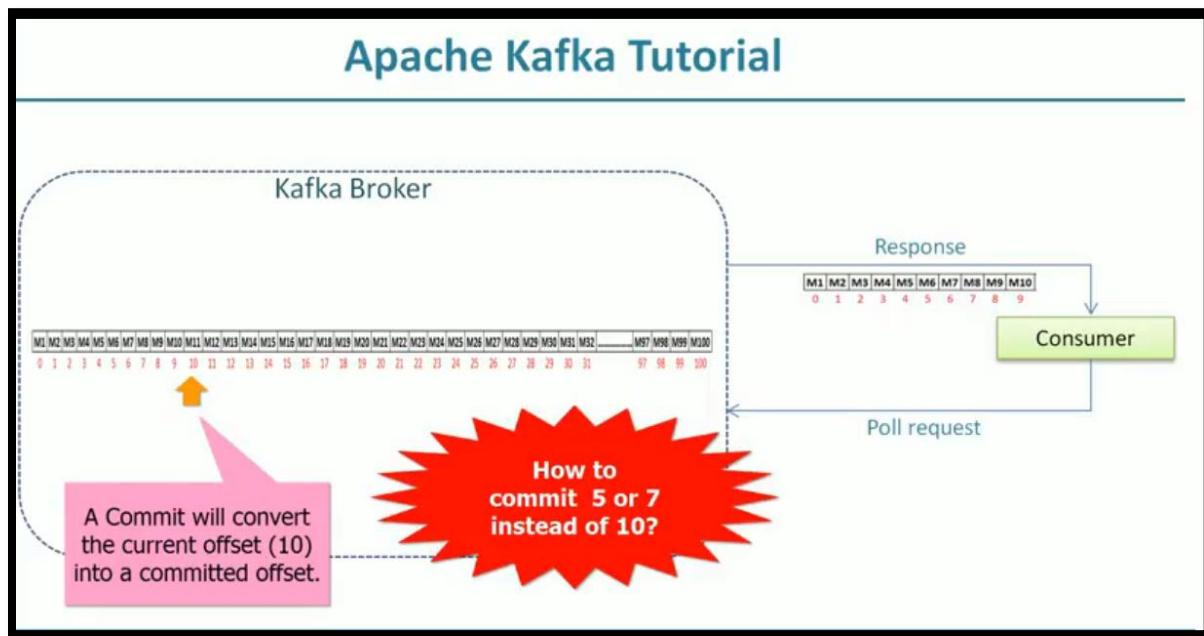
        Properties props = new Properties();
        props.put("bootstrap.servers", "localhost:9092,localhost:9093");
        props.put("group.id", groupName);
        props.put("key.deserializer", "org.apache.kafka.common.serialization.StringDeserializer");
        props.put("value.deserializer", "SupplierDeserializer");
        // for manual commit we need to set auto commit to false
        props.put("enable.auto.commit", "false");

        KafkaConsumer<String, Supplier> consumer = null;

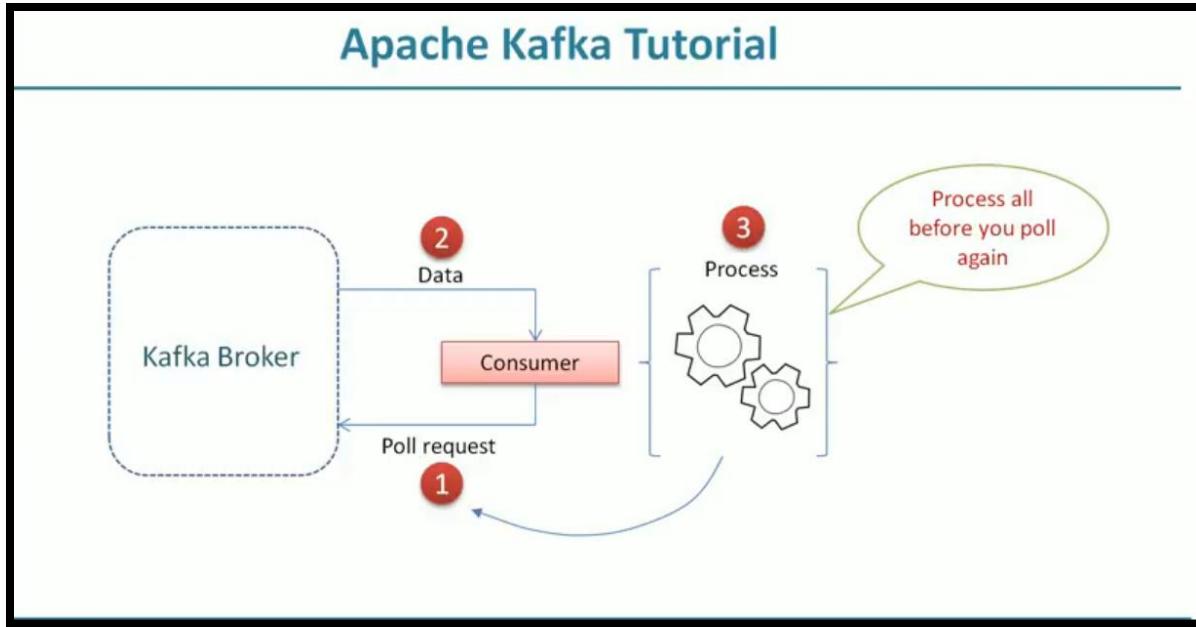
        try {
            consumer = new KafkaConsumer<>(props);
            consumer.subscribe(Arrays.asList(topicName));

            while (true){
                ConsumerRecords<String, Supplier> records = consumer.poll(100);
                for (ConsumerRecord<String, Supplier> record : records){
                    System.out.println("Supplier id= " + String.valueOf(record.value().getID()) + " Supplier Name = " + record.value().getName() + " Supplier Start Date = " + record.value().getStartDate().toString());
                }
                consumer.commitAsync();
            }
        }catch(Exception ex){
            ex.printStackTrace();
        }finally{
            consumer.commitSync();
            consumer.close();
        }
    }
}
```

Rebalance Listener



Suppose we are getting lots of data using the poll method and it is going to take some reasonable time to complete the processing of all of the records



If you are taking lots of time to process the records you have two types of risks

- Delay in next poll – if you don't poll for a long, group coordinator might think that you are dead and trigger group rebalancing activity
- Rebalance is triggered - Rebalance activity for some other reason while you were processing the extensive records

In both the cases your current partition will be taken away from you and assign to some other consumer in such case you will want to commit whatever you had already processed before the ownership of partition is taken from you and the new owner of the partition is supposed to start consuming it from the last committed offset

Apache Kafka Tutorial

I am starting a Re-balance.



Coordinator

Hold On. Let me commit an offset.



Consumer

Apache Kafka Tutorial

- ➡ 1. How to commit a particular offset?
- 2. How to know that a rebalance is triggered?

Apache Kafka Tutorial

Example code

1. Maintain an offset of processed record
- 2. Commit when rebalance is triggered

Apache Kafka Tutorial

ConsumerRebalanceListener

1. onPartitionsRevoked
2. onPartitionsAssigned

We can commit on onPartitionsRevoked

```
import java.util.*;  
import org.apache.kafka.clients.producer.*;  
public class RandomProducer {  
  
    public static void main(String[] args) throws InterruptedException{  
  
        String topicName = "RandomProducerTopic";
```

```

String msg;

Properties props = new Properties();
props.put("bootstrap.servers", "localhost:9092,localhost:9093");
props.put("key.serializer", "org.apache.kafka.common.serialization.StringSerializer");
props.put("value.serializer", "org.apache.kafka.common.serialization.StringSerializer");

Producer<String, String> producer = new KafkaProducer <>(props);
Random rg = new Random();
Calendar dt = Calendar.getInstance();
dt.set(2016,1,1);
try{
    while(true){
        for (int i=0;i<100;i++){
            msg = dt.get(Calendar.YEAR)+"-"+dt.get(Calendar.MONTH)+"-"+dt.get(Calendar.DATE) + "," +
rg.nextInt(1000);
            producer.send(new ProducerRecord<String, String>(topicName,0,"Key",msg).get();
            msg = dt.get(Calendar.YEAR)+"-"+dt.get(Calendar.MONTH)+"-"+dt.get(Calendar.DATE) + "," +
rg.nextInt(1000);
            producer.send(new ProducerRecord<String, String>(topicName,1,"Key",msg).get();
        }
        dt.add(Calendar.DATE,1);
        System.out.println("Data Sent for " + dt.get(Calendar.YEAR) + "-" + dt.get(Calendar.MONTH) + "-"+
" + dt.get(Calendar.DATE) );
    }
}
catch(Exception ex){
    System.out.println("Intrupted");
}
finally{
    producer.close();
}

}
}

```

Here we are sending the same data to two partitions

```

import java.util.*;
import org.apache.kafka.clients.consumer.*;
import org.apache.kafka.common.*;

public class RandomConsumer{

    public static void main(String[] args) throws Exception{

```

```
String topicName = "RandomProducerTopic";
KafkaConsumer<String, String> consumer = null;

String groupName = "RG";
Properties props = new Properties();
props.put("bootstrap.servers", "localhost:9092,localhost:9093");
props.put("group.id", groupName);
props.put("key.deserializer", "org.apache.kafka.common.serialization.StringDeserializer");
props.put("value.deserializer", "org.apache.kafka.common.serialization.StringDeserializer");
props.put("enable.auto.commit", "false");

consumer = new KafkaConsumer<>(props);
RebalanceListner rebalanceListner = new RebalanceListner(consumer);

consumer.subscribe(Arrays.asList(topicName),rebalanceListner);
try{
    while (true){
        ConsumerRecords<String, String> records = consumer.poll(100);
        for (ConsumerRecord<String, String> record : records){
            //System.out.println("Topic:"+ record.topic() + " Partition:" + record.partition() + " Offset:"
+ record.offset() + " Value:" + record.value());
            // Do some processing and save it to Database
            rebalanceListner.addOffset(record.topic(), record.partition(),record.offset());
        }
        //consumer.commitSync(rebalanceListner.getCurrentOffsets());
    }
}catch(Exception ex){
    System.out.println("Exception.");
    ex.printStackTrace();
}
finally{
    consumer.close();
}
}

}
```

Apache Kafka Tutorial

ConsumerRebalanceListener

1. Maintain a list of offsets that are processed and ready to be committed.
2. Commit the offsets when Partitions are going away.

```
import java.util.*;
import org.apache.kafka.clients.consumer.*;
import org.apache.kafka.common.*;

public class RebalanceListner implements ConsumerRebalanceListener {
    private KafkaConsumer consumer;
    private Map<TopicPartition, OffsetAndMetadata> currentOffsets = new HashMap();

    public RebalanceListner(KafkaConsumer con){
        this.consumer=con;
    }

    public void addOffset(String topic, int partition, long offset){
        currentOffsets.put(new TopicPartition(topic, partition),new OffsetAndMetadata(offset,"Commit"));
    }

    public Map<TopicPartition, OffsetAndMetadata> getCurrentOffsets(){
        return currentOffsets;
    }

    public void onPartitionsAssigned(Collection<TopicPartition> partitions) {
        System.out.println("Following Partitions Assigned ....");
        for(TopicPartition partition: partitions)
            System.out.println(partition.partition()+"","");
    }

    public void onPartitionsRevoked(Collection<TopicPartition> partitions) {
        System.out.println("Following Partitions Revoked ....");
    }
}
```

```

for(TopicPartition partition: partitions)
    System.out.println(partition.partition() + ",");

System.out.println("Following Partitions committed ....");
for(TopicPartition tp: currentOffsets.keySet())
    System.out.println(tp.partition());

    consumer.commitSync(currentOffsets);
    currentOffsets.clear();
}
}

```

The image shows three terminal windows side-by-side, each displaying Kafka log output. The leftmost window shows a continuous stream of log entries starting with 'Data Sent for 2016-2-29'. The middle window shows log entries for 'Following Partitions Revoked', 'Following Partitions committed', and 'Following Partitions Assigned'. The rightmost window shows log entries for 'SupplierConsumer', 'SupplierProducer', and 'SynchronousProducer' with an input prompt 'Enter number: 4'. A yellow oval containing the text 'What if you kill one of these?' has two arrows pointing from it to the top two windows.

```

JavaProj : bash
File Edit View Scrollback Bookmarks Settings Help
Data Sent for 2016-2-29
Data Sent for 2016-2-30
Data Sent for 2016-2-31
Data Sent for 2016-3-1
Data Sent for 2016-3-2
Data Sent for 2016-3-3
Data Sent for 2016-3-4
Data Sent for 2016-3-5
Data Sent for 2016-3-6
Data Sent for 2016-3-7
Data Sent for 2016-3-8
Data Sent for 2016-3-9
Data Sent for 2016-3-10
Data Sent for 2016-3-11
Data Sent for 2016-3-12
Data Sent for 2016-3-13
Data Sent for 2016-3-14
Data Sent for 2016-3-15
Data Sent for 2016-3-16
Data Sent for 2016-3-17
Data Sent for 2016-3-18
Data Sent for 2016-3-19
Data Sent for 2016-3-20
Data Sent for 2016-3-21
Data Sent for 2016-3-22
Data Sent for 2016-3-23
Data Sent for 2016-3-24
Data Sent for 2016-3-25
Data Sent for 2016-3-26
Data Sent for 2016-3-27
>Data Sent for 2016-3-28
[root@localhost javaProj]# 

javaProj : sbt
File Edit View Scrollback Bookmarks Settings Help
Following Partitions Revoked ....
Following Partitions committed ....
Following Partitions Assigned ....
0,
1,
Following Partitions Revoked ....
0,
1,
Following Partitions committed ....
0,
1,
Following Partitions Assigned ....
0,
[]

javaProj : sbt
File Edit View Scrollback Bookmarks Settings Help
[8] SupplierConsumer
[9] SupplierProducer
[10] SynchronousProducer
>Enter number: 4

[info] Running RandomConsumer
SLF4J: Failed to load class "org.slf4j.impl.StaticLoggerBinder".
SLF4J: Defaulting to no-operation (NOP) logger implementation
SLF4J: See http://www.slf4j.org/codes.html#staticLoggerBinder for further details.
Following Partitions Revoked ....
Following Partitions committed ....
Following Partitions Assigned ....
1,

```

Exactly one processing

Apache Kafka Tutorial

Why a consumer group

1. Allows you to parallel process a topic.
2. Automatically manages partition assignment.
3. Detects entry/exit/failure of a consumer and perform partition rebalancing activity.

Apache Kafka Tutorial

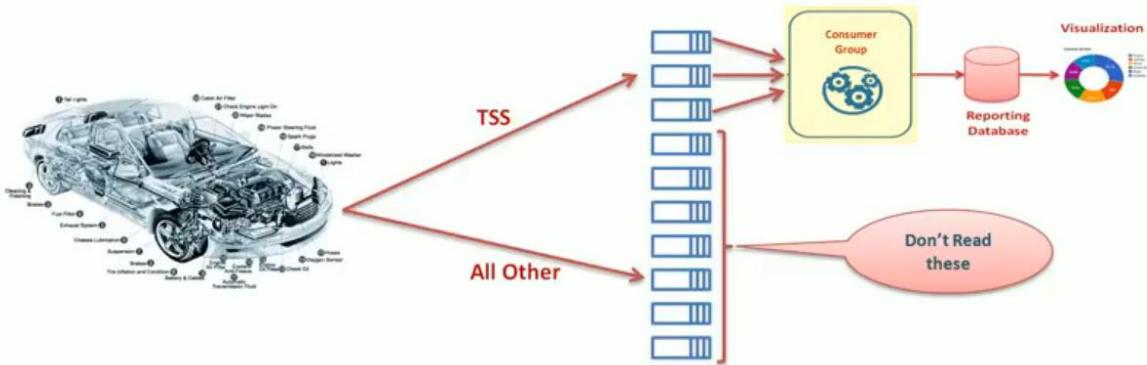
Automatic Partition Assignment

You don't know which partition goes to which consumer.

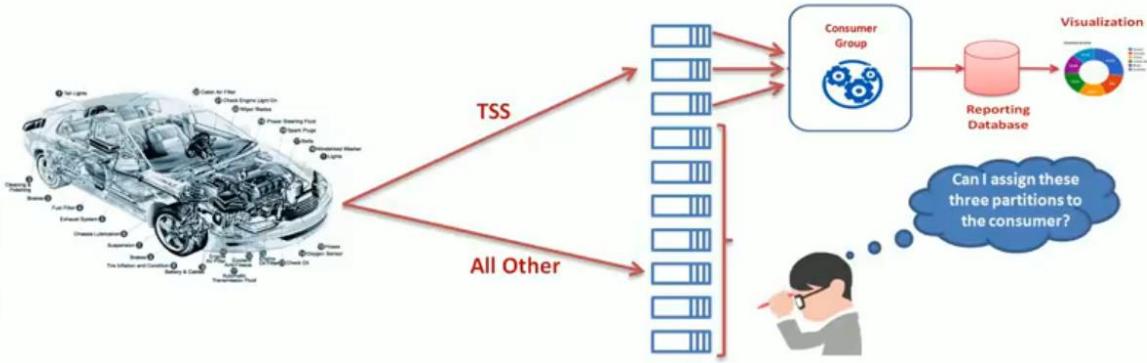
There are two strategies

1. Range
2. Round robin

Apache Kafka Tutorial



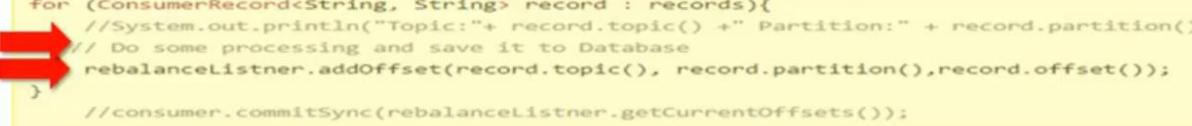
Apache Kafka Tutorial



```

    consumer.poll(100);
    for (ConsumerRecord<String, String> record : records){
        //System.out.println("Topic:" + record.topic() + " Partition:" + record.partition()
        // Do some processing and save it to Database
        rebalanceListner.addOffset(record.topic(), record.partition(), record.offset());
    }
    //consumer.commitSync(rebalanceListner.getCurrentOffsets());
}
catch(Exception ex){
    System.out.println("Exception.");
    ex.printStackTrace();
}
finally{
    consumer.close();
}

```



Problem – since we are saving the record to database and then adding offset to listener, these two steps are not atomic, it is possible after saving it to database consumer crashes before committing the offset, so we need to create these two steps as single atomic transaction

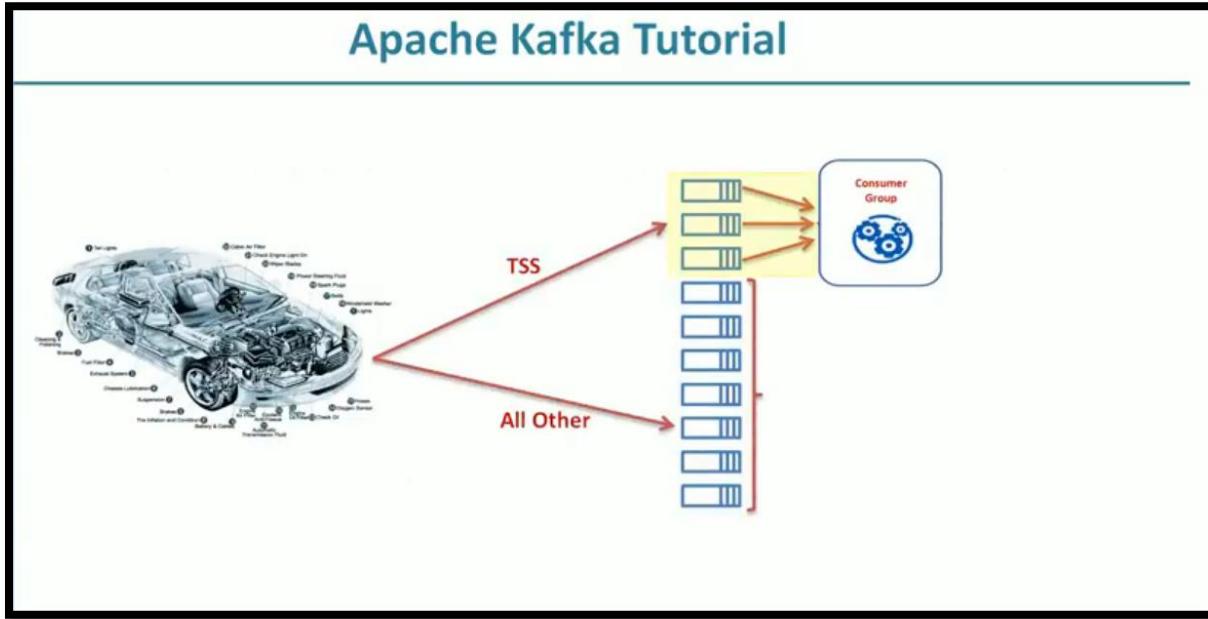
Apache Kafka Tutorial

We need to take control of following

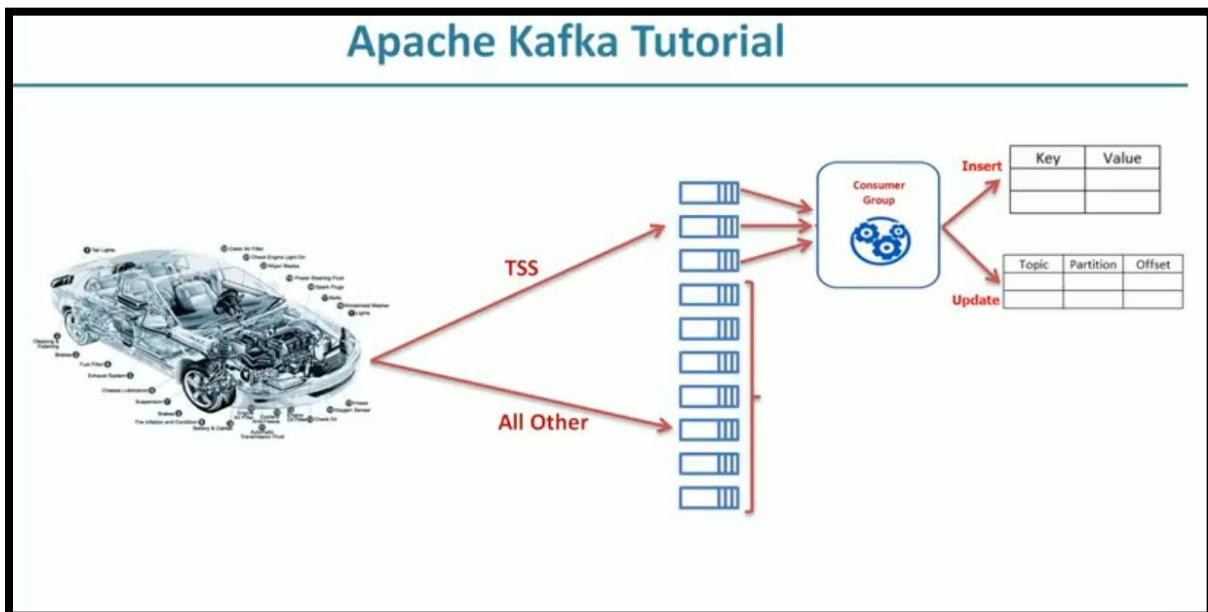
1. Partition assignment.
2. Committed Offset

- We don't want kafka to assign partition to different consumer, we want to take control and assign partition

- We don't want kafka to store the offset



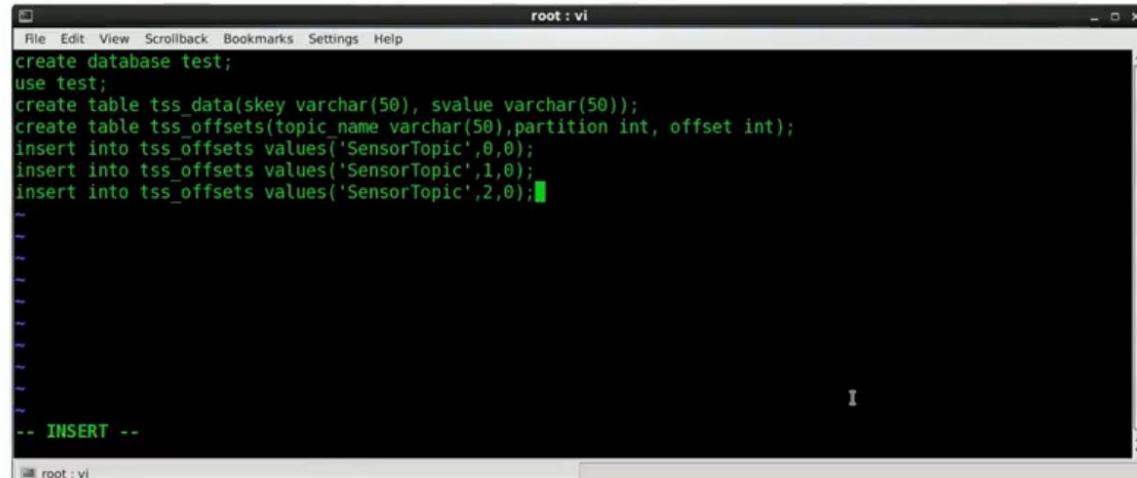
We need to create a consumer that will assign these 3 partitions to itself then it will start reading messages from these 3 partitions and insert each message into a database table, it will not commit the offset back to kafka instead it will update the current offset into another database table, the insert and update statement will be part of the single transaction so either both will complete or both will fail



Apache Kafka Tutorial

1. Create a Database.
2. Create a table to insert TSS data.
3. Create a table to update TSS offsets.
→
4. Insert three rows for 3 TSS partitions.

Apache Kafka Tutorial



```
root : vi
File Edit View Scrollback Bookmarks Settings Help
create database test;
use test;
create table tss_data(skey varchar(50), svalue varchar(50));
create table tss_offsets(topic_name varchar(50),partition int, offset int);
insert into tss_offsets values('SensorTopic',0,0);
insert into tss_offsets values('SensorTopic',1,0);
insert into tss_offsets values('SensorTopic',2,0);
-
-
-
-
-
-
-
-- INSERT --
root : vi
```

Consumer will insert the data into tss_data

Apache Kafka Tutorial



```

root : mysql
File Edit View Scrollback Bookmarks Settings Help
Query OK, 1 row affected (0.00 sec)

mysql> select * from tss_data;
Empty set (0.00 sec)

mysql> select * from tss_offset;
ERROR 1146 (42S02): Table 'test.tss_offset' doesn't exist
mysql> select * from tss_offsets;
+-----+-----+-----+
| topic_name | partition | offset |
+-----+-----+-----+
| SensorTopic | 0 | 0 |
| SensorTopic | 1 | 0 |
| SensorTopic | 2 | 0 |
+-----+-----+-----+
3 rows in set (0.00 sec)

mysql>

```

Initial offset set to 0, consumer will start reading from the beginning and keep updating this table

```

import java.util.*;
import org.apache.kafka.clients.consumer.*;
import org.apache.kafka.common.*;
import java.sql.*;

public class SensorConsumer{

    public static void main(String[] args) throws Exception{

        String topicName = "SensorTopic";
        KafkaConsumer<String, String> consumer = null;
        int rCount;

        Properties props = new Properties();
        props.put("bootstrap.servers", "localhost:9092,localhost:9093");
        props.put("key.deserializer", "org.apache.kafka.common.serialization.StringDeserializer");
        props.put("value.deserializer", "org.apache.kafka.common.serialization.StringDeserializer");
        props.put("enable.auto.commit", "false");

        consumer = new KafkaConsumer<>(props);
        TopicPartition p0 = new TopicPartition(topicName, 0);
        TopicPartition p1 = new TopicPartition(topicName, 1);
        TopicPartition p2 = new TopicPartition(topicName, 2);

        consumer.assign(Arrays.asList(p0,p1,p2));
        System.out.println("Current position p0=" + consumer.position(p0))
    }
}

```

```

        + " p1=" + consumer.position(p1)
        + " p2=" + consumer.position(p2));

    //adjusted offset to appropriate position
    consumer.seek(p0, getOffsetFromDB(p0));
    consumer.seek(p1, getOffsetFromDB(p1));
    consumer.seek(p2, getOffsetFromDB(p2));
    System.out.println("New positions po=" + consumer.position(p0)
        + " p1=" + consumer.position(p1)
        + " p2=" + consumer.position(p2));

    System.out.println("Start Fetching Now");
    try{
        do{
            ConsumerRecords<String, String> records = consumer.poll(1000);
            System.out.println("Record polled " + records.count());
            rCount = records.count();
            for (ConsumerRecord<String, String> record : records){

                saveAndCommit(consumer,record);
            }
        }while (rCount>0);
    }catch(Exception ex){
        System.out.println("Exception in main.");
    }
    finally{
        consumer.close();
    }
}

private static long getOffsetFromDB(TopicPartition p){
    long offset = 0;
    try{
        Class.forName("com.mysql.jdbc.Driver");
        Connection
con=DriverManager.getConnection("jdbc:mysql://localhost:3306/test","root","pandey");

        String sql = "select offset from tss_offsets where topic_name='" + p.topic() + "' and partition="
+ p.partition();
        Statement stmt=con.createStatement();
        ResultSet rs = stmt.executeQuery(sql);
        if (rs.next())
            offset = rs.getInt("offset");
        stmt.close();
        con.close();
    }catch(Exception e){
        System.out.println("Exception in getOffsetFromDB");
    }
}

```

```
        return offset;
    }

    private static void saveAndCommit(KafkaConsumer<String, String> c, ConsumerRecord<String, String> r){
        System.out.println("Topic=" + r.topic() + " Partition=" + r.partition() + " Offset=" + r.offset() + " Key=" + r.key() + " Value=" + r.value());
        try{
            Class.forName("com.mysql.jdbc.Driver");
            Connection con=DriverManager.getConnection("jdbc:mysql://localhost:3306/test","root","pandey");
            con.setAutoCommit(false);

            String insertSQL = "insert into tss_data values(?,?)";
            PreparedStatement psInsert = con.prepareStatement(insertSQL);
            psInsert.setString(1,r.key());
            psInsert.setString(2,r.value());

            String updateSQL = "update tss_offsets set offset=? where topic_name=? and partition=?";
            PreparedStatement psUpdate = con.prepareStatement(updateSQL);
            psUpdate.setLong(1,r.offset()+1);
            psUpdate.setString(2,r.topic());
            psUpdate.setInt(3,r.partition());

            psInsert.executeUpdate();
            psUpdate.executeUpdate();
            con.commit();
            con.close();
        }catch(Exception e){
            System.out.println("Exception in saveAndCommit");
        }
    }
}
```

Apache Kafka Tutorial

The screenshot shows two terminal windows side-by-side. The left window displays the output of a Kafka producer command, showing log entries for keys and partitions. The right window shows the output of a Kafka consumer command, displaying log entries for topics and offsets.

Left Terminal (Producer Log):

```
javaProj : bash
File Edit View Scrollback Bookmarks Settings Help
Key = SSP3 Partition = 6
Key = SSP4 Partition = 3
Key = SSP5 Partition = 8
Key = SSP6 Partition = 3
Key = SSP7 Partition = 9
Key = SSP8 Partition = 6
Key = SSP9 Partition = 7
Key = TSS Partition = 1
Key = TSS Partition = 0
Key = TSS Partition = 0
Key = TSS Partition = 2
Key = TSS Partition = 0
Key = TSS Partition = 1
Key = TSS Partition = 0
Key = TSS Partition = 0
Key = TSS Partition = 1
Key = TSS Partition = 1
SimpleProducer Completed.
[success] Total time: 16 s, completed Jan 13, 2017 9:33:24 PM
```

Right Terminal (Consumer Log):

```
javaProj : bash
File Edit View Scrollback Bookmarks Settings Help
^ SLF4J: Failed to load class "org.slf4j.impl.StaticLoggerBinder".
SLF4J: Defaulting to no-operation (NOP) logger implementation
SLF4J: See http://www.slf4j.org/codes.html#StaticLoggerBinder for further details.
Current position p0=5 p1=4 p2=1
New positions p0=0 p1=0 p2=0
Start Fetching Now
Record polled 6
Topic=SensorTopic Partition=0 Offset=0 Key=TSS Value=5001
Topic=SensorTopic Partition=0 Offset=1 Key=TSS Value=5002
Topic=SensorTopic Partition=0 Offset=2 Key=TSS Value=5004
Topic=SensorTopic Partition=0 Offset=3 Key=TSS Value=5006
Topic=SensorTopic Partition=0 Offset=4 Key=TSS Value=5007
Topic=SensorTopic Partition=2 Offset=0 Key=TSS Value=5003
Record polled 4
Topic=SensorTopic Partition=1 Offset=0 Key=TSS Value=5000
Topic=SensorTopic Partition=1 Offset=1 Key=TSS Value=5005
Topic=SensorTopic Partition=1 Offset=2 Key=TSS Value=5008
Topic=SensorTopic Partition=1 Offset=3 Key=TSS Value=5009
```

Apache Kafka Tutorial

The screenshot shows two terminal windows. The left window displays the output of a Kafka producer command. The right window shows a MySQL database session where a query is run against a table named 'tss_offsets'.

Left Terminal (Producer Log):

```
javaProj : bash
File Edit View Scrollback Bookmarks Settings Help
Key = SSP3 Partition = 6
Key = SSP4 Partition = 3
Key = SSP5 Partition = 8
Key = SSP6 Partition = 3
Key = SSP7 Partition = 9
Key = SSP8 Partition = 6
Key = SSP9 Partition = 7
Key = TSS Partition = 1
Key = TSS Partition = 0
Key = TSS Partition = 0
Key = TSS Partition = 2
Key = TSS Partition = 0
Key = TSS Partition = 1
Key = TSS Partition = 0
Key = TSS Partition = 0
Key = TSS Partition = 1
Key = TSS Partition = 1
SimpleProducer Completed.
[success] Total time: 16 s, completed Jan 13, 2017 9:33:24 PM
```

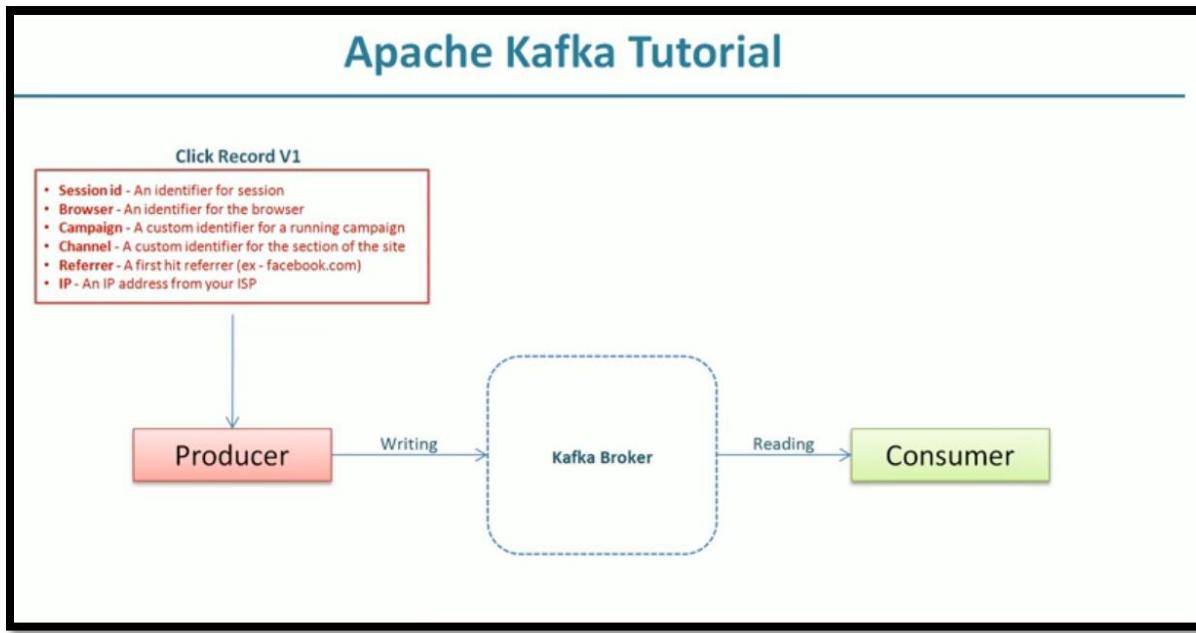
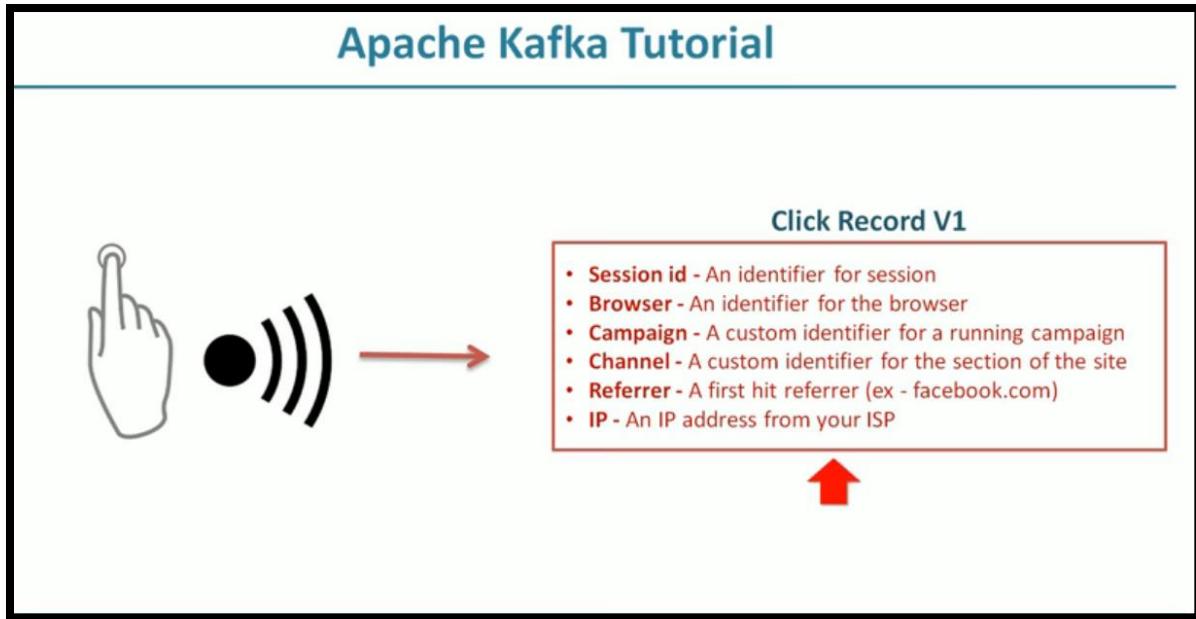
Right Terminal (MySQL Query):

```
javaProj : bash
File Edit View Scrollback Bookmarks Settings Help
mysql> use test;
Reading table information for completion of table and column names
You can turn off this feature to get a quicker startup with -A
Database changed
mysql> select * from tss_offsets;
+-----+-----+-----+
| topic_name | partition | offset |
+-----+-----+-----+
| SensorTopic | 0 | 5 |
| SensorTopic | 1 | 4 |
| SensorTopic | 2 | 1 |
+-----+-----+-----+
3 rows in set (0.00 sec)
```

A red arrow points to the value '5' in the 'offset' column of the first row of the MySQL query result.

Schema Evolution

Schema evolution dealing with changes in message records over the period of time



Apache Kafka Tutorial

Click Record V1

- Session id - An identifier for session
- Browser - An identifier for the browser
- Campaign - A custom identifier for a running campaign
- Channel - A custom identifier for the section of the site
- Referrer - A first hit referrer (ex - facebook.com)
- IP - An IP address from your ISP

Producer

Serializer

Writing

Kafka Broker

Reading

Consumer

Deserializer

Apache Kafka Tutorial

Click Record V1

- Session id - An identifier for session
- Browser - An identifier for the browser
- Campaign - A custom identifier for a running campaign
- Channel - A custom identifier for the section of the site
- Referrer - A first hit referrer (ex - facebook.com)
- IP - An IP address from your ISP

Click Record V2

- Session id - An identifier for session
- Browser - An identifier for the browser
- Campaign - A custom identifier for a running campaign
- Channel - A custom identifier for the section of the site
- Referrer - A first hit referrer (ex - facebook.com)
- Entry URL - A first hit referrer URL
- IP - An IP address from your ISP
- Language - An identifier for the language
- OS ID - An identifier for the operating system

Apache Kafka Tutorial

Click Record V1

- Session id - An identifier for session
- Browser - An identifier for the browser
- Campaign - A custom identifier for a running campaign
- Channel - A custom identifier for the section of the site
- Referrer - A first hit referrer (ex - facebook.com)
- IP - An IP address from your ISP

Producer

Serializer

Writing

Kafka Broker

Consumer

Deserializer

Reading

Do I need to change all of
my current Producers?

Apache Kafka Tutorial

Click Record V1

- Session id - An identifier for session
- Browser - An identifier for the browser
- Campaign - A custom identifier for a running campaign
- Channel - A custom identifier for the section of the site
- Referrer - A first hit referrer (ex - facebook.com)
- IP - An IP address from your ISP

Producer

Serializer

Writing

Kafka Broker

Consumer

Deserializer

I don't want to change
this immediately.

Click Record V2

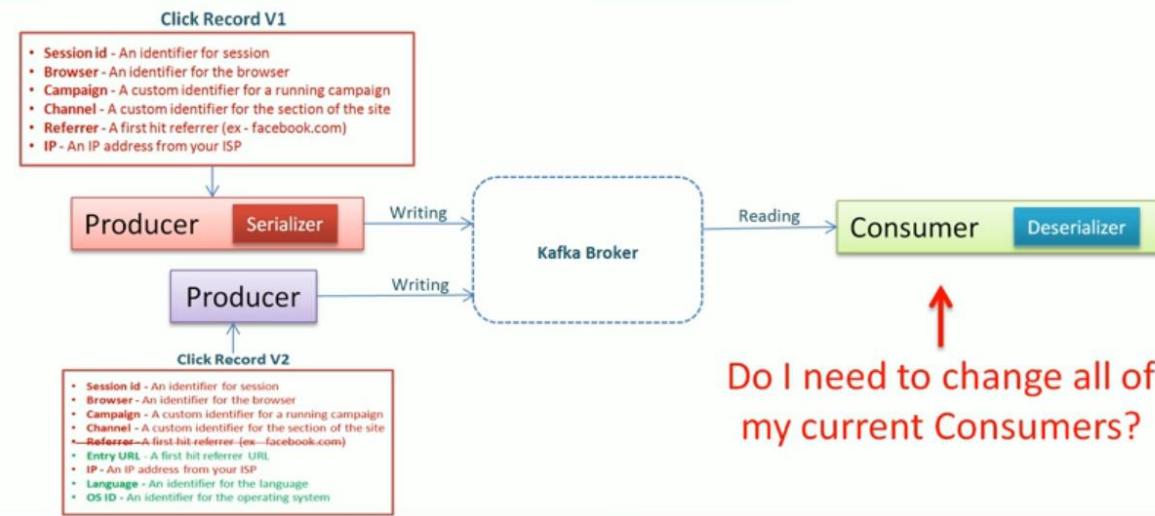
- Session id - An identifier for session
- Browser - An identifier for the browser
- Campaign - A custom identifier for a running campaign
- Channel - A custom identifier for the section of the site
- Referrer - A first hit referrer (ex - facebook.com)
- Entry URL - A first hit referrer URL
- IP - An IP address from your ISP
- Language - An identifier for the language
- OS ID - An identifier for the operating system

Producer

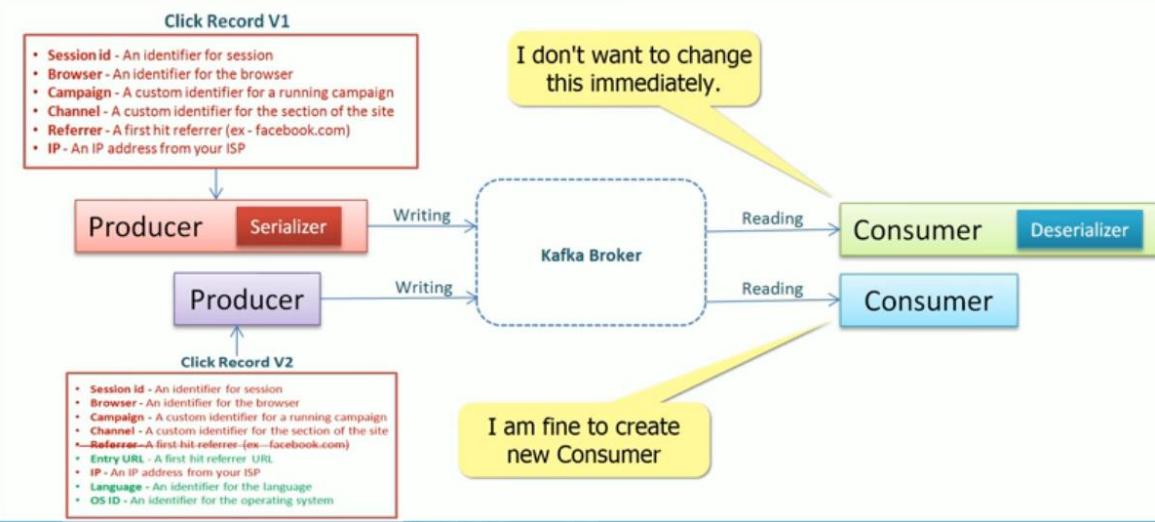
Writing

I am fine to create
new Producer

Apache Kafka Tutorial

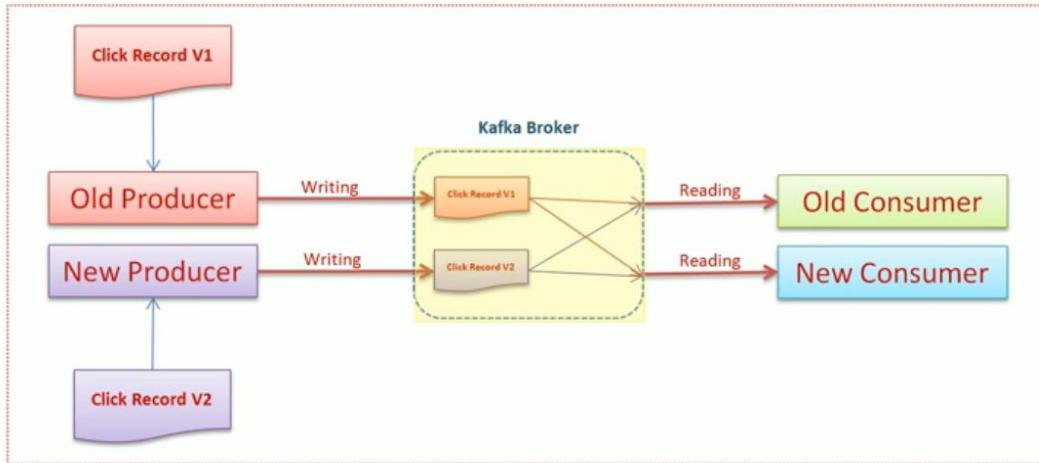


Apache Kafka Tutorial



Apache Kafka Tutorial

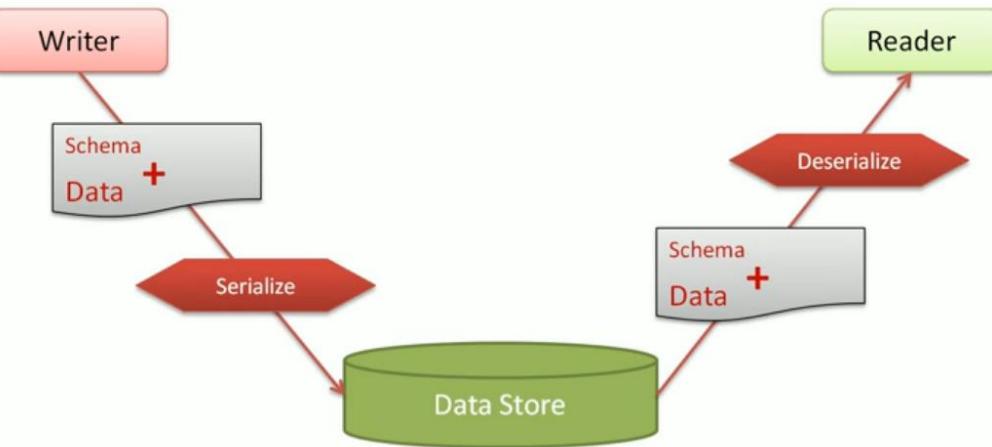
Schema Evolution



We want to support both old and new schemas simultaneously

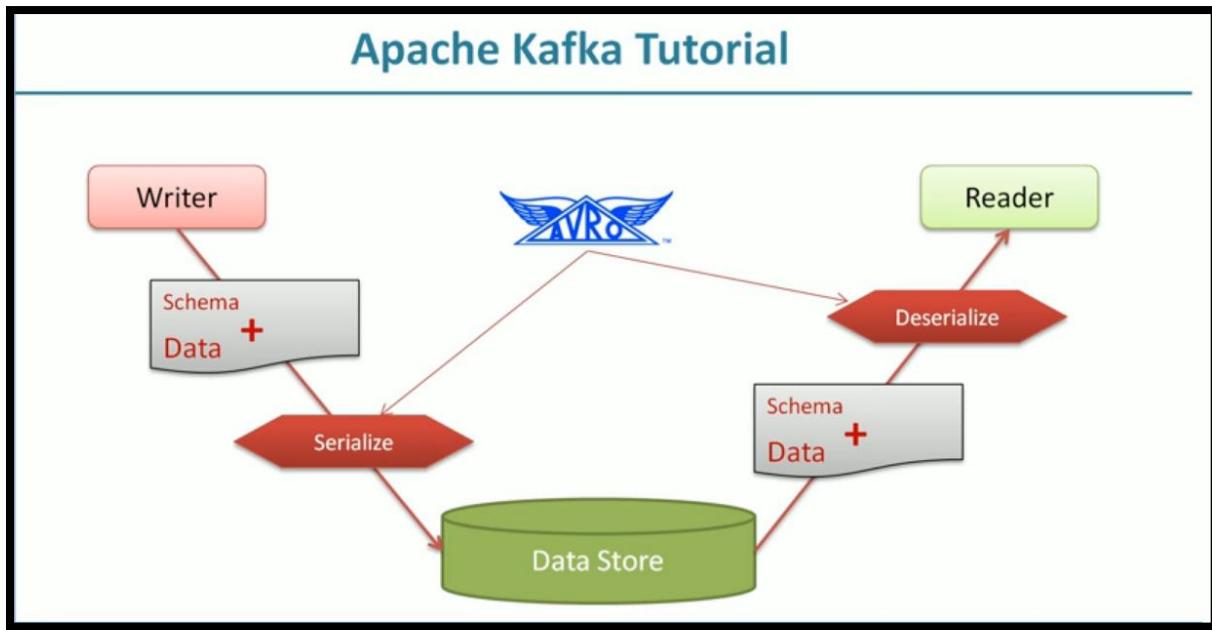
We need to have a combination of old and new producer as well as combination of old and new consumer, Kafka should be able to store both type of messages on the same topic. Consumer should be read both types of messages without any error

Apache Kafka Tutorial



The Industry solution to handle schema evolution is to include schema information with the data when someone is writing data should write both schema and data and when someone wants to read the data they first read schema and then read data based on the schema

There are prebuilt and reusable serialization system to help us and simplify the whole process of translation messages according to schema and embedding schema information in the message records



AVRO is one of them it is most popular serialization system for Hadoop and its ecosystem, Kafka follows the same approach and uses Avro to handle the schema evolution problem

See below URL

<http://www-us.apache.org/dist/avro/stable/java/>

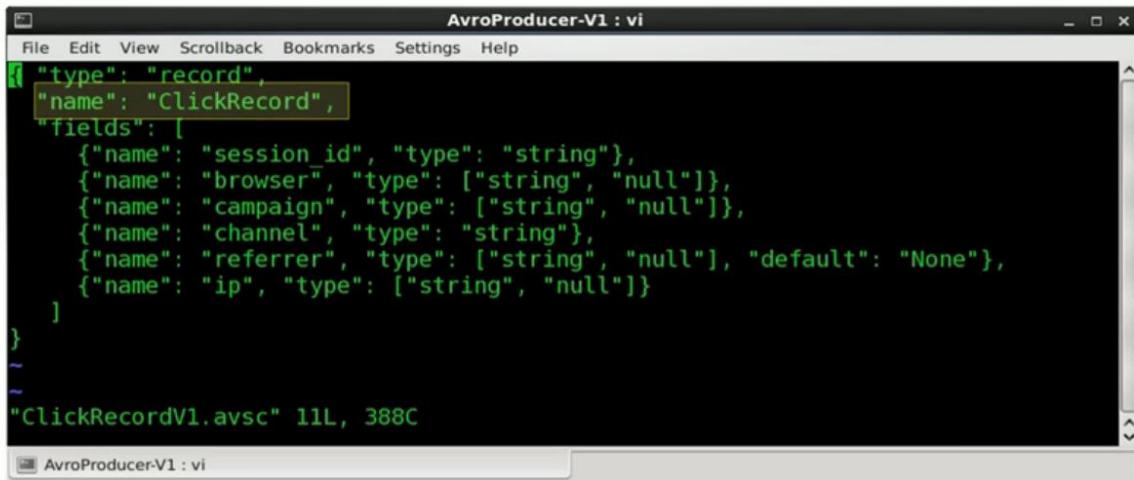
<http://avro.apache.org/releases.html#Download>

Apache Kafka Tutorial

A Typical Serialization System offers following

- - 1. Allows you to define a schema for your data.
 - 2. Generates code for your schema. (Optional in Avro)
 - 3. Provide APIs to serialize your data according to the schema and embed schema information in the data.
 - 4. Provide APIs to extract schema information and deserialize your data based on the schema.

Apache Kafka Tutorial



```
"type": "record",
"name": "ClickRecord",
"fields": [
    {"name": "session_id", "type": "string"},
    {"name": "browser", "type": ["string", "null"]},
    {"name": "campaign", "type": ["string", "null"]},
    {"name": "channel", "type": "string"},
    {"name": "referrer", "type": ["string", "null"], "default": "None"},
    {"name": "ip", "type": ["string", "null"]}
]
}

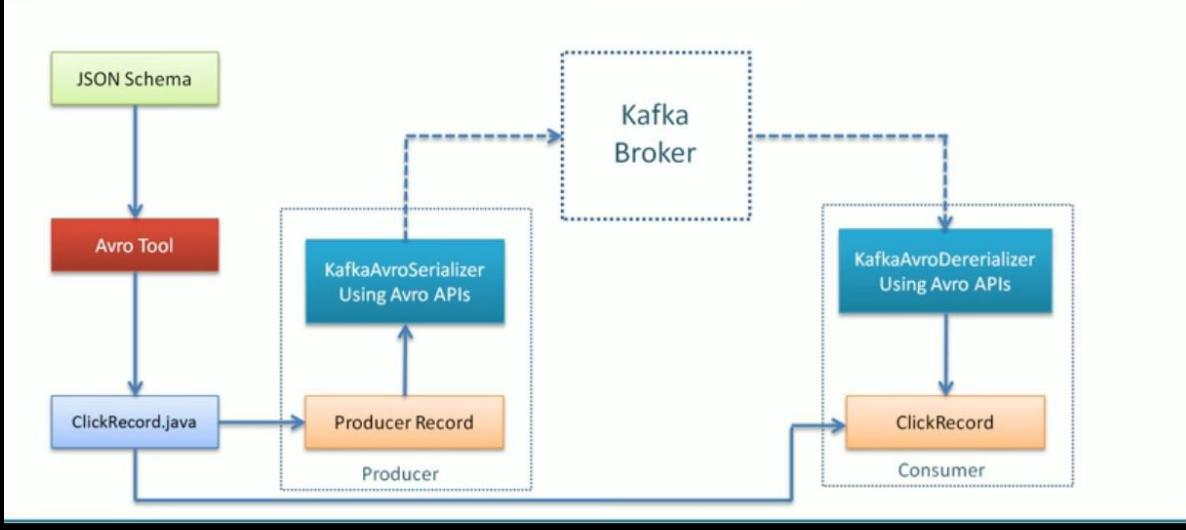
"ClickRecordV1.avsc" 11L, 388C
```

Schema definition

Apache Kafka Tutorial

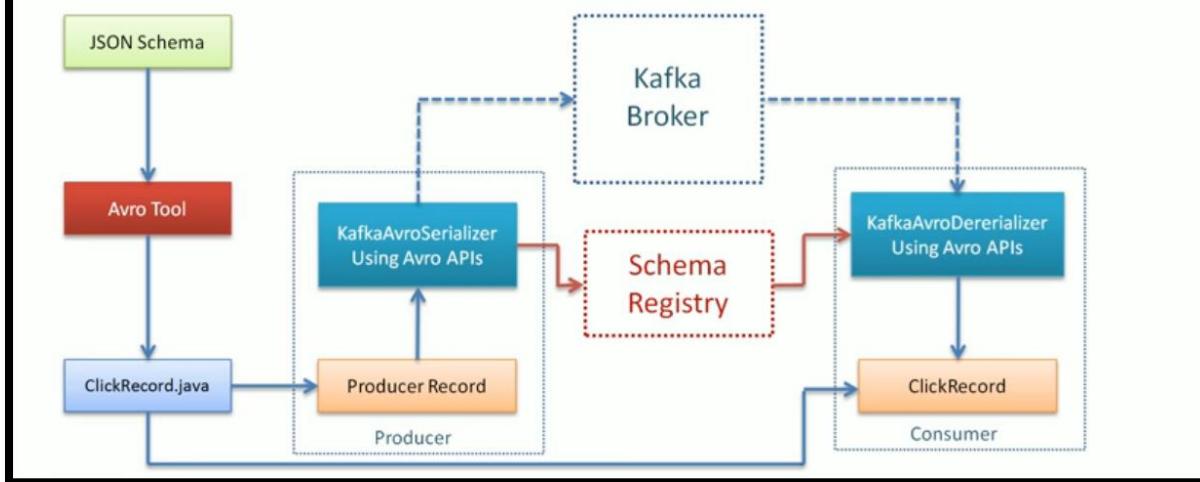
```
File Edit View Scrollback Bookmarks Settings Help
[root@localhost AvroProducer-V1]# ll
total 33780
-rw-r--r--. 1 root root 34583098 Jan 22 21:18 avro-tools-1.8.1.jar
-rw-r--r--. 1 root root      388 Jan 22 21:16 ClickRecordV1.avsc
[root@localhost AvroProducer-V1]# java -jar avro-tools-1.8.1.jar compile schema ClickRecordV1.avsc .
Input files to compile:
ClickRecordV1.avsc
log4j:WARN No appenders could be found for logger (AvroVelocityLogChute).
log4j:WARN Please initialize the log4j system properly.
log4j:WARN See http://logging.apache.org/log4j/1.2/faq.html#noconfig for more info.
[root@localhost AvroProducer-V1]# ll
total 33800
-rw-r--r--. 1 root root 34583098 Jan 22 21:18 avro-tools-1.8.1.jar
-rw-r--r--. 1 root root     16474 Jan 22 22:07 ClickRecord.java ←
-rw-r--r--. 1 root root      388 Jan 22 21:16 ClickRecordV1.avsc
[root@localhost AvroProducer-V1]#
```

Apache Kafka Tutorial



The deserializer should now know the schema without knowing the schema it can't deserialize raw bytes back to an object that's where the schema registry is useful the kafka avro serializer will store the schema definition into the schema registry and include an ID of the schema into the message record when kafka avro deserializer sees the message it takes the schema id from the message and get the schema definition from the schema registry, once we have the schema details and message bytes it is simple to deserialize them.

Apache Kafka Tutorial



Apache Kafka Tutorial

How to use Avro in Kafka

- ➡ • Define an Avro Schema for your message record.
- Generate a source code for your Avro Schema.
- Create a producer and use KafkaAvroSerializer.
- Create a consumer and use KafkaAvroDeserializer.

```
import java.util.*;
import org.apache.kafka.clients.producer.*;
public class AvroProducer {

    public static void main(String[] args) throws Exception{
        String topicName = "AvroClicks";
        String msg;
```

```

Properties props = new Properties();
props.put("bootstrap.servers", "localhost:9092,localhost:9093");
props.put("key.serializer", "org.apache.kafka.common.serialization.StringSerializer");
props.put("value.serializer", "io.confluent.kafka.serializers.KafkaAvroSerializer");
//  

props.put("schema.registry.url", "http://localhost:8081");

Producer<String, ClickRecord> producer = new KafkaProducer <>(props);
ClickRecord cr = new ClickRecord();
try{
    cr.setSessionId("10001");
    cr.setChannel("HomePage");
    cr.setIp("192.168.0.1");

    producer.send(new ProducerRecord<String,
ClickRecord>(topicName,cr.getSessionId().toString(),cr).get();

    System.out.println("Complete");
}
catch(Exception ex){
    ex.printStackTrace(System.out);
}
finally{
    producer.close();
}

}
}

```

ClickRecordV1.avsc

```

{ "type":  

"record",  

    "name": "ClickRecord",  

    "fields": [  

        {"name": "session_id", "type": "string"},  

        {"name": "browser", "type": ["string", "null"]},  

        {"name": "campaign", "type": ["string", "null"]},  

        {"name": "channel", "type": "string"},  

        {"name": "referrer", "type": ["string", "null"], "default": "None"},  

        {"name": "ip", "type": ["string", "null"]}  

    ]  

}

```

```

import java.util.*;
import org.apache.kafka.clients.consumer.*;

public class AvroConsumer{

    public static void main(String[] args) throws Exception{
        String topicName = "AvroClicks";

        String groupName = "RG";
        Properties props = new Properties();
        props.put("bootstrap.servers", "localhost:9092,localhost:9093");
        props.put("group.id", groupName);
        props.put("key.deserializer", "org.apache.kafka.common.serialization.StringDeserializer");
        props.put("value.deserializer", "io.confluent.kafka.serializers.KafkaAvroDeserializer");
        props.put("schema.registry.url", "http://localhost:8081");
        //this is mandatory for avro read
        props.put("specific.avro.reader", "true");

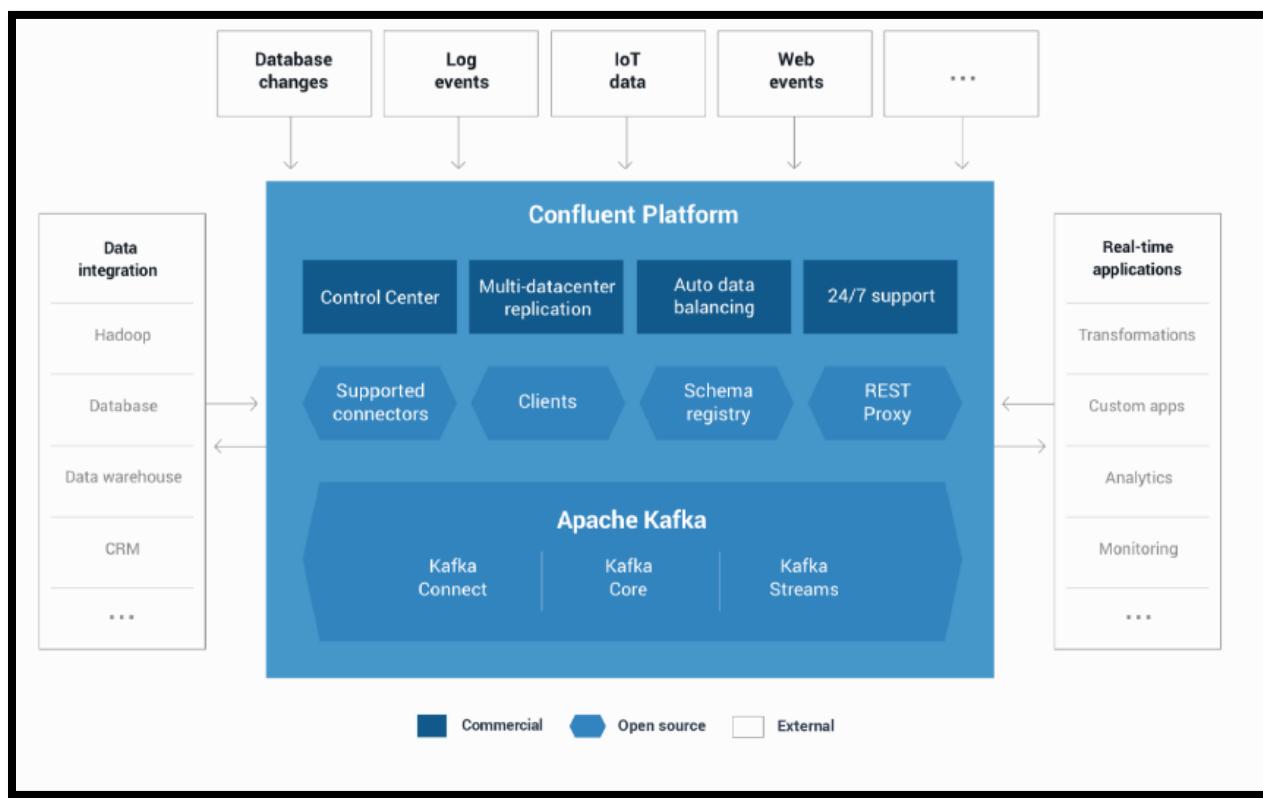
        KafkaConsumer<String, ClickRecord> consumer = new KafkaConsumer<>(props);
        consumer.subscribe(Arrays.asList(topicName));
        try{
            while (true){
                ConsumerRecords<String, ClickRecord> records = consumer.poll(100);
                for (ConsumerRecord<String, ClickRecord> record : records){
                    System.out.println("Session id=" + record.value().getSessionId()
                        + " Channel=" + record.value().getChannel()
                        + " Referrer=" + record.value().getReferrer());
                }
            }
        }catch(Exception ex){
            ex.printStackTrace();
        }
        finally{
            consumer.close();
        }
    }
}

```

For that we need to download confluent platform

Confluent Platform

The Confluent Platform makes it easy build real-time data pipelines and streaming applications. By integrating data from multiple sources and locations into a single, central streaming platform for your company, the Confluent Platform lets you focus on how to derive business value from your data rather than worrying about the underlying mechanics such as how data is being transported or massaged between various systems. Specifically, the Confluent Platform simplifies connecting data sources to Kafka, building applications with Kafka, as well as securing, monitoring, and managing your Kafka infrastructure.

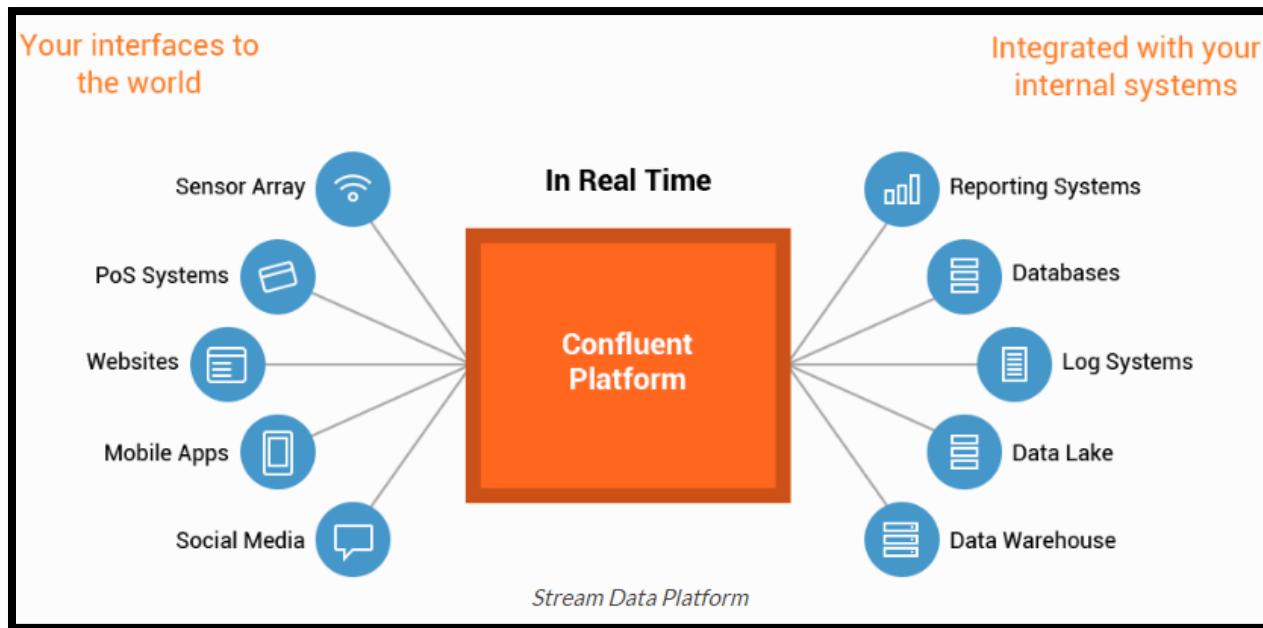


As its backbone, the Confluent Platform leverages Apache Kafka, the most popular open source distributed streaming platform. It has three key capabilities:

- It lets you publish and subscribe to streams of records
- It lets you store streams of records in a fault tolerant way
- It lets you process streams of records

In confluent platform Apache Kafka package together with additional components, all these elements together with apache Kafka make it a powerful and flexible streaming platform

The Confluent Platform is a streaming platform that enables you to organize and manage data from many different sources with one reliable, high performance system.



Confluent Schema Registry

One of the most difficult challenges with loosely coupled systems is ensuring compatibility of data and code as the system grows and evolves. With a messaging service like Kafka, services that interact with each other must agree on a common format, called a *schema*, for messages. In many systems, these formats are ad hoc, only implicitly defined by the code, and often are duplicated across each system that uses that message type.

As requirements change, it becomes necessary to evolve these formats. With only an ad-hoc definition, it is very difficult for developers to determine what the impact of their change might be.

The [Confluent Schema Registry](#) enables safe, zero downtime evolution of schemas by centralizing the management of schemas written for the [Avro](#) serialization system. It tracks all versions of schemas used for every topic in Kafka and only allows evolution of schemas according to user-defined compatibility settings. This gives developers confidence that they can safely modify schemas as necessary without worrying that doing so will break a different service they may not even be aware of.

The Schema Registry also includes plugins for Kafka clients that handle schema storage and retrieval for Kafka messages that are sent in the Avro format. This integration is seamless – if you are already using Kafka with Avro data, using the Schema Registry only requires including the serializers with your application and changing one setting.

The Confluent Schema Registry is available as open source software under the [Apache License v2.0 license](#). Source code at <https://github.com/confluentinc/schema-registry>.

The Confluent Schema Registry is available as open source software under the [Apache License v2.0 license](#). Source code at <https://github.com/confluentinc/schema-registry>.

Installation Step on Centos7

Please follow below link for install confluent platform

<http://docs.confluent.io/3.1.2/installation.html>

First install Confluent's public key, which is used to sign packages in the yum repository.

```
$ sudo rpm --import http://packages.confluent.io/rpm/3.1/archive.key
```

Add the repository to your **/etc/yum.repos.d/** directory in a file named **confluent.repo**

If you are using RHEL/Centos/Oracle 7

```
[Confluent.dist]
name=Confluent repository (dist)
baseurl=http://packages.confluent.io/rpm/3.1/
gpgcheck=1
gpgkey=http://packages.confluent.io/rpm/3.1/archive.key
enabled=1

[Confluent]
name=Confluent repository
baseurl=http://packages.confluent.io/rpm/3.1
gpgcheck=1
gpgkey=http://packages.confluent.io/rpm/3.1/archive.key
enabled=1
```

It is recommended to clear the yum caches before proceeding:

```
$ sudo yum clean all
```

The repository is now ready for use.

You can install Confluent Open Source with:

```
$ sudo yum install confluent-platform-oss-2.11
```

The number at the end of the package name specifies the Scala version. Currently supported versions are 2.11 (recommended) and 2.10.

Here 2.11 is Scala version

```
[root@mac92 yum.repos.d]# sudo yum install confluent-platform-oss-2.11
Loaded plugins: fastestmirror, langpacks
Repository HDP-UTILS-1.1.0.20 is listed more than once in the configuration
Confluent | 2.9 kB 00:00:00
Confluent.dist | 2.9 kB 00:00:00
HDP-2.3 | 2.9 kB 00:00:00
HDP-2.3.4.0 | 2.9 kB 00:00:00
HDP-UTILS-1.1.0.20 | 2.9 kB 00:00:00
Updates-ambari-2.2.0.0 | 2.9 kB 00:00:00
base | 3.6 kB 00:00:00
cloudera-cdh5 | 951 B 00:00:00
cloudera-manager | 2.9 kB 00:00:00
epel/x86_64/metalink | 6.1 kB 00:00:00
epel | 4.3 kB 00:00:00
extras | 3.4 kB 00:00:00
mesosphere | 2.9 kB 00:00:00
mesosphere-noarch | 2.9 kB 00:00:00
mysql-connectors-community | 2.5 kB 00:00:00
```

mysql-tools-community	2.5 kB 00:00:00
mysql56-community	2.5 kB 00:00:00
updates	3.4 kB 00:00:00
(1/19): HDP-2.3.4.0/primary_db	61 kB 00:00:00
(2/19): HDP-2.3/primary_db	61 kB 00:00:00
(3/19): Updates-ambari-2.2.0.0/primary_db	5.8 kB 00:00:00
(4/19): HDP-UTILS-1.1.0.20/primary_db	28 kB 00:00:00
(5/19): cloudera-manager/primary_db	11 kB 00:00:00
(6/19): Confluent/primary_db	17 kB 00:00:00
(7/19): Confluent.dist/primary_db	8.5 kB 00:00:00
(8/19): base/7/x86_64/group_gz	155 kB 00:00:00
(9/19): extras/7/x86_64/primary_db	121 kB 00:00:00
(10/19): mesosphere/x86_64/primary_db	35 kB 00:00:00
(11/19): mesosphere-noarch/primary_db	2.5 kB 00:00:00
(12/19): mysql-connectors-community/x86_64/primary_db 00:00:00	13 kB
(13/19): mysql-tools-community/x86_64/primary_db 00:00:00	32 kB
(14/19): mysql56-community/x86_64/primary_db	159 kB 00:00:00
(15/19): base/7/x86_64/primary_db	5.6 MB 00:00:01
(16/19): epel/x86_64/group_gz	170 kB 00:00:01
(17/19): updates/7/x86_64/primary_db	2.2 MB 00:00:01
(18/19): epel/x86_64/updateinfo	732 kB 00:00:03
(19/19): epel/x86_64/primary_db	4.5 MB 00:00:10

cloudera-cdh5/primary	43 kB 00:00:00
Determining fastest mirrors	
* base: centos.excellmedia.net	
* epel: epel.mirror.net.in	
* extras: centos.excellmedia.net	
* updates: centos.excellmedia.net	
cloudera-cdh5	146/146
Resolving Dependencies	
--> Running transaction check	
---> Package confluent-platform-oss-2.11.noarch 0:3.1.2-1 will be installed	
--> Processing Dependency: confluent-kafka-connect-elasticsearch >= 3.1.2 for package: confluent-platform-oss-2.11-3.1.2-1.noarch	
--> Processing Dependency: confluent-kafka-connect-hdfs >= 3.1.2 for package: confluent-platform-oss-2.11-3.1.2-1.noarch	
--> Processing Dependency: confluent-camus >= 3.1.2 for package: confluent-platform-oss-2.11-3.1.2-1.noarch	
--> Processing Dependency: confluent-kafka-connect-jdbc >= 3.1.2 for package: confluent-platform-oss-2.11-3.1.2-1.noarch	
--> Processing Dependency: confluent-schema-registry >= 3.1.2 for package: confluent-platform-oss-2.11-3.1.2-1.noarch	
--> Processing Dependency: confluent-kafka-2.11 >= 0.10.1.1 for package: confluent-platform-oss-2.11-3.1.2-1.noarch	
--> Processing Dependency: confluent-kafka-rest >= 3.1.2 for package: confluent-platform-oss-2.11-3.1.2-1.noarch	
--> Running transaction check	
---> Package confluent-camus.noarch 0:3.1.2-1 will be installed	
---> Package confluent-kafka-2.11.noarch 0:0.10.1.1-1 will be installed	

```
--> Package confluent-kafka-connect-elasticsearch.noarch 0:3.1.2-1 will be installed  
  
--> Processing Dependency: confluent-common for package: confluent-kafka-connect-elasticsearch-3.1.2-1.noarch  
  
--> Package confluent-kafka-connect-hdfs.noarch 0:3.1.2-1 will be installed  
  
--> Package confluent-kafka-connect-jdbc.noarch 0:3.1.2-1 will be installed  
  
--> Package confluent-kafka-rest.noarch 0:3.1.2-1 will be installed  
  
--> Processing Dependency: confluent-rest-utils for package: confluent-kafka-rest-3.1.2-1.noarch  
  
--> Package confluent-schema-registry.noarch 0:3.1.2-1 will be installed  
  
--> Running transaction check  
  
--> Package confluent-common.noarch 0:3.1.2-1 will be installed  
  
--> Package confluent-rest-utils.noarch 0:3.1.2-1 will be installed  
  
--> Finished Dependency Resolution
```

Dependencies Resolved

Package	Arch	Version	Repository	Size
<hr/>				
Installing:				
confluent-platform-oss-2.11	noarch	3.1.2-1	Confluent	6.6 k
Installing for dependencies:				
confluent-camus	noarch	3.1.2-1	Confluent	19 M
confluent-common	noarch	3.1.2-1	Confluent	2.0 M
confluent-kafka-2.11	noarch	0.10.1.1-1	Confluent	36 M

confluent-kafka-connect-elasticsearch M	noarch	3.1.2-1	Confluent	4.3
confluent-kafka-connect-hdfs	noarch	3.1.2-1	Confluent	86 M
confluent-kafka-connect-jdbc	noarch	3.1.2-1	Confluent	6.0 M
confluent-kafka-rest	noarch	3.1.2-1	Confluent	16 M
confluent-rest-utils	noarch	3.1.2-1	Confluent	7.0 M
confluent-schema-registry	noarch	3.1.2-1	Confluent	26 M
<hr/>				
Transaction Summary				
<hr/> <hr/>				
Install 1 Package (+9 Dependent packages)				
<hr/> <hr/>				
Total download size: 203 M				
Installed size: 229 M				
Is this ok [y/d/N]: y				
Downloading packages:				
(1/10): confluent-common-3.1.2-1.noarch.rpm			2.0 MB	00:00:04
(2/10): confluent-kafka-2.11-0.10.1.1-1.noarch.rpm			36 MB	00:00:24
(3/10): confluent-camus-3.1.2-1.noarch.rpm			19 MB	00:00:33
(4/10): confluent-kafka-connect-elasticsearch-3.1.2-1.noarch.rpm 00:00:07			4.3 MB	
(5/10): confluent-kafka-connect-jdbc-3.1.2-1.noarch.rpm 00:00:07			6.0 MB	
(6/10): confluent-kafka-rest-3.1.2-1.noarch.rpm			16 MB	00:00:09

(7/10): confluent-platform-oss-2.11-3.1.2-1.noarch.rpm	6.6 kB 00:00:00
(8/10): confluent-rest-utils-3.1.2-1.noarch.rpm	7.0 MB 00:00:12
(9/10): confluent-schema-registry-3.1.2-1.noarch.rpm	26 MB 00:00:25
(10/10): confluent-kafka-connect-hdfs-3.1.2-1.noarch.rpm 00:01:00	86 MB
<hr/>	
Total	2.2 MB/s 203 MB 00:01:33
Running transaction check	
Running transaction test	
Transaction test succeeded	
Running transaction	
Installing : confluent-common-3.1.2-1.noarch	1/10
Installing : confluent-rest-utils-3.1.2-1.noarch	2/10
Installing : confluent-schema-registry-3.1.2-1.noarch	3/10
Installing : confluent-kafka-rest-3.1.2-1.noarch	4/10
Installing : confluent-kafka-connect-jdbc-3.1.2-1.noarch	5/10
Installing : confluent-kafka-connect-elasticsearch-3.1.2-1.noarch	6/10
Installing : confluent-kafka-connect-hdfs-3.1.2-1.noarch	7/10
Installing : confluent-camus-3.1.2-1.noarch	8/10
Installing : confluent-kafka-2.11-0.10.1.1-1.noarch	9/10
Installing : confluent-platform-oss-2.11-3.1.2-1.noarch	10/10
Verifying : confluent-kafka-2.11-0.10.1.1-1.noarch	1/10
Verifying : confluent-platform-oss-2.11-3.1.2-1.noarch	2/10
Verifying : confluent-schema-registry-3.1.2-1.noarch	3/10

Verifying :confluent-kafka-rest-3.1.2-1.noarch	4/10
Verifying :confluent-kafka-connect-jdbc-3.1.2-1.noarch	5/10
Verifying :confluent-kafka-connect-elasticsearch-3.1.2-1.noarch	6/10
Verifying :confluent-camus-3.1.2-1.noarch	7/10
Verifying :confluent-rest-utils-3.1.2-1.noarch	8/10
Verifying :confluent-kafka-connect-hdfs-3.1.2-1.noarch	9/10
Verifying :confluent-common-3.1.2-1.noarch	10/10

Installed:

confluent-platform-oss-2.11.noarch 0:3.1.2-1

Dependency Installed:

confluent-camus.noarch 0:3.1.2-1

confluent-common.noarch 0:3.1.2-1

confluent-kafka-2.11.noarch 0:0.10.1.1-1
0:3.1.2-1

confluent-kafka-connect-elasticsearch.noarch

confluent-kafka-connect-hdfs.noarch 0:3.1.2-1

confluent-kafka-connect-jdbc.noarch 0:3.1.2-1

confluent-kafka-rest.noarch 0:3.1.2-1

confluent-rest-utils.noarch 0:3.1.2-1

confluent-schema-registry.noarch 0:3.1.2-1

Complete!

[root@mac92 yum.repos.d]#

Components are installed under **/etc/kafka**, **/etc/kafka-rest**, **/etc/schema-registry** and **/etc/camus**.

Start zookeeper, Kafka and schema registry

```
$ zookeeper-server-start /etc/kafka/zookeeper.properties
$ kafka-server-start /etc/kafka/server.properties
$ schema-registry-start /etc/schema-registry/schema-registry.properties
```

Schema registry default listen on 8081 check schema-registry.properties file

```
[root@mac92 schema-registry]# schema-registry-start /etc/schema-registry/schema-registry.properties
[2017-02-08 09:17:49,139] INFO SchemaRegistryConfig values:
metric.reporters = []
kafkastore.sasl.kerberos.kinit.cmd = /usr/bin/kinit
response.mediatype.default = application/vnd.schemaregistry.v1+json
kafkastore.ssl.trustmanager.algorithm = PKIX
authentication.realm =
ssl.keystore.type = JKS
kafkastore.topic = _schemas
metrics.jmx.prefix = kafka.schema.registry
kafkastore.ssl.enabled.protocols = TLSv1.2,TLSv1.1,TLSv1
kafkastore.topic.replication.factor = 3
ssl.truststore.password =
kafkastore.timeout.ms = 500
host.name = mac92.cybage.com
kafkastore.bootstrap.servers = []
schema.registry.zk.namespace = schema_registry
kafkastore.sasl.kerberos.ticket.renew.window.factor = 0.8
kafkastore.sasl.kerberos.service.name =
ssl.endpoint.identification.algorithm =
compression.enable = false
kafkastore.ssl.truststore.type = JKS
avro.compatibility.level = backward
kafkastore.ssl.protocol = TLS
kafkastore.ssl.provider =
kafkastore.ssl.truststore.location =
response.mediatype.preferred = [application/vnd.schemaregistry.v1+json,
application/vnd.schemaregistry+json, application/json]
kafkastore.ssl.keystore.type = JKS
ssl.truststore.type = JKS
kafkastore.ssl.truststore.password =
access.control.allow.origin =
ssl.truststore.location =
ssl.keystore.password =
port = 8081
kafkastore.ssl.keystore.location =
master.eligibility = true
ssl.client.auth = false
kafkastore.ssl.keystore.password =
```

```
kafkastore.security.protocol = PLAINTEXT
ssl.trustmanager.algorithm =
authentication.method = NONE
request.logger.name = io.confluent.rest-utils.requests
ssl.key.password =
kafkastore.zk.session.timeout.ms = 30000
kafkastore.sasl.mechanism = GSSAPI
kafkastore.sasl.kerberos.ticket.renew.jitter = 0.05
kafkastore.ssl.key.password =
zookeeper.set.acl = false
authentication.roles = [*]
metrics.num.samples = 2
ssl.protocol = TLS
kafkastore.ssl.keymanager.algorithm = SunX509
kafkastore.connection.url = localhost:2181
debug = false
listeners = [http://0.0.0.0:8081]
ssl.provider =
ssl.enabled.protocols = []
shutdown.graceful.ms = 1000
ssl.keystore.location =
ssl.cipher.suites = []
kafkastore.ssl.endpoint.identification.algorithm =
kafkastore.ssl.cipher.suites =
access.control.allow.methods =
kafkastore.sasl.kerberos.min.time.before.relogin = 60000
ssl.keymanager.algorithm =
metrics.sample.window.ms = 30000
kafkastore.init.timeout.ms = 60000
(io.confluent.kafka.schemaregistry.rest.SchemaRegistryConfig:169)
[2017-02-08 09:17:50,800] INFO Initializing KafkaStore with broker endpoints:
PLAINTEXT://mac92.cybage.com:9092 (io.confluent.kafka.schemaregistry.storage.KafkaStore:127)
[2017-02-08 09:17:50,815] WARN Creating the schema topic _schemas using a replication factor of 1,
which is less than the desired one of 3. If this is a production environment, it's crucial to add more
brokers and increase the replication factor of the topic.
(io.confluent.kafka.schemaregistry.storage.KafkaStore:258)
[2017-02-08 09:17:51,663] INFO Initialized last consumed offset to -1
(io.confluent.kafka.schemaregistry.storage.KafkaStoreReaderThread:122)
[2017-02-08 09:17:51,664] INFO [kafka-store-reader-thread-_schemas], Starting
(io.confluent.kafka.schemaregistry.storage.KafkaStoreReaderThread:70)
[2017-02-08 09:17:51,881] INFO Wait to catch up until the offset of the last message at 0
(io.confluent.kafka.schemaregistry.storage.KafkaStore:343)
[2017-02-08 09:17:52,113] INFO Created schema registry namespace localhost:2181/schema_registry
(io.confluent.kafka.schemaregistry.storage.KafkaSchemaRegistry:238)
[2017-02-08 09:17:52,179] INFO Successfully elected the new master:
{"host":"mac92.cybage.com","port":8081,"master_eligibility":true,"version":1}
(io.confluent.kafka.schemaregistry.zookeeper.ZookeeperMasterElector:83)
[2017-02-08 09:17:52,229] INFO Successfully elected the new master:
```

```
{"host":"mac92.cybage.com","port":8081,"master_eligibility":true,"version":1}
(io.confluent.kafka.schemaregistry.zookeeper.ZookeeperMasterElector:83)
[2017-02-08 09:17:52,339] INFO Logging initialized @4316ms (org.eclipse.jetty.util.log:186)
[2017-02-08 09:17:52,409] INFO Adding listener: http://0.0.0.0:8081 (io.confluent.rest.Application:174)
[2017-02-08 09:17:52,470] INFO jetty-9.2.12.v20150709 (org.eclipse.jetty.server.Server:327)
[2017-02-08 09:17:53,395] INFO HV000001: Hibernate Validator 5.1.2.Final
(org.hibernate.validator.internal.util.Version:27)
[2017-02-08 09:17:53,625] INFO Started o.e.j.s.ServletContextHandler@5c153b9e{/null,AVAILABLE}
(org.eclipse.jetty.server.handler.ContextHandler:744)
[2017-02-08 09:17:53,666] INFO Started
NetworkTrafficServerConnector@784c3487{HTTP/1.1}{0.0.0.0:8081}
(org.eclipse.jetty.server.NetworkTrafficServerConnector:266)
[2017-02-08 09:17:53,668] INFO Started @5645ms (org.eclipse.jetty.server.Server:379)
[2017-02-08 09:17:53,669] INFO Server started, listening for requests...
(io.confluent.kafka.schemaregistry.rest.SchemaRegistryMain:45)
```

To test it we can use console command first

```
kafka-console-producer \
  --broker-list localhost:9092 --topic test \
  --property
value.schema='{"type":"record","name":"myrecord","fields":[{"name":"f1","type":"string"}]}'
```

Once started, the process will wait for you to enter messages, one per line, and will send them immediately when you hit the Enter key. Try entering a couple of messages:

```
{"f1": "value1"}
{"f1": "value2"}
{"f1": "value3"}
```

```
[root@mac92 ~]# kafka-console-producer --broker-list localhost:9092 --topic test --
property value.schema='{"type":"record","name":"myrecord","fields":[{"name":"f1","type":"string"}]}'
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/usr/share/java/kafka-serde-tools/slf4j-log4j12-
1.7.6.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/usr/share/java/schema-registry/slf4j-log4j12-
1.7.6.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.slf4j.impl.Log4jLoggerFactory]
{"f1": "value1"}
{"f1": "value2"}
{"f1": "value3"}
```

Now start consumer in separate window

```
[root@mac92 ~]# kafka-console-consumer --topic test \
>     --zookeeper localhost:2181 \
>     --from-beginning
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/usr/share/java/kafka-serde-tools/slf4j-log4j12-
1.7.6.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/usr/share/java/schema-registry/slf4j-log4j12-
1.7.6.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.slf4j.impl.Log4jLoggerFactory]
Using the ConsoleConsumer with old consumer is deprecated and will be removed in a future major
release. Consider using the new consumer by passing [bootstrap-server] instead of [zookeeper].
{"f1":"value1"}
{"f1":"value1"}
{"f1":"value2"}
```

You should see all the messages you created in the previous step written to the console in the same format.

The consumer does not exit after reading all the messages so it can listen for and process new messages as they are published. Try keeping the consumer running and repeating step 5 – you will see messages delivered to the consumer immediately after you hit Enter for each message in the producer.

When you're done, shut down the consumer with Ctrl+C.

Now let's try to produce data to the same topic using an incompatible schema. We'll run the producer with nearly the same command, but change the schema to expect plain integers.

```
kafka-console-producer \
    --broker-list localhost:9092 --topic test \
    --property value.schema='{"type":"int"}'
```

Now if you enter an integer and hit enter, you should see the following (expected) exception:

```
org.apache.kafka.common.errors.SerializationException: Error registering Avro schema: "int"
Caused by: io.confluent.kafka.schemaregistry.client.rest.exceptions.RestClientException:
Schema being registered is incompatible with the latest schema; error code: 409
    at
io.confluent.kafka.schemaregistry.client.rest.utils.RestUtils.httpRequest(RestUtils.java:146)
    at
io.confluent.kafka.schemaregistry.client.rest.utils.RestUtils.registerSchema(RestUtils.java:17
4)
    at
io.confluent.kafka.schemaregistry.client.CachedSchemaRegistryClient.registerAndGetId(CachedSche
maRegistryClient.java:51)
    at
io.confluent.kafka.schemaregistry.client.CachedSchemaRegistryClient.register(CachedSchemaRegis
```

```
tryClient.java:89)
    at
io.confluent.kafka.serializers.AbstractKafkaAvroSerializer.serializeImpl(AbstractKafkaAvroSerializer.java:49)
    at
io.confluent.kafka.formatter.AvroMessageReader.readMessage(AvroMessageReader.java:155)
    at kafka.tools.ConsoleProducer$.main(ConsoleProducer.scala:94)
    at kafka.tools.ConsoleProducer.main(ConsoleProducer.scala)
```

It is because we have change the registered schema with topic

```
kafka-avro-console-producer \
>     --broker-list localhost:9092 --topic test \
>     --property value.schema='{"type":"int"}'
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/usr/share/java/kafka-serde-tools/slf4j-log4j12-
1.7.6.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/usr/share/java/schema-registry/slf4j-log4j12-
1.7.6.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.slf4j.impl.Log4jLoggerFactory]
123
org.apache.kafka.common.errors.SerializationException: Error registering Avro schema: "int"
Caused by: io.confluent.kafka.schemaregistry.client.rest.exceptions.RestClientException: Schema being
registered is incompatible with the latest schema; error code: 409
    at io.confluent.kafka.schemaregistry.client.rest.RestService.sendHttpRequest(RestService.java:170)
    at io.confluent.kafka.schemaregistry.client.rest.RestService.httpRequest(RestService.java:187)
    at io.confluent.kafka.schemaregistry.client.rest.RestService.registerSchema(RestService.java:238)
    at io.confluent.kafka.schemaregistry.client.rest.RestService.registerSchema(RestService.java:230)
    at io.confluent.kafka.schemaregistry.client.rest.RestService.registerSchema(RestService.java:225)
    at
io.confluent.kafka.schemaregistry.client.CachedSchemaRegistryClient.registerAndGetId(CachedSchema
RegistryClient.java:59)
    at
io.confluent.kafka.schemaregistry.client.CachedSchemaRegistryClient.register(CachedSchemaRegistryCli
ent.java:91)
    at
io.confluent.kafka.serializers.AbstractKafkaAvroSerializer.serializeImpl(AbstractKafkaAvroSerializer.java:
72)
    at io.confluent.kafka.formatter.AvroMessageReader.readMessage(AvroMessageReader.java:158)
    at kafka.tools.ConsoleProducer$.main(ConsoleProducer.scala:55)
    at kafka.tools.ConsoleProducer.main(ConsoleProducer.scala)
[root@mac92 ~]#
```

Apache Kafka Tutorial

```

AvroConsumer-V1 : sbt
File Edit View Scrollback Bookmarks Settings Help
[root@localhost AvroConsumer-V1]# sbt run
[info] Set current project to AvroTest (in build file:/root/sbt_proj/javaProj/AvroConsumer-V1/)
[info] Running AvroConsumer
log4j:WARN No appenders could be found for logger (org.apache.kafka.clients.consumer.ConsumerConfig).
log4j:WARN Please initialize the log4j system properly.
log4j:WARN See http://logging.apache.org/log4j/1.2/faq.html#noconfig for more info.
Session id=10001 Channel=HomePage Referrer=null

```



```

AvroProducer-V1 : sbt
File Edit View Scrollback Bookmarks Settings Help
[root@localhost AvroProducer-V1]# sbt run
[info] Set current project to AvroTest (in build file:/root/sbt_proj/javaProj/AvroProducer-V1/)
[info] Running AvroProducer
log4j:WARN No appenders could be found for logger (org.apache.kafka.clients.producer.ProducerConfig).
log4j:WARN Please initialize the log4j system properly.
log4j:WARN See http://logging.apache.org/log4j/1.2/faq.html#noconfig for more info.
Complete
[success] Total time: 3 s, completed Jan 26, 2017 2:20:20 PM
[root@localhost AvroProducer-V1]#

```

Now change the schema, name and type should be the same as old schema

Apache Kafka Tutorial

```

AvroProducer-V2 : vi
File Edit View Scrollback Bookmarks Settings Help
{"type": "record",
 "name": "ClickRecord",
 "fields": [
     {"name": "session_id", "type": "string"},
     {"name": "browser", "type": ["string", "null"]},
     {"name": "campaign", "type": ["string", "null"]},
     {"name": "channel", "type": "string"},
     {"name": "entry_url", "type": ["string", "null"], "default": "None"},
     {"name": "ip", "type": ["string", "null"]},
     {"name": "language", "type": ["string", "null"], "default": "None"},
     {"name": "os", "type": ["string", "null"], "default": "None"}
]

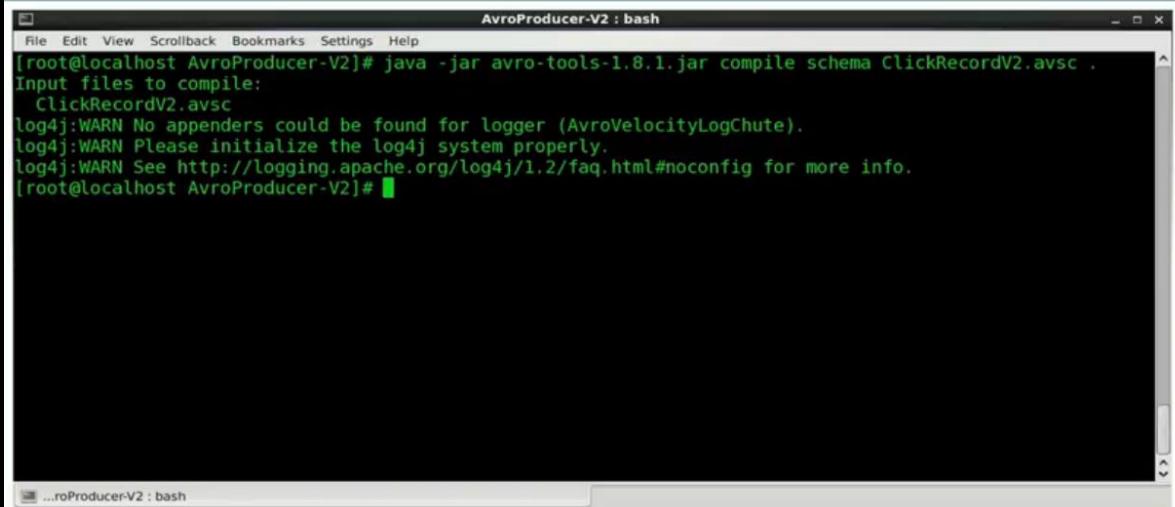
```

You cannot change the schema randomly there are some rule for avro schema resolution check below link

<https://avro.apache.org/docs/current/spec.html#Schema+Resolution>

Generate class for new schema

Apache Kafka Tutorial



```
File Edit View Scrollback Bookmarks Settings Help
[root@localhost AvroProducer-V2]# java -jar avro-tools-1.8.1.jar compile schema ClickRecordV2.avsc .
Input files to compile:
ClickRecordV2.avsc
log4j:WARN No appenders could be found for logger (AvroVelocityLogChute).
log4j:WARN Please initialize the log4j system properly.
log4j:WARN See http://logging.apache.org/log4j/1.2/faq.html#noconfig for more info.
[root@localhost AvroProducer-V2]#
```

Create new producer

```
import java.util.*;
import org.apache.kafka.clients.producer.*;
public class ClickRecordProducerV2 {

    public static void main(String[] args) throws Exception{

        String topicName = "AvroClicks";
        String msg;

        Properties props = new Properties();
        props.put("bootstrap.servers", "localhost:9092,localhost:9093");
        props.put("key.serializer", "org.apache.kafka.common.serialization.StringSerializer");
        props.put("value.serializer", "io.confluent.kafka.serializers.KafkaAvroSerializer");
        props.put("schema.registry.url", "http://localhost:8081");

        Producer<String, ClickRecord> producer = new KafkaProducer<>(props);
        ClickRecord cr = new ClickRecord();
        try{
            cr.setSessionId("10001");
            cr.setChannel("HomePage");
            cr.setIp("192.168.0.1");
            cr.setLanguage("Spanish");
            cr.setOs("Mac");
            cr.setEntryUrl("http://facebook.com/myadd");
        }
    }
}
```

```

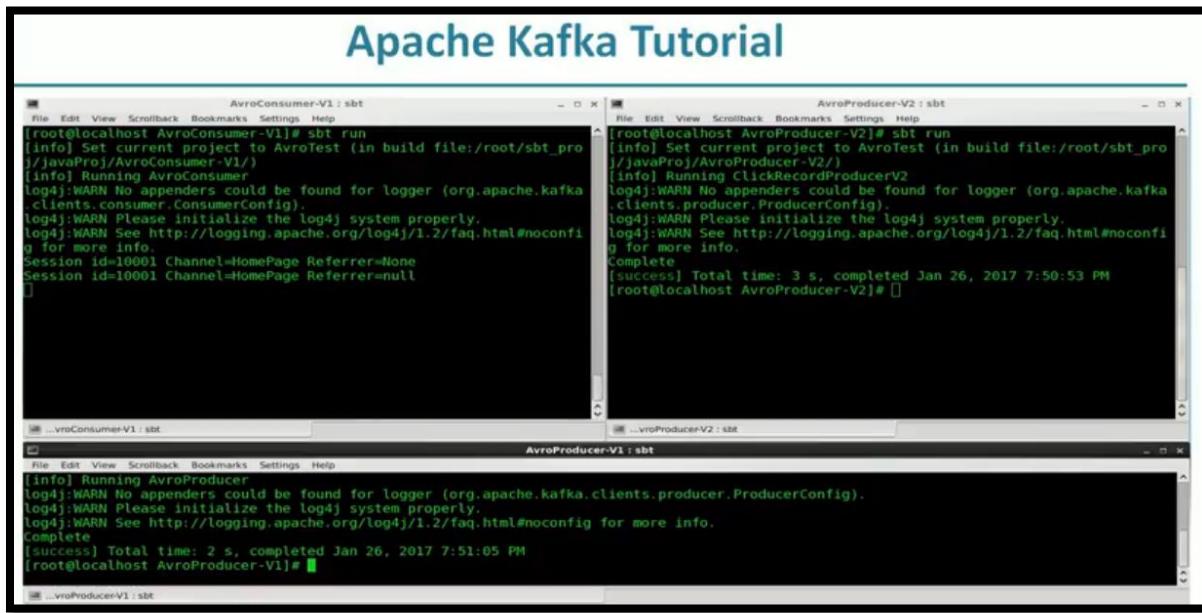
        producer.send(new ProducerRecord<String,
ClickRecord>(topicName,cr.getSessionId().toString(),cr)).get();

        System.out.println("Complete");
    }
    catch(Exception ex){
        ex.printStackTrace(System.out);
    }
    finally{
        producer.close();
    }

}
}

```

Code is similar just we use new ClickRecord.java class



Old consumer can read both schema but it will produce null values for new added field we can create new consumer to read new fields

```

import java.util.*;
import org.apache.kafka.clients.consumer.*;

public class ClickRecordConsumerV2{

```

```
public static void main(String[] args) throws Exception{

    String topicName = "AvroClicks";

    String groupName = "RG";
    Properties props = new Properties();
    props.put("bootstrap.servers", "localhost:9092,localhost:9093");
    props.put("group.id", groupName);
    props.put("key.deserializer", "org.apache.kafka.common.serialization.StringDeserializer");
    props.put("value.deserializer", "io.confluent.kafka.serializers.KafkaAvroDeserializer");
    props.put("schema.registry.url", "http://localhost:8081");
    props.put("specific.avro.reader", "true");

    KafkaConsumer<String, ClickRecord> consumer = new KafkaConsumer<>(props);
    consumer.subscribe(Arrays.asList(topicName));
    try{
        while (true){
            ConsumerRecords<String, ClickRecord> records = consumer.poll(100);
            for (ConsumerRecord<String, ClickRecord> record : records){
                System.out.println("Session id=" + record.value().getSessionId()
                    + " Channel=" + record.value().getChannel()
                    + " Entry URL=" + record.value().getEntryUrl()
                    + " Language=" + record.value().getLanguage());
            }
        }
    }catch(Exception ex){
        ex.printStackTrace();
    }
    finally{
        consumer.close();
    }
}
}
```

Apache Kafka Tutorial

Thing to Cover

- ➡ 1. Restful clients and interface to Python.
2. Integration with other systems like Apache Spark.
3. Kafka Connect and relevant prebuilt connectors.
4. Kafka Stream and APIs.
5. Cluster Architecture, Security, Operations, and Tuning.

<https://www.safaribooksonline.com/library/view/kafka-the-definitive/9781491936153/ch02.html>